

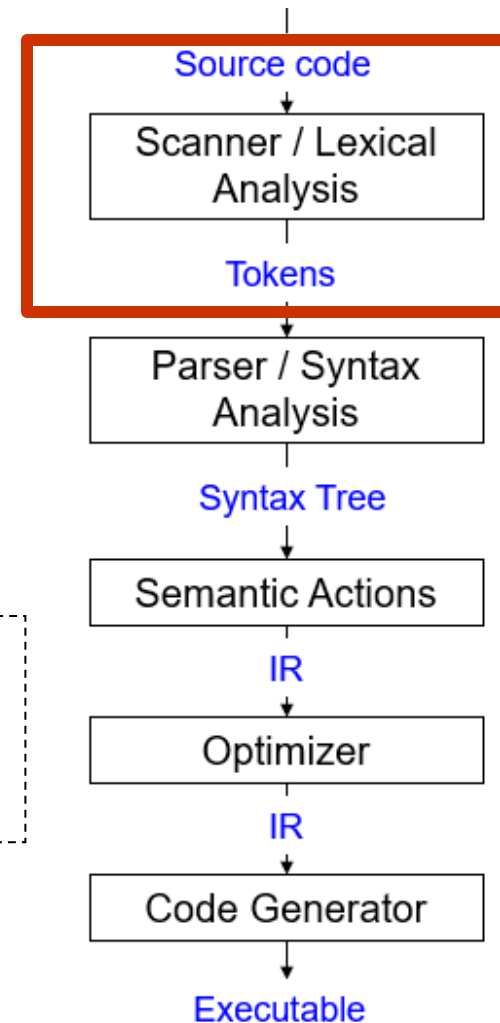
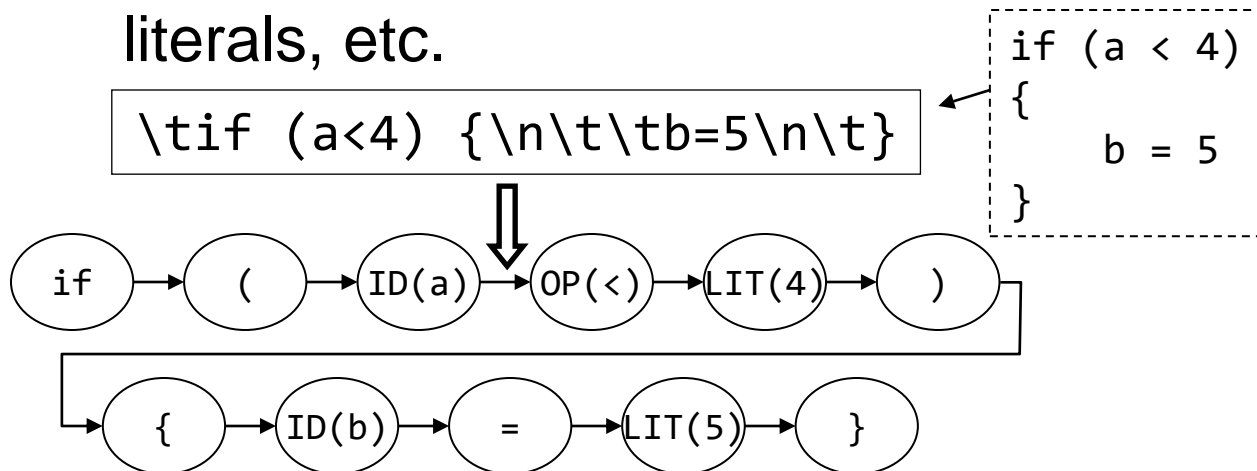
CS323: Compilers

Spring 2023

Week 2: Scanners

Scanner - Overview

- Also called lexers / lexical analyzers
- Recall: scanners
 - See program text as a stream of letters
 - break input stream up into a set of tokens: Identifiers, reserved words, literals, etc.



Scanner - Motivation

- Why have a separate scanner when you can combine this with syntax analyzer (parser)?
 - Simplicity of design
 - E.g. rid parser of handling whitespaces
 - Improve compiler efficiency
 - E.g. sophisticated buffering algorithms for reading input
 - Improve compiler portability
 - E.g. handling ^M character in Linux (CR+LF in Windows)

Scanner - Tasks

1. Divide the program text into *substrings* or *lexemes*
 - place dividers
2. Identify the *class* of the substring identified
 - Examples of predefined categories: Identifiers, keywords, operators, etc.
 - Identifier – *strings of letters or digits starting with a letter*
 - Integer – *non-empty string of digits*
 - Keyword – *“if”, “else”, “for”* etc.
 - Blankspace - *\t, \n, ‘ ‘*
 - Operator – *(,), <, =, etc.*
 - *Observation:* substrings can be categorized i.e. follow some pattern

Categorizing a Substring (English Text)

- What is the English language analogy for *class*?
 - Noun, Verb, Adjective, Article, etc.
 - In an English essay, each of these classes can have a set of strings.
 - Similarly, in a program, each class can have a set of substrings.

Exercise

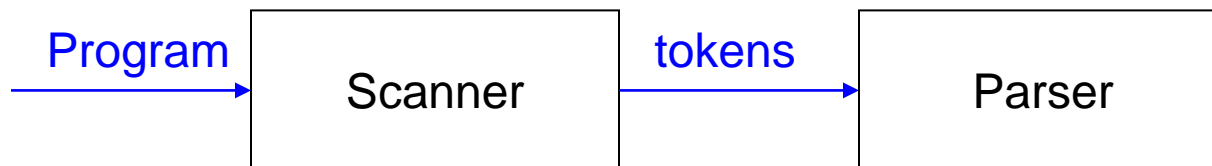

- How many tokens of class *identifier* exist in the code below?

```
for(int i=0;i<10;i++) {  
    printf("hello");  
}
```

Scanner Output

- A token corresponding to each lexeme
 - Token is a pair: <class, value>

A string / lexeme / substring of program text



E.g. `int x = 0;`

(Keyword, "int"),
(Identifier, "x"),
("="),
(Integer, "0"),
(";")

Scanners – interesting examples

- Fortran (white spaces are ignored)

DO 5 I = 1,25 ← DO Loop

DO 5 I = 1.25 ← Assignment statement

- PL/1 (keywords are not reserved)

DECLARE (ARG1, ARG2, . . ., ARGN);

- C++

Nested template: Quad<Square<Box>> b;

Stream input: std::cin >> bx;

Scanners – interesting examples (discussion)

- How did we go about recognizing tokens in previous examples?
 - Scan **left-to-right** till a token is identified
 - **One token at a time**: continue scanning the remaining text till the next token is identified...
 - So on...

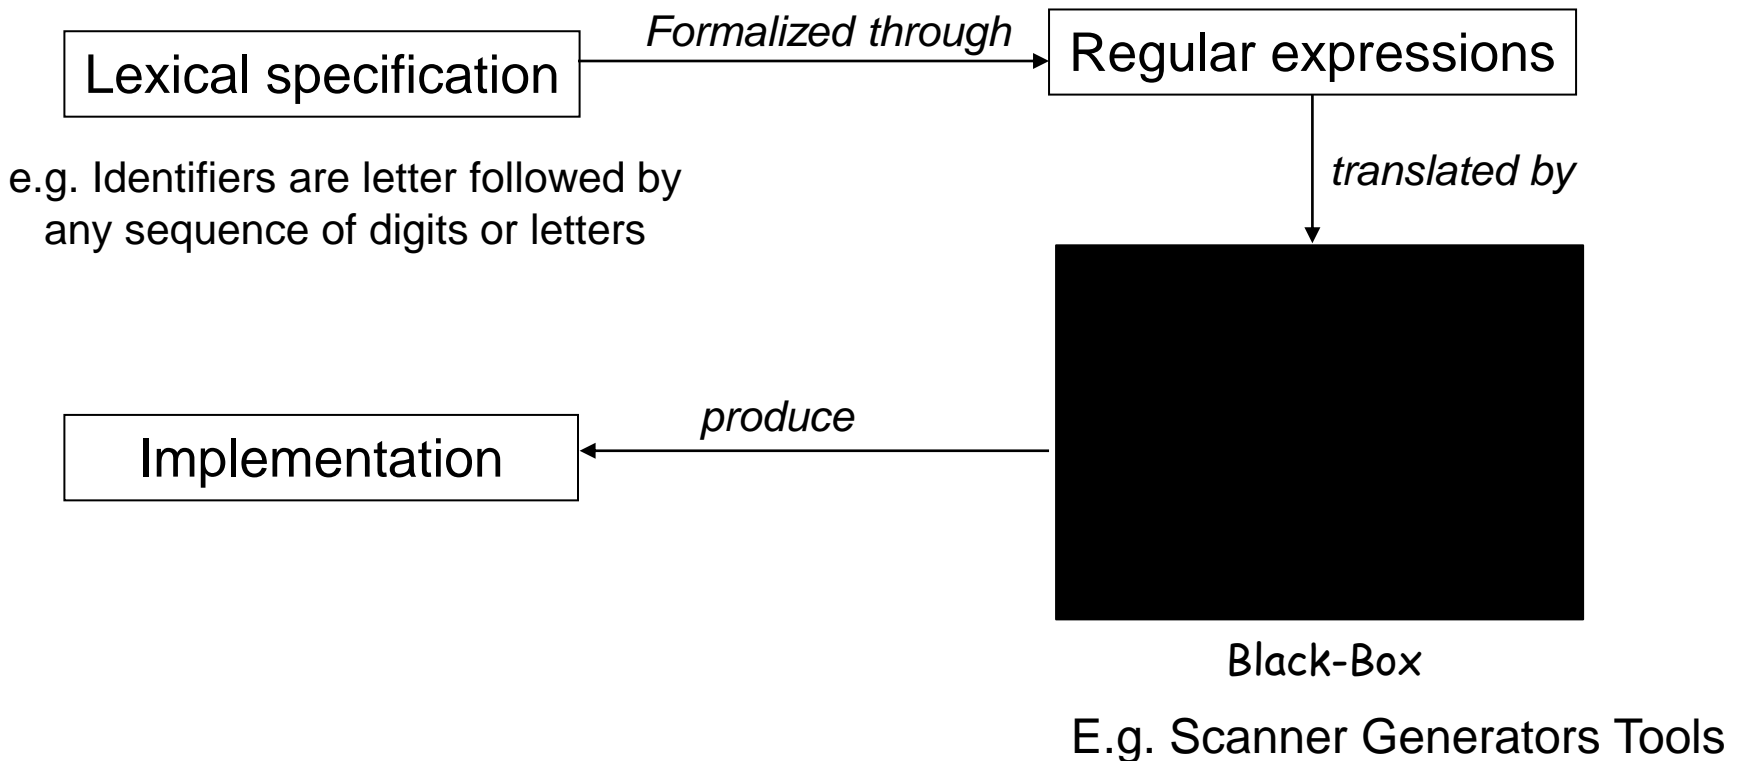
We always need to *look-ahead* to identify tokens

....but we want to minimize the amount of look-ahead done to simplify scanner implementation

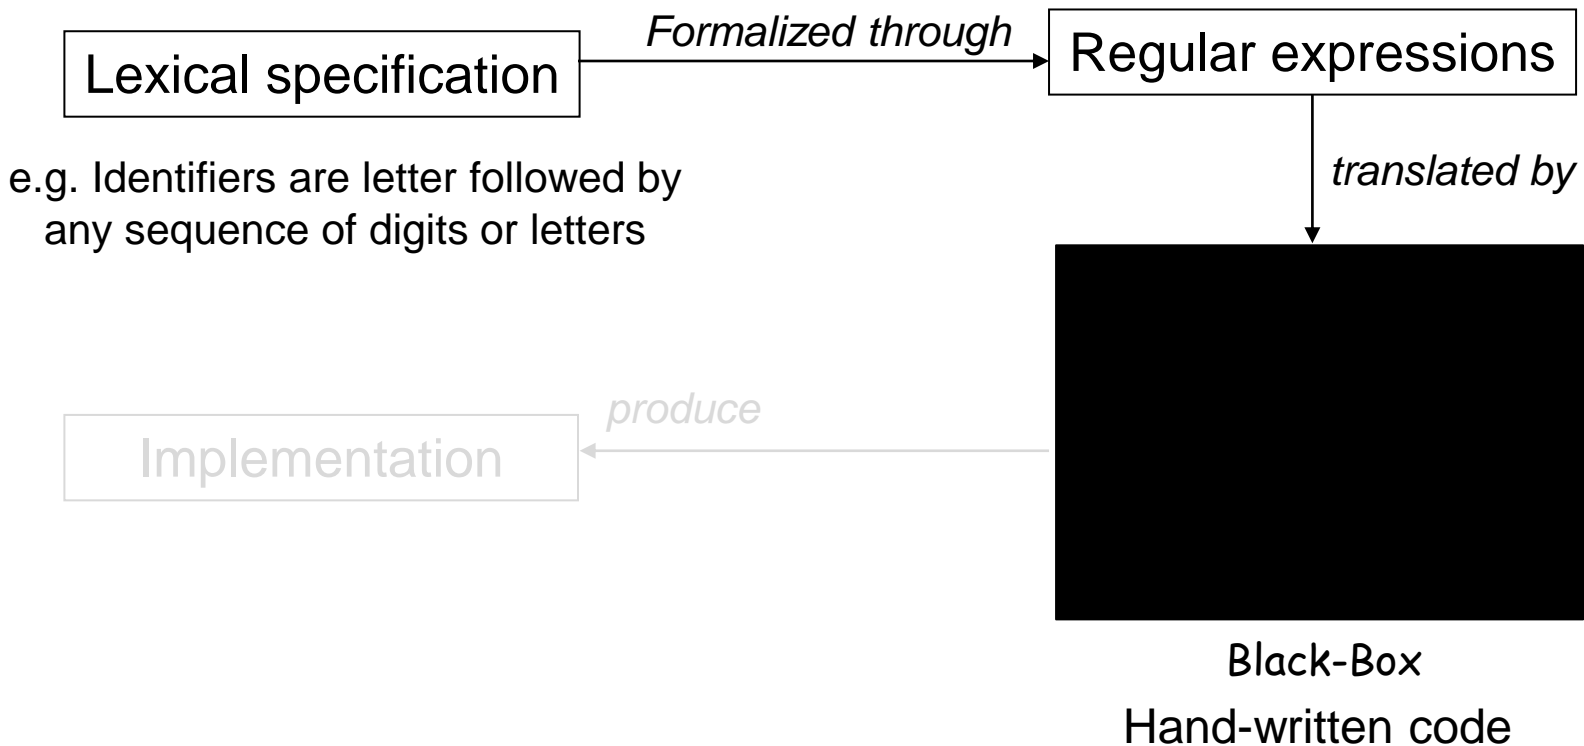
Scanners – what do we need to know?

1. How do we define tokens?
 - Regular expressions
2. How do we recognize tokens?
 - build code to find a lexeme that is a prefix and that belongs to one of the classes.
3. How do we write lexers?
 - E.g. use a lexer generator tool such as Flex

Scanner / Lexical Analyzer - flowchart



Scanner / Lexical Analyzer - flowchart



Scanner Generators

- Essentially, tools for converting regular expressions into scanners
 - Lex (Flex) generates C/C++ scanner program
 - ANTLR (ANother Tool for Language Recognition) generates Java program for translating program text (JFlex is a less popular option)
 - Pylexer is a Python-based lexical analyzer (**not a scanner generator**). *It just scans input, matches regexps, and tokenizes. Doesn't produce any program.*

Regular Expressions

- Used to define the structure of tokens
- Regular sets:
 - Informal:** a set of strings defined by regular expressions
 - Formal:** a language that can be defined by regular expressions

Start with a finite *character set* or *Vocabulary* (V). Strings are formed using this character set with the following rules:

Regular Expressions

- Strings are regular sets (with one element): pi 3.14159
 - So is the empty string: λ (ϵ instead)
- Concatenations of regular sets are regular: pi3.14159
 - To avoid ambiguity, can use () to group regexps together
- A choice between two regular sets is regular, using |: (pi|3.14159)
- 0 or more of a regular set is regular, using *: (pi)*
- other notation used for convenience:
 - Use **Not** to accept all strings except those in a regular set
 - Use ? to make a string optional: x? equivalent to (x| λ)
 - Use + to mean 1 or more strings from a set: x+ equivalent to xx*
 - Use [] to present a range of choices: [1-3] equivalent to (1|2|3)

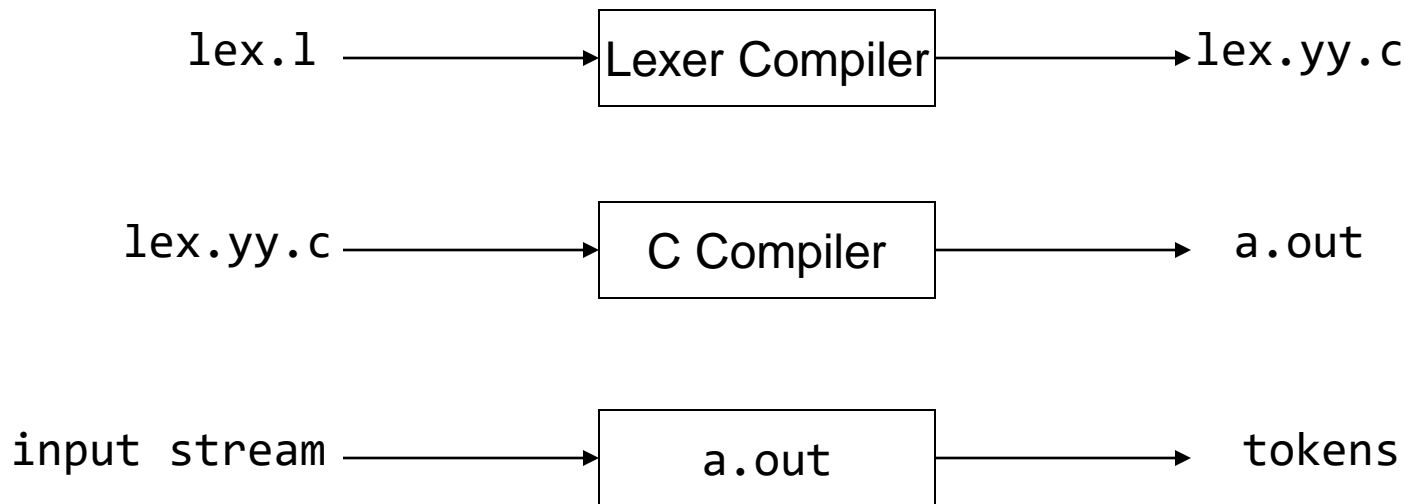
Regular Expressions for Lexical Specifications

- Digit: $D = (0|1|2|3|4|5|6|7|8|9)$
- Letter: $L = [A-Za-z]$ alternative definition: $[0-9]$
- Literals (integers or floats): $-?D+(.D^*)?$
- Identifiers: $(_|L)(_|L|D)^*$
- Comments (as in Micro): $//\text{Not}(\backslash n)^*\backslash n$
- More complex comments (delimited by `##`, can use `#` inside comment):
$$## \ (\ (#|\backslash) \text{Not}(\#))^* \ ##$$

Lex (Flex)

- Commonly used Unix scanner generator (superseded by Flex)
- Flex is a domain specific language for writing scanners
- Features:
 - Character classes : define sets of characters (e.g., digits)
 - Token definitions : `regex {action to take}`

Lex (Flex)



Lex (Flex)

- Format of lex.l (3 parts separated by %%)
format: **name definition**

Declarations

e.g. **DIGIT [0-9]**

Refer to DIGIT here
using {} braces {DIGIT}
expands to ([0-9])

%%

Translation rules

format: **pattern action**

e.g. **{DIGIT}+ {printf("INTLITERAL");}**

User code mentioned here copied as is to lex.yy.c
Auxiliary functions

Example: Lex (Flex)

```
DIGIT    [ 0-9 ]  
ID       [ a-z ] [ a-z 0-9 ] *
```

```
%%
```

```
{DIGIT}+ {  
    printf( "An integer: %s (%d)\n", yytext,  
        atoi( yytext ) );  
}
```

```
{DIGIT}+"."{DIGIT}* {  
    printf( "A float: %s (%g)\n", yytext,  
        atof( yytext ) );  
}
```

```
if|then|begin|end|procedure|function {  
    printf( "A keyword: %s\n", yytext );  
}
```

```
{ID}      printf( "An identifier: %s\n", yytext );
```

Lex (Flex)

- The order in which tokens are defined matters!
- Lex will match the longest possible token
 - “ifa” becomes ID(ifa), not IF ID(a)
- If two regexes both match, Lex uses the one defined first
 - “if” becomes IF, not ID(if)
- Use action blocks to process tokens as necessary
 - Convert integer/float literals to numbers
 - Remove quotes from string literals

Demo

Documentation

- [Flex \(manual web-version\):](#)
- [Lexical Analysis With Flex, for Flex 2.6.2: Top \(westes.github.io\)](#)
- [Lex - A Lexical Analyzer Generator \(compilertools.net\)](#)
- [ANTLR](#)

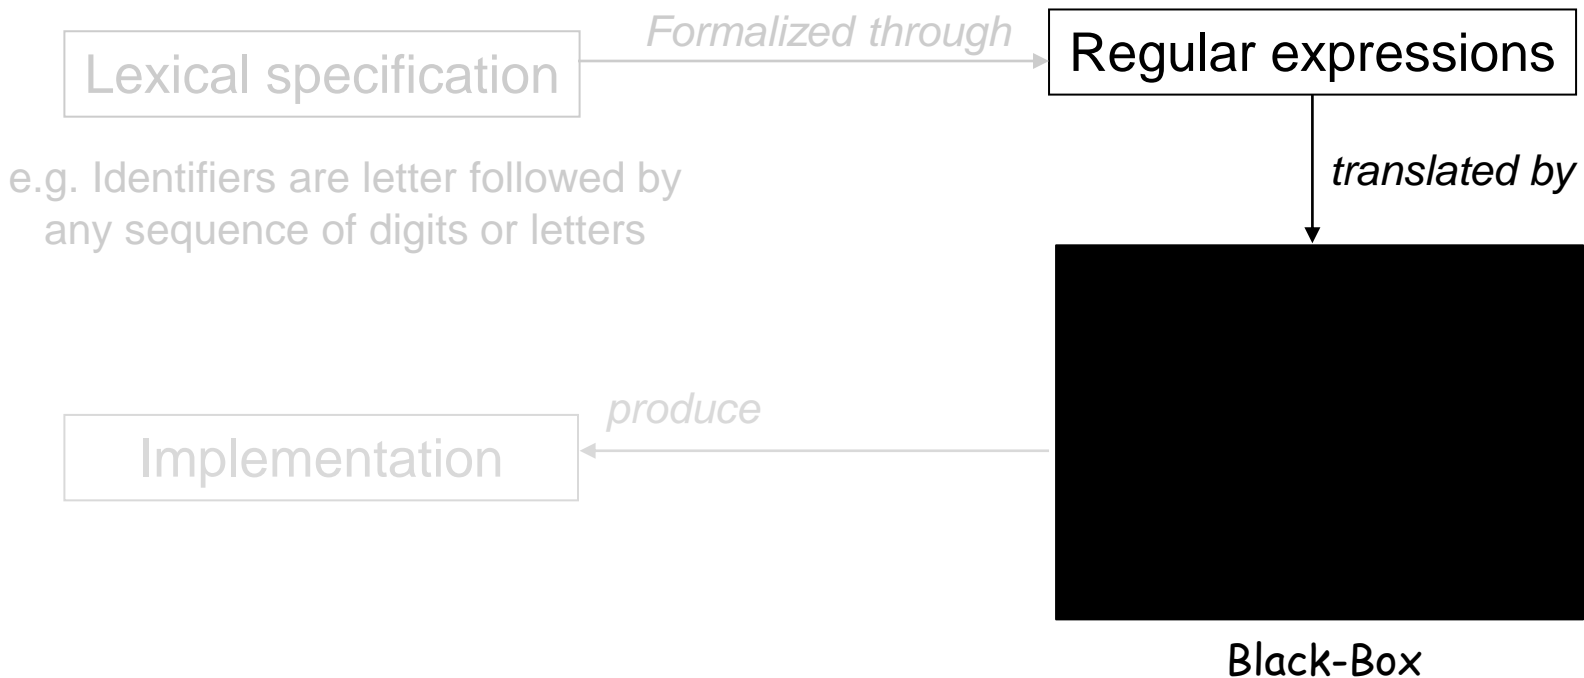
Summary

- We saw what it takes to write a scanner:
 - Specify how to identify token classes (using regexps)
 - Convert the **regexps to code** that identifies a *prefix* of the input program text as a *lexeme* matching one of the token classes
 - Use tools for automatic code generation (e.g. Flex / ANTLR)
 - *How do the tools convert regexps to code?* **Finite Automata**

OR

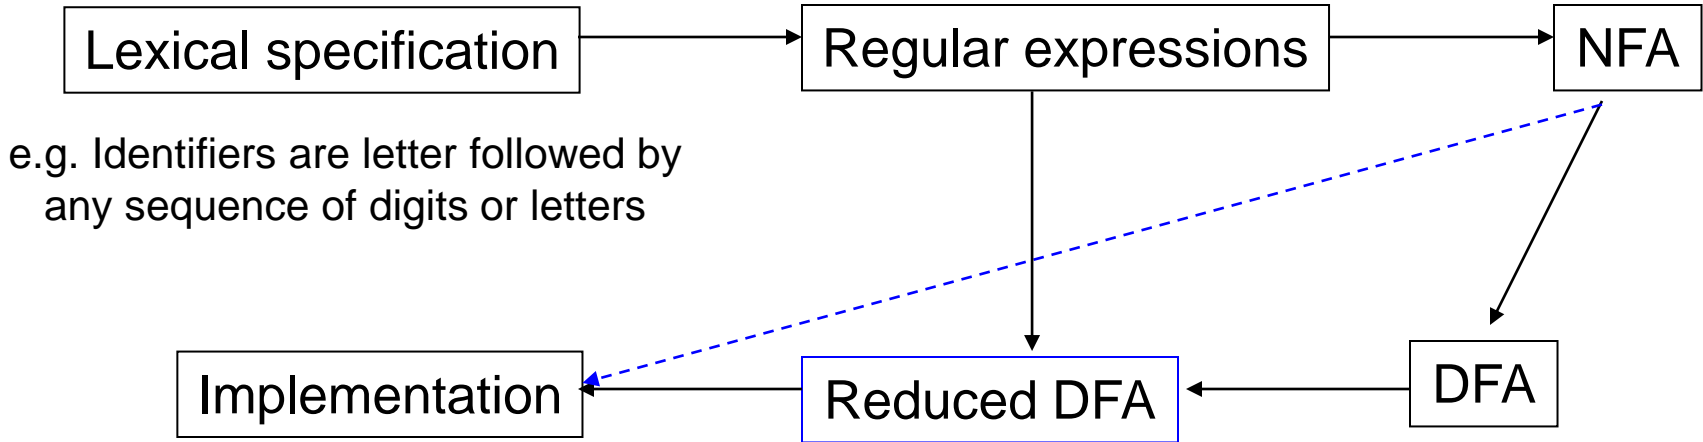
- Scanner code manually

Scanner- Implementation



How does a tool such as Flex generate code?

Scanner - flowchart




Finite Automata

- Another formal way to describe sets of strings (just like regular expressions)
- Also known as finite state machines / automata
- Reads a string, either recognizes it or not
- Two Features:
 - **State**: initial, matching / final / accepting, non-matching
 - **Transition**: a move from one state to another

Finite Automata


- Regular expressions and FA are equivalent*



Exercise: what is the equivalent regular expression for this FA?

λ transitions


- Transitions between states that aren't triggered by seeing another character
 - Can *optionally* take the transition, but do not have to
 - Can be used to link states together



Think of this as an arrow to a state without a label

Non-deterministic Finite Automata

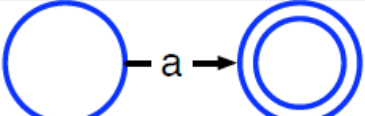
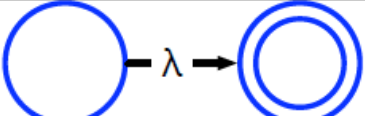
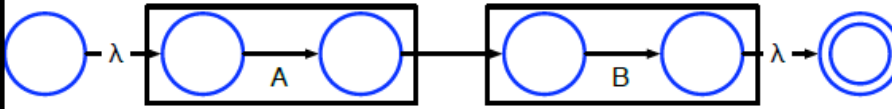
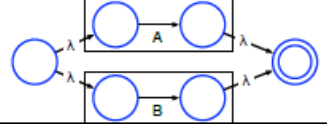
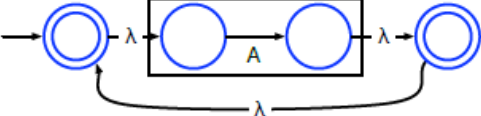
- A FA is non-deterministic if, from one state reading a single character could result in transition to multiple states (or has λ transitions)
- Sometimes regular expressions and NFAs have a close correspondence



\equiv

$a(bb)^+a$

Building a FA from a regexp

Expression	FA
a	
λ	
AB	
A B	
A*	


Mini-exercise: how do we build an FA that accepts Not(A)?

What about A? (? as in optional)

“Running” an NFA


- Intuition: take every possible path through an NFA
 - Think: parallel execution of NFA
 - Maintain a “pointer” that tracks the current state
 - Every time there is a choice, “split” the pointer, and have one pointer follow each choice
 - Track each pointer simultaneously
 - If a pointer gets stuck, stop tracking it
 - If any pointer reaches an accept state at the end of input, accept

Running an NFA - Example



- NFAs are concise but slow
- Example:
 - Running the NFA for input string abbb requires exploring all execution paths

Running an NFA - Example



- NFAs are concise but slow
- Example:
 - Running the NFA for input string abbb requires exploring all execution paths
 - **Optimization: run through the execution paths in parallel**
 - *Complicated. Can we do better?*


Deterministic Finite Automata

- Each possible input character read leads to at most one new state
 - Can convert NFAs to *deterministic* finite automata (DFAs)
 - No choices — never a need to “split” pointers
 - Initial idea: simulate NFA for all possible inputs, any time there is a new configuration of pointers, create a state to capture it
 - Pointers at states 1, 3 and 4 → new state {1, 3, 4}
 - Trying all possible inputs is impractical; instead, for any new state, explore all possible *next* states (that can be reached with a single character)
 - Process ends when there are no new states found
 - This can result in very large DFAs!

DFA reduction

- DFAs built from NFAs are not necessarily optimal
- May contain many more states than is necessary


$$(ab)^+ \equiv (ab)(ab)^*$$



DFA reduction

- DFAs built from NFAs are not necessarily optimal
- May contain many more states than is necessary

$$(ab)^+ \equiv (ab)(ab)^*$$



DFA reduction

- Intuition: merge equivalent states
 - Two states are equivalent if they have the same transitions to the same states
- Basic idea of optimization algorithm
 - Start with two big nodes, one representing all the final states, the other representing all other states
 - Successively split those nodes whose transitions lead to nodes in the original DFA that are in different nodes in the optimized DFA


Implementation

- While doing lexical analysis, we need extensions to regular expressions
 - Match as long a substring as possible
 - Handle errors
- Good algorithms for substring matching
 - Require only a single pass over the input
 - Using Tries
 - Few operations per character
 - Table look-up method


Implementation: Transition Tables

- A table encodes states and transitions of FA
 - 1 row per state
 - 1 column per character in the alphabet
 - Table entry: state (label)

State / Character	a	b	c
1	1	3	2
2	-	3	-
3	-	-	3




Example 1




NFA OR DFA?

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3
2	3	-	4

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	7	5	-

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	7	5	-
6,7	-	6,7	6,7

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	7	5	-
6,7	-	6,7	6,7
7	-	6	6

Example 1: NFA -> DFA




State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	7	5	-
6,7	-	6,7	6,7
7	-	6	6
6	-	7	7

Example 1: DFA




State	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	7	5	-
6,7	-	6,7	6,7
7	-	6	6
6	-	7	7

Example 2: NFA -> DFA




NFA OR DFA?

Example 2: NFA -> DFA




State/ char	0	1	Final ?	Comments
A	{A, B}	A	No	In state A, on seeing input 0, we have a choice to go to either state A or B

Example 2: NFA -> DFA




State/ char	0	1	Final ?	Comments
A	{A, B}	A	No	In state A, on seeing input 0, FA gives us a choice to go to either state A or state B
A,B	{A,B,C}	A	No	In state A,B we have two component states A and B. From A on input 0, FA takes us to states A and B. From B on 0 FA takes us to C. So, the set of states reachable from A,B on input 0 is A,B,C. Similarly, for input 1, from A FA takes us to A. From B on input 1, FA gets stuck in an error state.

Example 2: NFA -> DFA



State/ char	0	1	Final ?	Comments
A	{A, B}	A	No	In state A, on seeing input 0, FA gives us a choice to go to either state A or state B
A,B	{A,B,C}	A	No	In state A,B we have two component states A and B. From A on input 0, FA takes us to states A and B. From B on 0 FA takes us to C. So, the set of states reachable from A,B on input 0 is A,B,C. Similarly, for input 1, from A FA takes us to A. From B on input 1, FA gets stuck in an error state.
A,B,C	{A,B,C}	A	Yes	One of the component states C is final in the FA. Hence, A,B,C is a final state.


Example 2: NFA -> DFA



State/ char	0	1	Final ?	Comments
A	{A, B}	A	No	In state A, on seeing input 0, FA gives us a choice to go to either state A or state B
A,B	{A, B, C}	A	No	In state A, B we have two component states A and B. From B on 0 FA takes us to C. So, the set of states reachable from A,B on input 0 is A,B,C. Similarly, for input 1, from A FA takes us to A. From B on input 1, FA gets stuck in an error state.
A,B,C	{A,B,C}	A	Yes	One of the component states C is final in the FA. Hence, A,B,C is a final state.


Should we consider states B and C in the table?

Example 2: DFA



State/ char	0	1	Final ?
A	{A, B}	A	No
A,B	{A,B,C}	A	No
A,B,C	{A,B,C}	A	Yes

Example 1: DFA



State	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	7	5	-
6,7	-	6,7	6,7
7	-	6	6
6	-	7	7

What states can be merged?

Example 1: Reduced DFA

What states can be merged?

State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	7	5	-
6,7	-	6,7	6,7
7	-	6	6
6	-	7	7

Example: Reduced DFA

What states can be merged?

Definition 8 (Equivalence of states) *Let $M = (Q, \Sigma, \delta, q_0, F)$ be a DFA. We say that two states $p, q \in Q$ are equivalent, and we write it $p \equiv q$, if for every string $x \in \Sigma^*$ the state that M reaches from p given x is accepting if and only if the state that M reaches from q given x is accepting.*

Definition 8 pic source: <https://people.eecs.berkeley.edu/~luca/cs172/notemindfa.pdf>

State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	7	5	-
6,7	-	6,7	6,7
7	-	6	6
6	-	7	7

Example: Reduced DFA

What states can be merged?

6 and 7

State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7	4	-
3,4	6,7	3,4	5
5	6_7_M	5	-
6,7	-	6,7	6,7
6_7_M	-	6_7_M	6_7_M

Example: Reduced DFA

What states can be merged?

6,7 and 6_7_M

State / Char	a	b	c
1	2	-	3
2	3	-	4
3	-	3,4	5
4	6,7_6_7_M	4	-
3,4	6,7_6_7_M	3,4	5
5	6,7_6_7_M	5	-
6,7_6_7_M	-	6,7_6_7_M	6,7_6_7_M

Example: Reduced DFA

What states can be merged?

4 and 5

State / Char	a	b	c
1	2	-	3
2	3	-	4_5_M
3	-	3,4	4_5_M
4_5_M	6,7_6_7_M	4_5_M	-
3,4	6,7_6_7_M	3,4	4_5_M
6,7_6_7_M	-	6,7_6_7_M	6,7_6_7_M

Example: Reduced DFA

