

CS406: Compilers

Spring 2022

Week 9: IR Code for Functions, Local Optimizations

Slides Acknowledgements: Milind Kulkarni

Functions Typical Syntax and Usage

```
FUNCTION VOID bar(INT x, FLOAT y) BEGIN
```

```
END
```

Keywords

Return type

comma separated parameter declarations.

Declarations (string or variable decl) followed by statement declarations

```
FUNCTION void foo() BEGIN
```

```
INT a;
```

```
FLOAT b;
```

```
...
```

```
bar(a, b);
```

← Calls function bar

```
END
```

Terms

```
void foo() {  
    int a, b;  
    ...  
    bar(a, b);  
}
```

```
void bar(int x, int y) {  
    ...  
}
```

- foo is the *caller*
- bar is the *callee*
- a, b are the *actual parameters* to bar
- x, y are the *formal parameters* of bar
- Shorthand:
 - *argument* = actual parameter
 - *parameter* = formal parameter

Different Kinds of Parameters

- Value
- Reference
- Result
- Value-Reference
- Read-only
- Call-by-Name

Value parameters

- “Call-by-value”
- Used in C, Java, default in C++
- Passes the value of an argument to the function
- Makes a copy of argument when function is called
- Advantages? Disadvantages?

Advantage: ‘side-effect’ free – caller can be sure that the argument is not modified by the callee

Disadvantage: Not efficient for larger sized arguments.

Value parameters

```
int x = 1;  
void main () {  
    foo(x, x);  
    print(x);  
}
```

```
void foo(int y, int z) {  
    y = 2;  
    z = 3;  
    print(x);  
}
```

- What do the print statements print?

Reference parameters

- “Call-by-reference”
- Optional in Pascal (use “var” keyword) and C++ (use “&”)
- Pass the *address* of the argument to the function
- If an argument is an expression, evaluate it, place it in memory and then pass the address of the memory location
- Advantages? Disadvantages?

Advantage: Efficiency – for larger sized arguments

Disadvantage: results in clumsy code at times (e.g. check for null pointers)

Reference parameters

```
int x = 1;
void main () {
    foo(x, x);
    print(x);
}
```

```
void foo(int &y, int &z) {
    y = 2;
    z = 3;
    print(x);
    print(y);
}
```

- What do the print statements print?

Result Parameters

- To capture the return value of a function
- Copied at the end of function into arguments of the caller
- E.g. output ports in Verilog module definitions

Result Parameters

```
int x = 1
void main () {
    foo(x, x);
    print(x);
}
```

- What do the print statements print?

```
void foo(int y, result int z) {
    y = 2;
    z = 3;
    print(x);
}
```

Value-Result Parameters

- “Copy-in copy-out”
- Evaluate argument expression, copy to parameters
- After subroutine is done, copy values of parameters back into arguments
- Results are often similar to pass-by-reference, but there are some subtle situations where they are different

Value-Result Parameters

```
int x = 1
void main () {
    foo(x, x);
    print(x);
}
```

- What do the print statements print?

```
void foo(int y, value result int z)
{
    y = 2;
    z = 3;
    print(x);
}
```

Read-only Parameters

- Used when callee will not change value of parameters
- Read-only restriction must be enforced by compiler
- E.g. `const` parameter in C/C++
- Enforcing becomes tricky when in the presence of aliasing and control flow. E.g.

```
void foo(readonly int x, int y) {  
    int * p;  
    if (...) p = &x else p = &y  
    *p = 4  
}
```

Call-by-name Parameters

- The arguments are passed to the function before evaluation
 - Usually, we evaluate the arguments before passing them
- Not used in many languages, but Haskell uses a variant

```
int x = 1
void main () {
    foo(x+2);
    print(x);
}
```

```
void foo(int y) {
    z = y + 3; //expands to z = x + 2 + 3
    print(z);
}
```

Call-by-name Parameters

- Why is this useful?
 - E.g. to analyze certain properties of a program/function – termination

```
void main () {  
    foo(bar());  
}
```

```
void foo(int y) {  
    z = 3;  
    if(z > 3)  
        z = y + z;  
}
```

- Even if bar has an infinite loop, the program terminates.

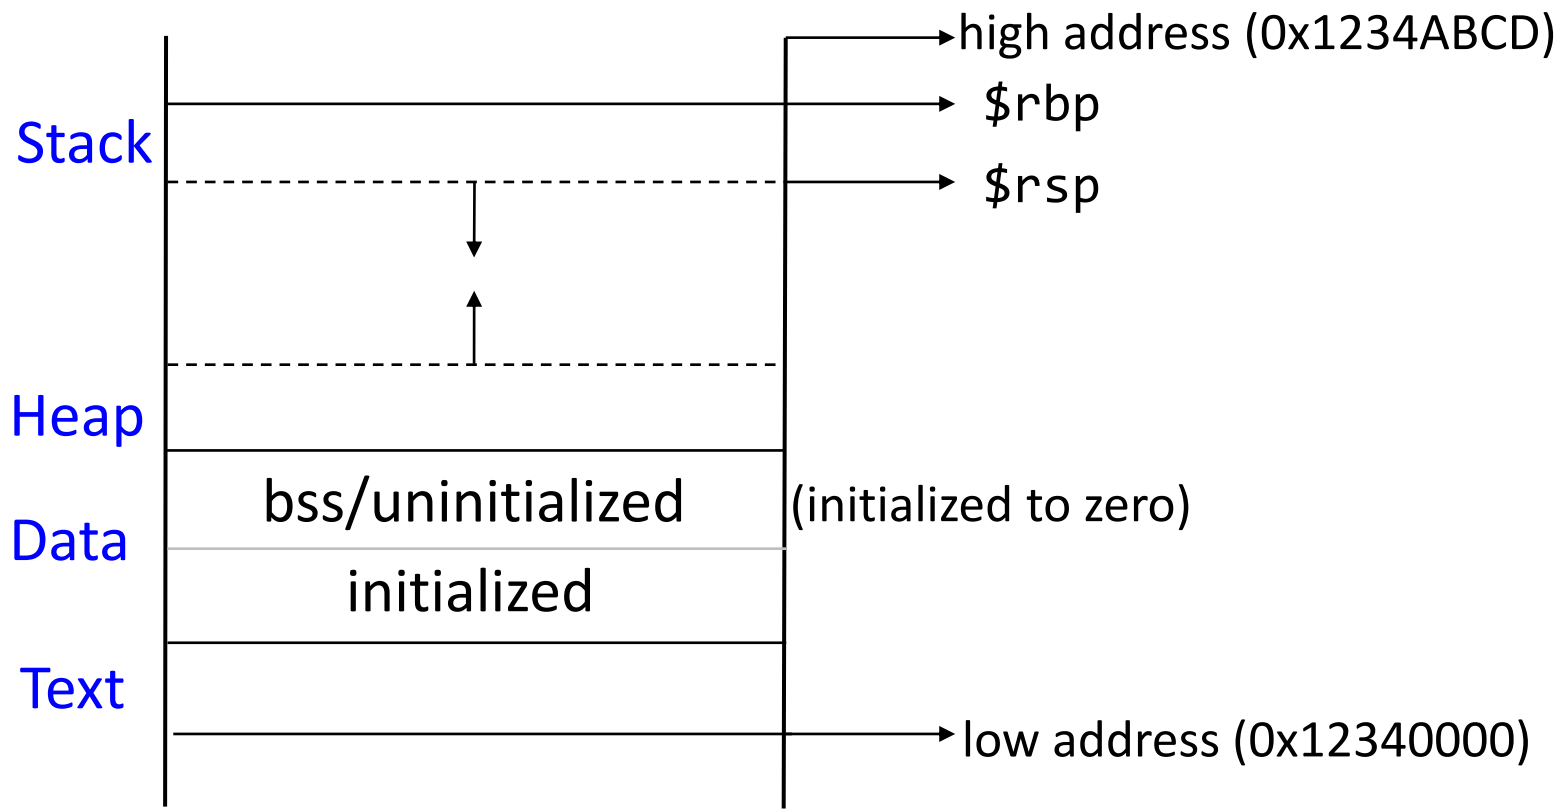
Program Layout in Memory

- Compiler assumes a *runtime environment* for execution of the program.
- A C/C++ program in Linux OS has 4 segments of memory
 - Every memory location is a *box* holding *data/instruction*

Program Layout in Memory

- A program's memory space is divided into four segments:
 1. Text
 - source code of the program
 2. Data
 - Broken into *uninitialized* and *initialized* segments; contains space for global and static variables. E.g. `int x = 7; int y;`
 3. Heap
 - Memory allocated using `malloc/calloc/realloc/new`
 4. Stack
 - Function arguments, return values, local variables, [special registers](#).

Program Layout in Memory



Activation

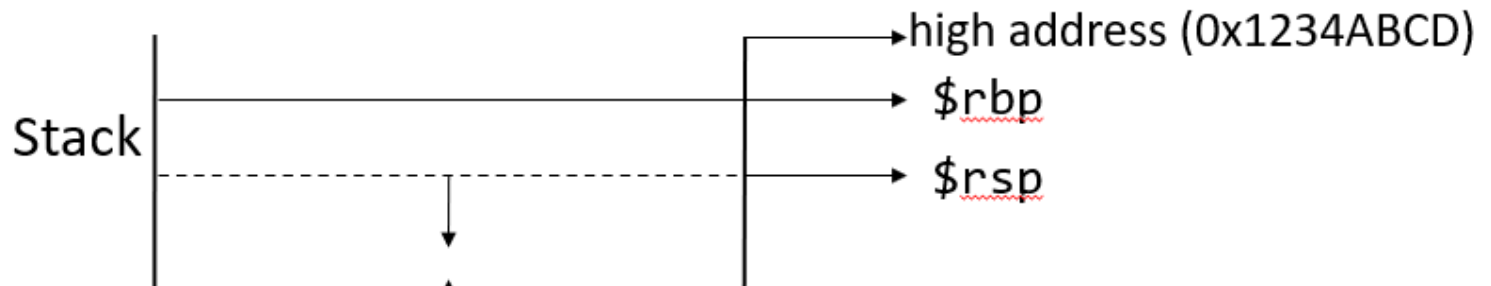
- A function call or invocation is termed an *activation*
- Calls to functions in a program form *activation tree*
 - Postorder traversal of the tree shows return sequence i.e. the order in which control returns from functions
 - Preorder traversal of the tree shows calling sequence
- In a sequential program, at any point in time, *control of execution is in any one activation*
 - All the ancestors of that activation are active i.e. have not returned

Activation

- Activations are managed through the help of *control stack*
- A function call (activation) results in allocating a chunk of memory called *activation record* or *frame* on the stack (also called *stack frame*)

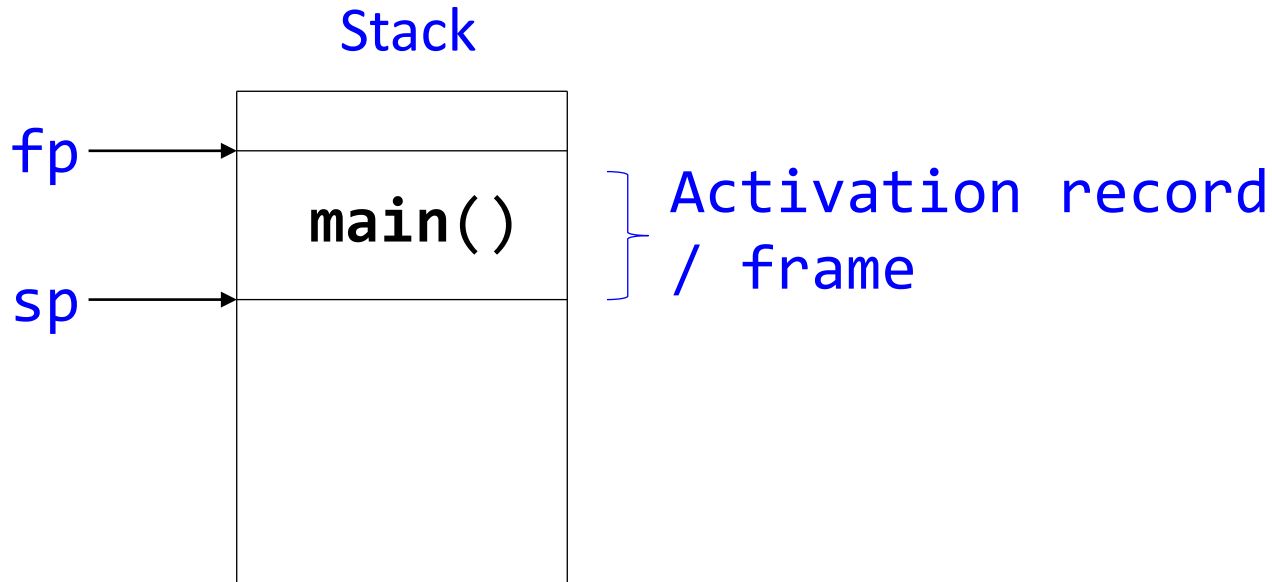
Activation Record

- A *sub-segment* of memory on the stack
 - **Special registers** `$rbp` and `$rsp` track the bottom and top of the stack frame. These are the names in x86 architecture.



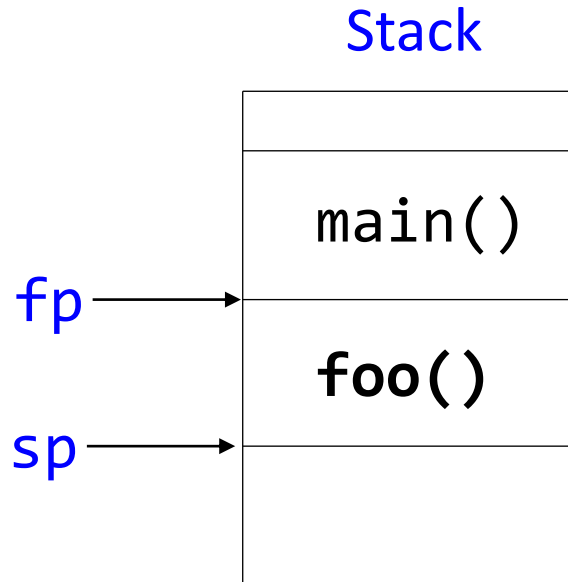
- `$rbp` – base pointer or frame pointer (**fp**)
- `$rsp` – stack pointer (**sp**)

Activation Record - Example



```
→ main() {  
    ...  
    foo();  
    ...  
}  
  
foo() {  
    bar();  
    ...  
    baz();  
}
```

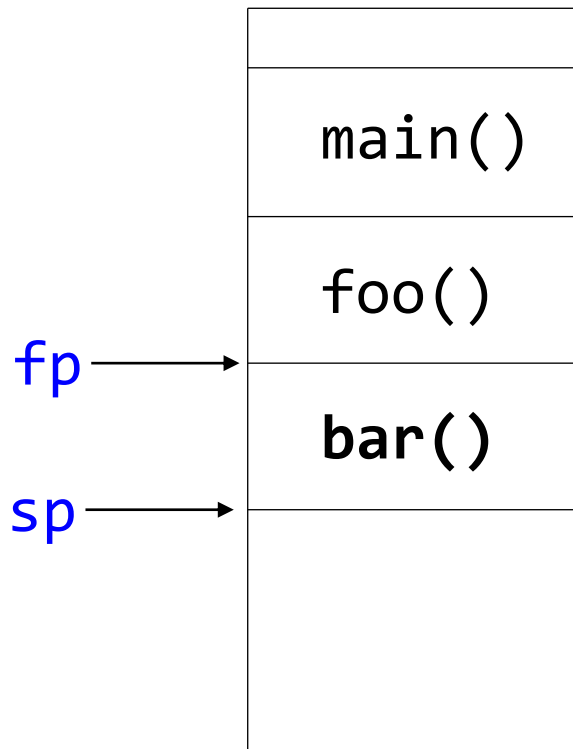
Activation Record - Example



```
main() {  
    ...  
    → foo();  
    ...  
}  
  
foo() {  
    bar();  
    ...  
    baz();  
}
```

Activation Record - Example

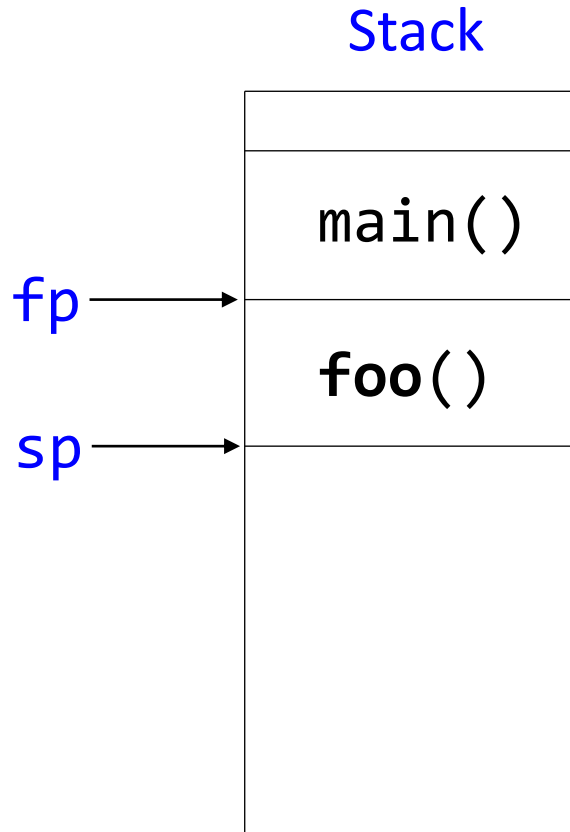
Stack



```
main() {  
    ...  
    foo();  
    ...  
}
```

```
foo() {  
    → bar();  
    ...  
    baz();  
}
```


Activation Record - Example



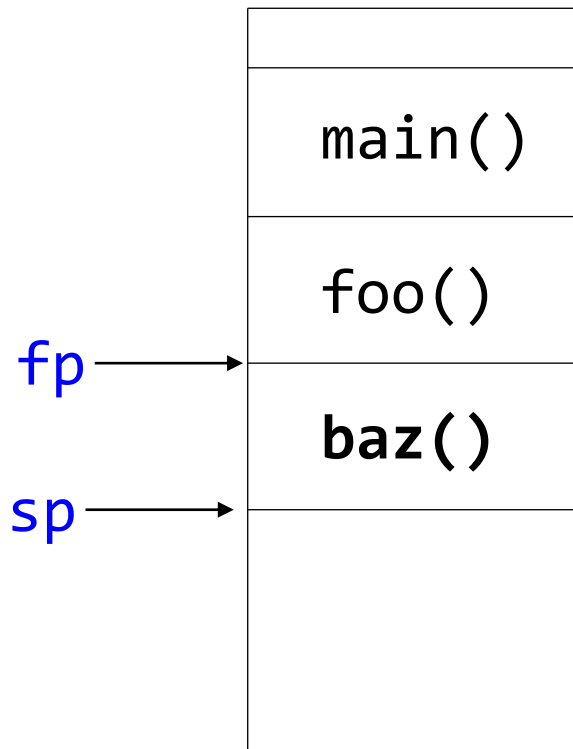
```
main() {  
    ...  
    foo();  
    ...  
}
```

```
foo() {  
    bar();  
    ...  
    baz();  
}
```



Activation Record - Example

Stack



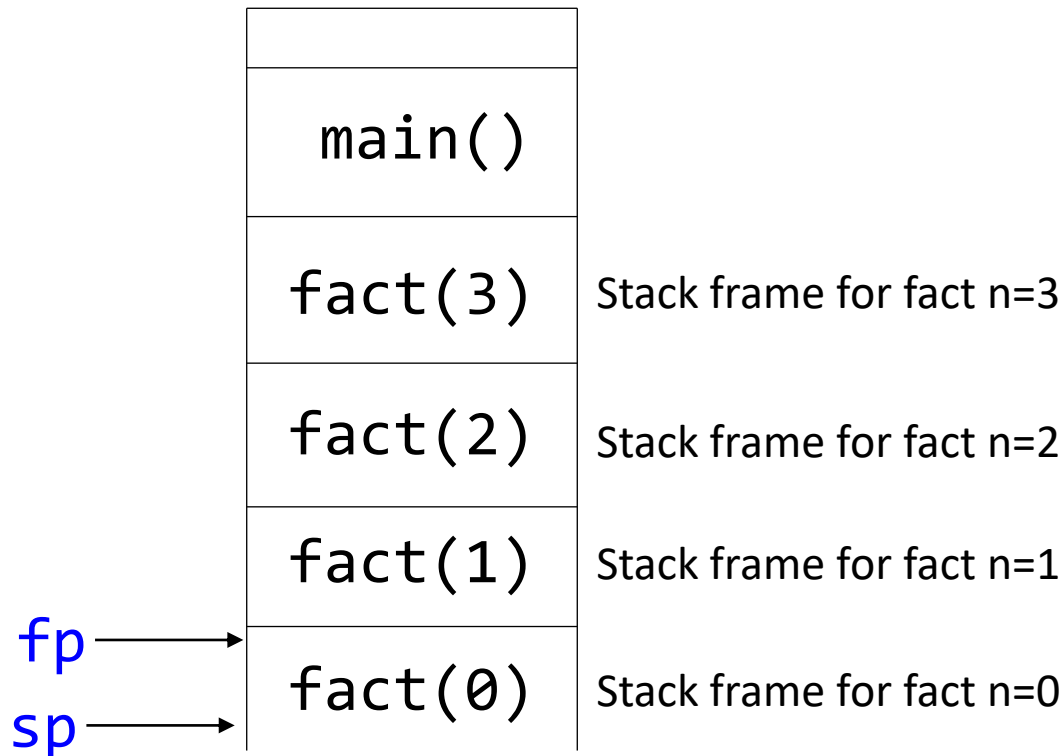
```
main() {  
    ...  
    foo();  
    ...  
}
```

```
foo() {  
    bar();  
    ...
```

```
    → baz();  
}
```

Activation Record – Example (Recursive Functions)


Stack



```
main() {  
    ...  
    fact(3);  
    ...  
}  
  
fact(int n) {  
    if (n=0) return 1  
    return n*fact(n-1)  
}
```

Activation Record – Example (Recursive Functions)

Stack

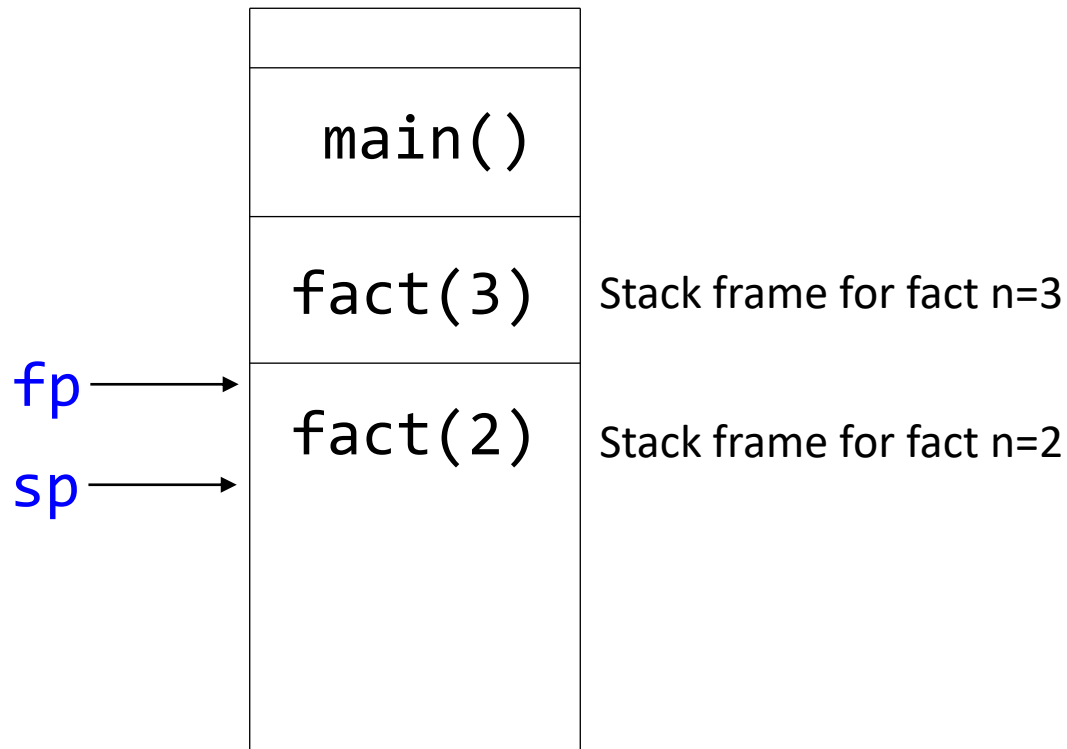


```
main() {  
    ...  
    fact(3);  
    ...  
}  
  
fact(int n) {  
    if (n=0) return 1  
    return n*fact(n-1)  
}
```

Stack frame for n=0 popped off. 1 Returned.

Activation Record – Example (Recursive Functions)

Stack

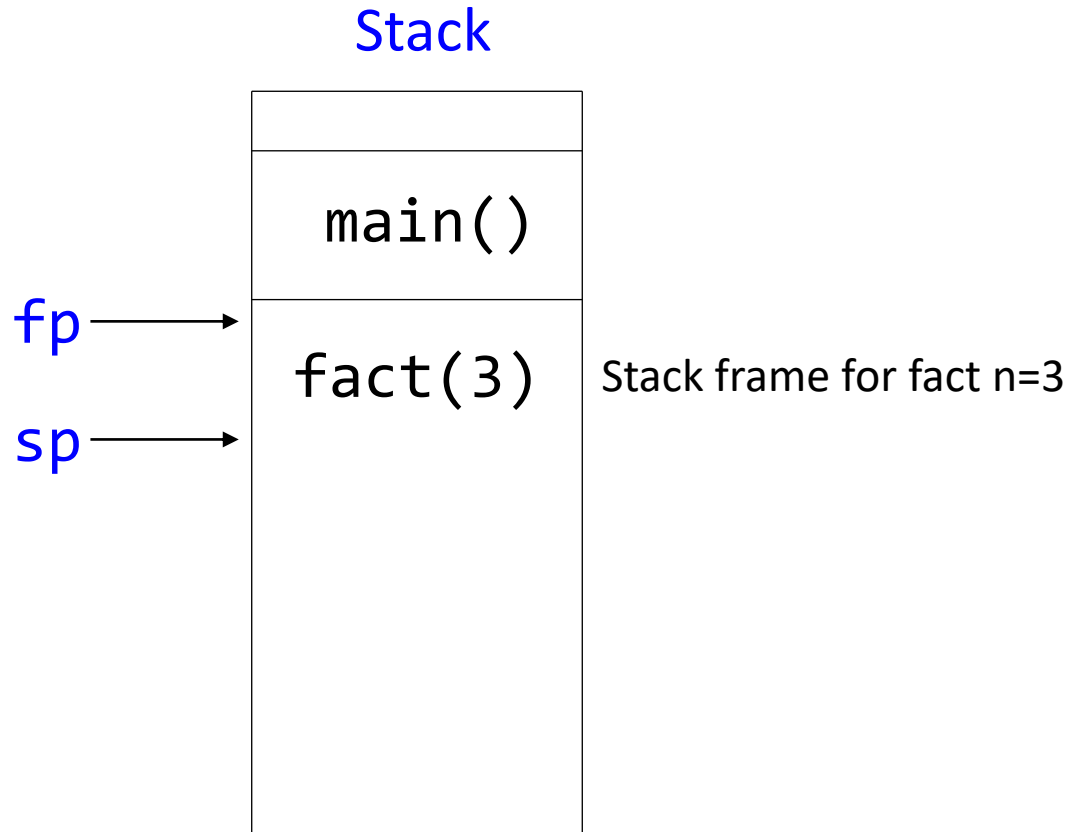


```
main() {  
    ...  
    fact(3);  
    ...  
}
```

```
fact(int n) {  
    if (n=0) return 1  
    return n*fact(n-1)  
}
```

Stack frame for n=1 popped off. 1 Returned.

Activation Record – Example (Recursive Functions)

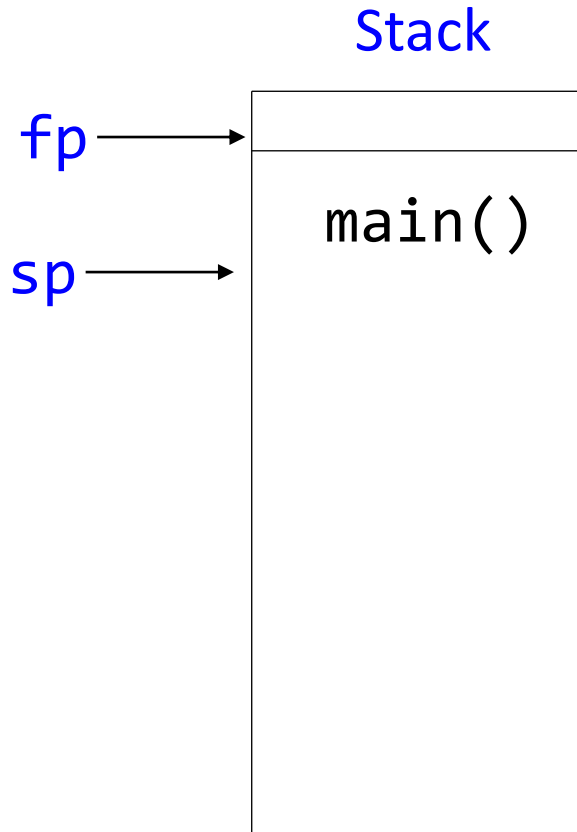


```
main() {  
    ...  
    fact(3);  
    ...  
}
```

```
fact(int n) {  
    if (n=0) return 1  
    return n*fact(n-1)  
}
```

Stack frame for n=2 popped off. 2 Returned.

Activation Record – Example (Recursive Functions)



```
main() {  
    ...  
    fact(3);  
    ...  
}  
  
fact(int n) {  
    if (n=0) return 1  
    return n*fact(n-1)  
}
```

Stack frame for n=3 popped off. 6 Returned.

Activation Record

- What happens when a function is called?
 1. fp and sp get adjusted
 2. Memory for the activation record is allocated on stack
 - The size of the memory allocated depends on local variables used by the called function (consult function's symbol table for this)
 3. Each invocation of a function has its own instantiation of local variables
- When the function call returns:
 - Memory for the activation record is destroyed when the function returns

Activation Record

- What is stored in the activation record?

Depends on the language being implemented:

- Temporaries
 - Local vars
 - Saved registers
 - Return address, previous fp
 - Return value
 - Actual Params
- Who stores this information?
 - Caller } together execute *calling sequence* and *return*
 - Callee } *sequence*

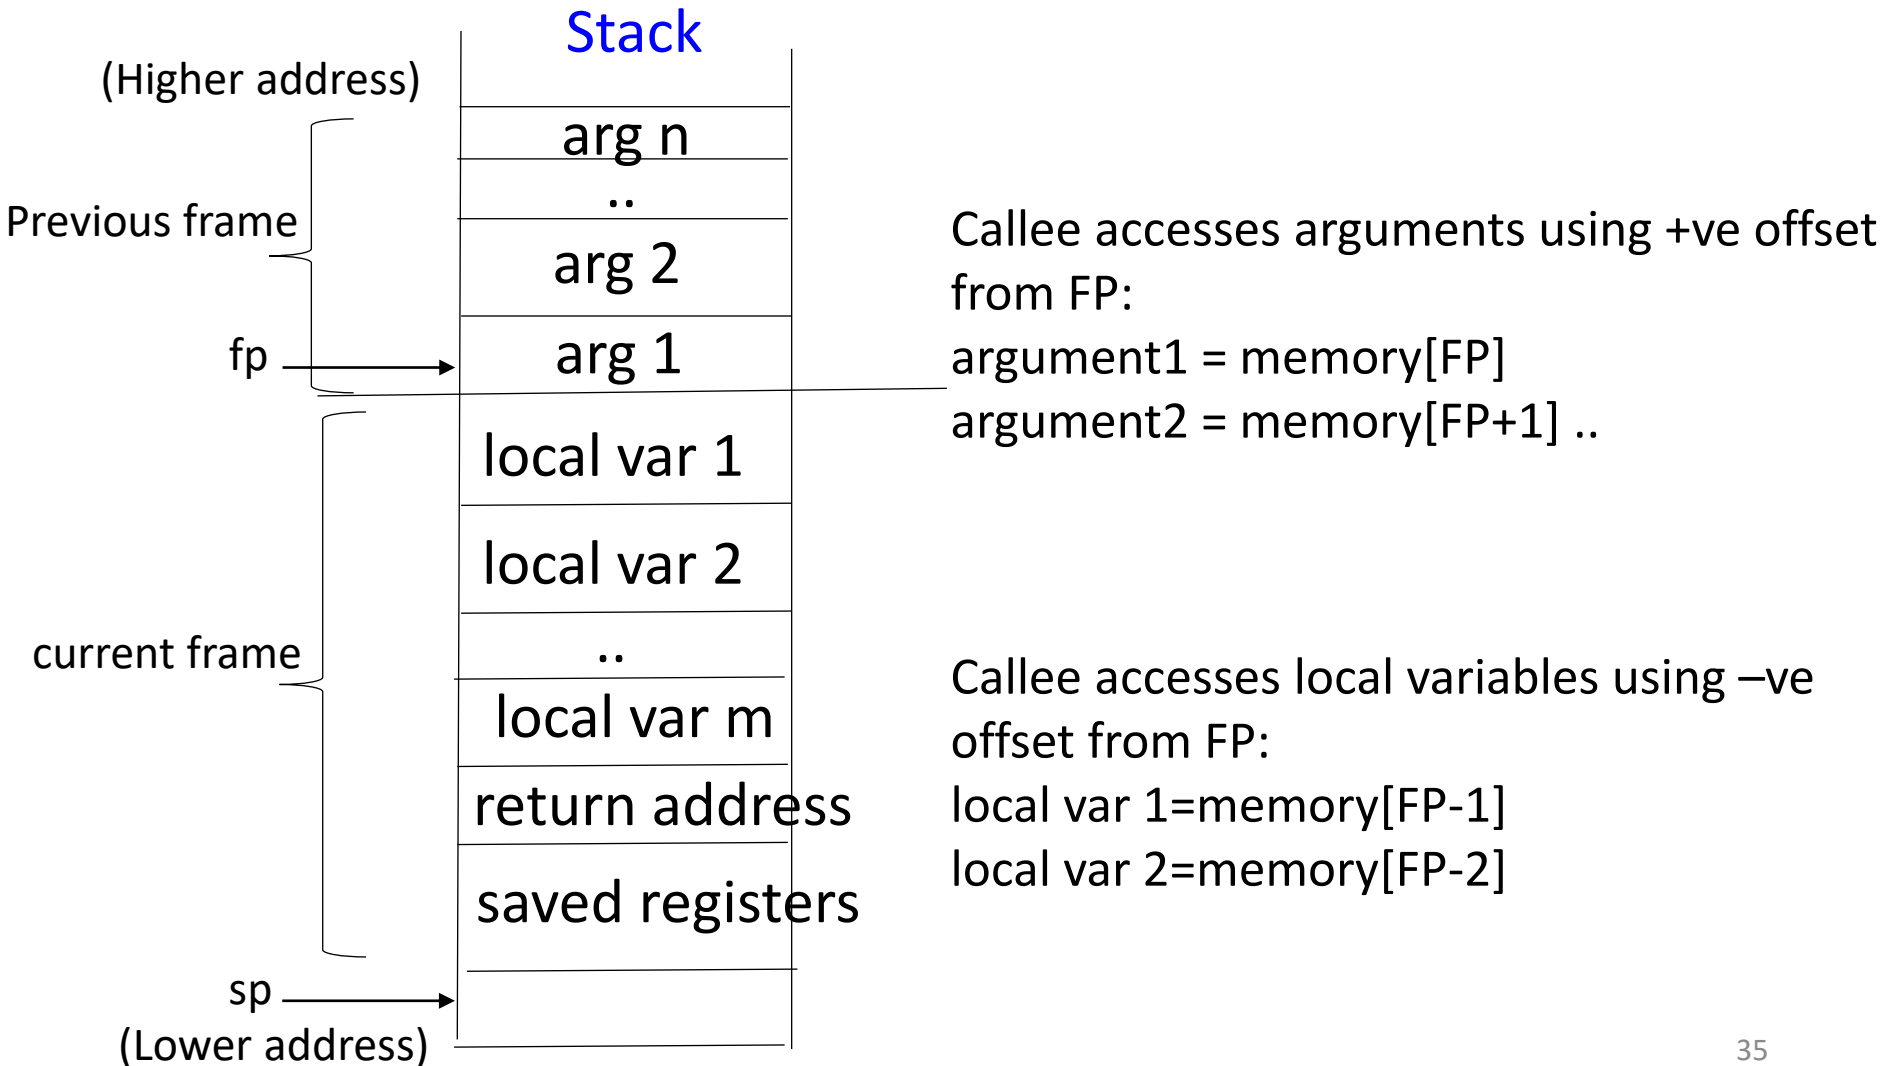
Application Binary Interface (ABI)

- How is data organized on the activation record?
 - ABI is the specification on how data is provided to functions
 - Caller saves or callee saves
 - ABI is meant to deliver interoperability between different compilers
 - Compile the function using one compiler to create an object code, Link object code with other code compiled using a different compiler



form the *calling convention*

Typical Activation Record



Function call: Peeking at Activation Record

- When `main` calls function `foo`

1. The following are pushed on to the stack:

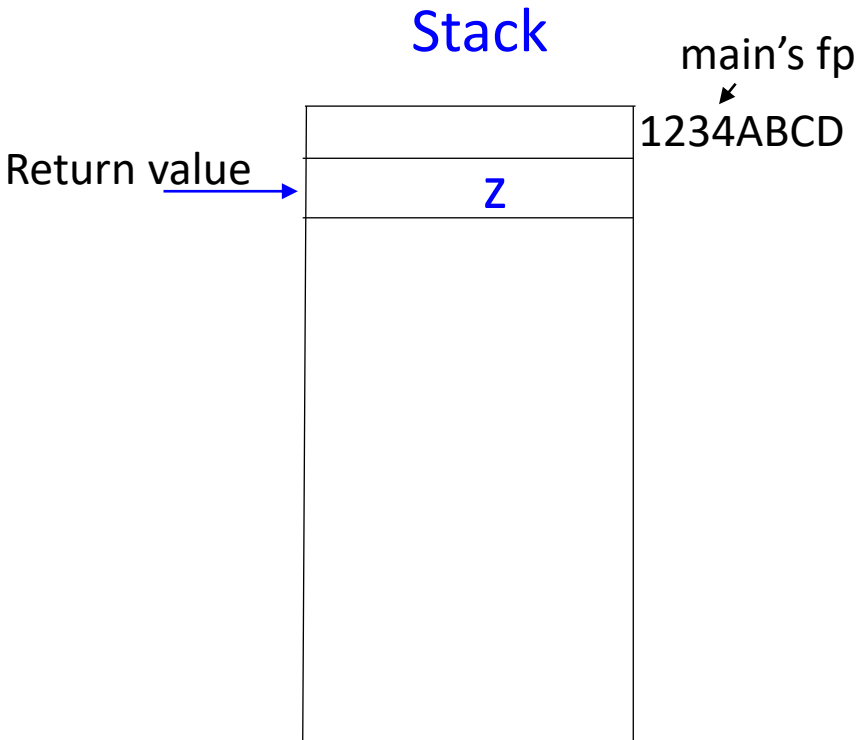
1. `foo`'s arguments
2. Space to hold `foo`'s return value
3. Address of the next instruction executed (in `main`) when `foo` returns (return address)
4. Current value of `$rbp` (frame pointer)

```
main() {  
    ...  
    foo();  
    ...  
}
```

`$rsp` is automatically updated (decremented) to point to current top of the stack.

2. `$rbp` is assigned the value of `$rsp`

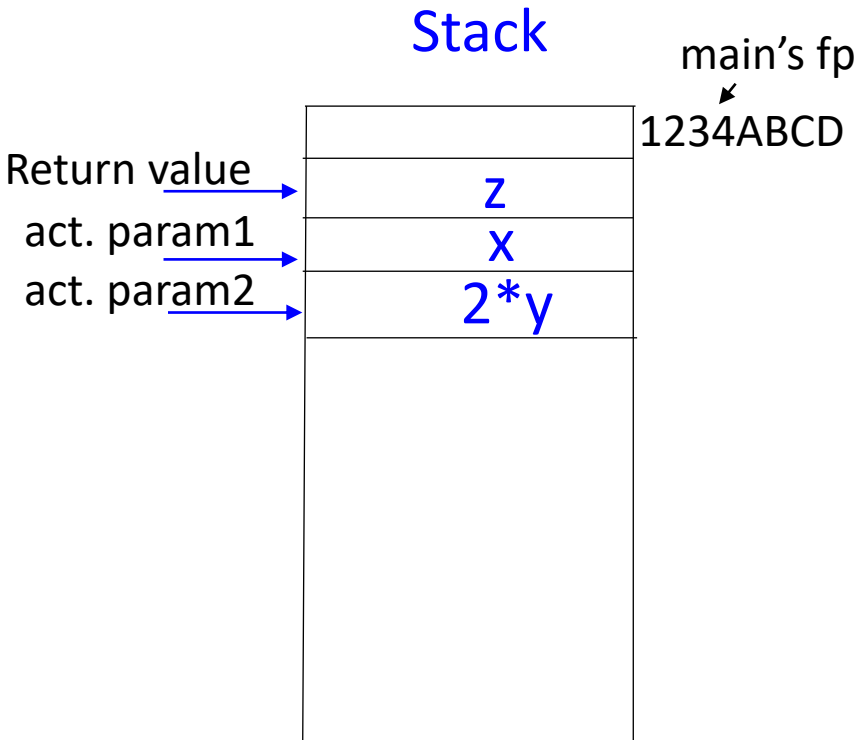
Function call: Peeking at Activation Record



```
main() {  
    z=foo(x, 2*y);  
    return;  
}
```

```
int foo(int a, int b) {  
    int l1, l2  
    l1=a;  
    l2=b;  
    return l1+l2;  
}
```

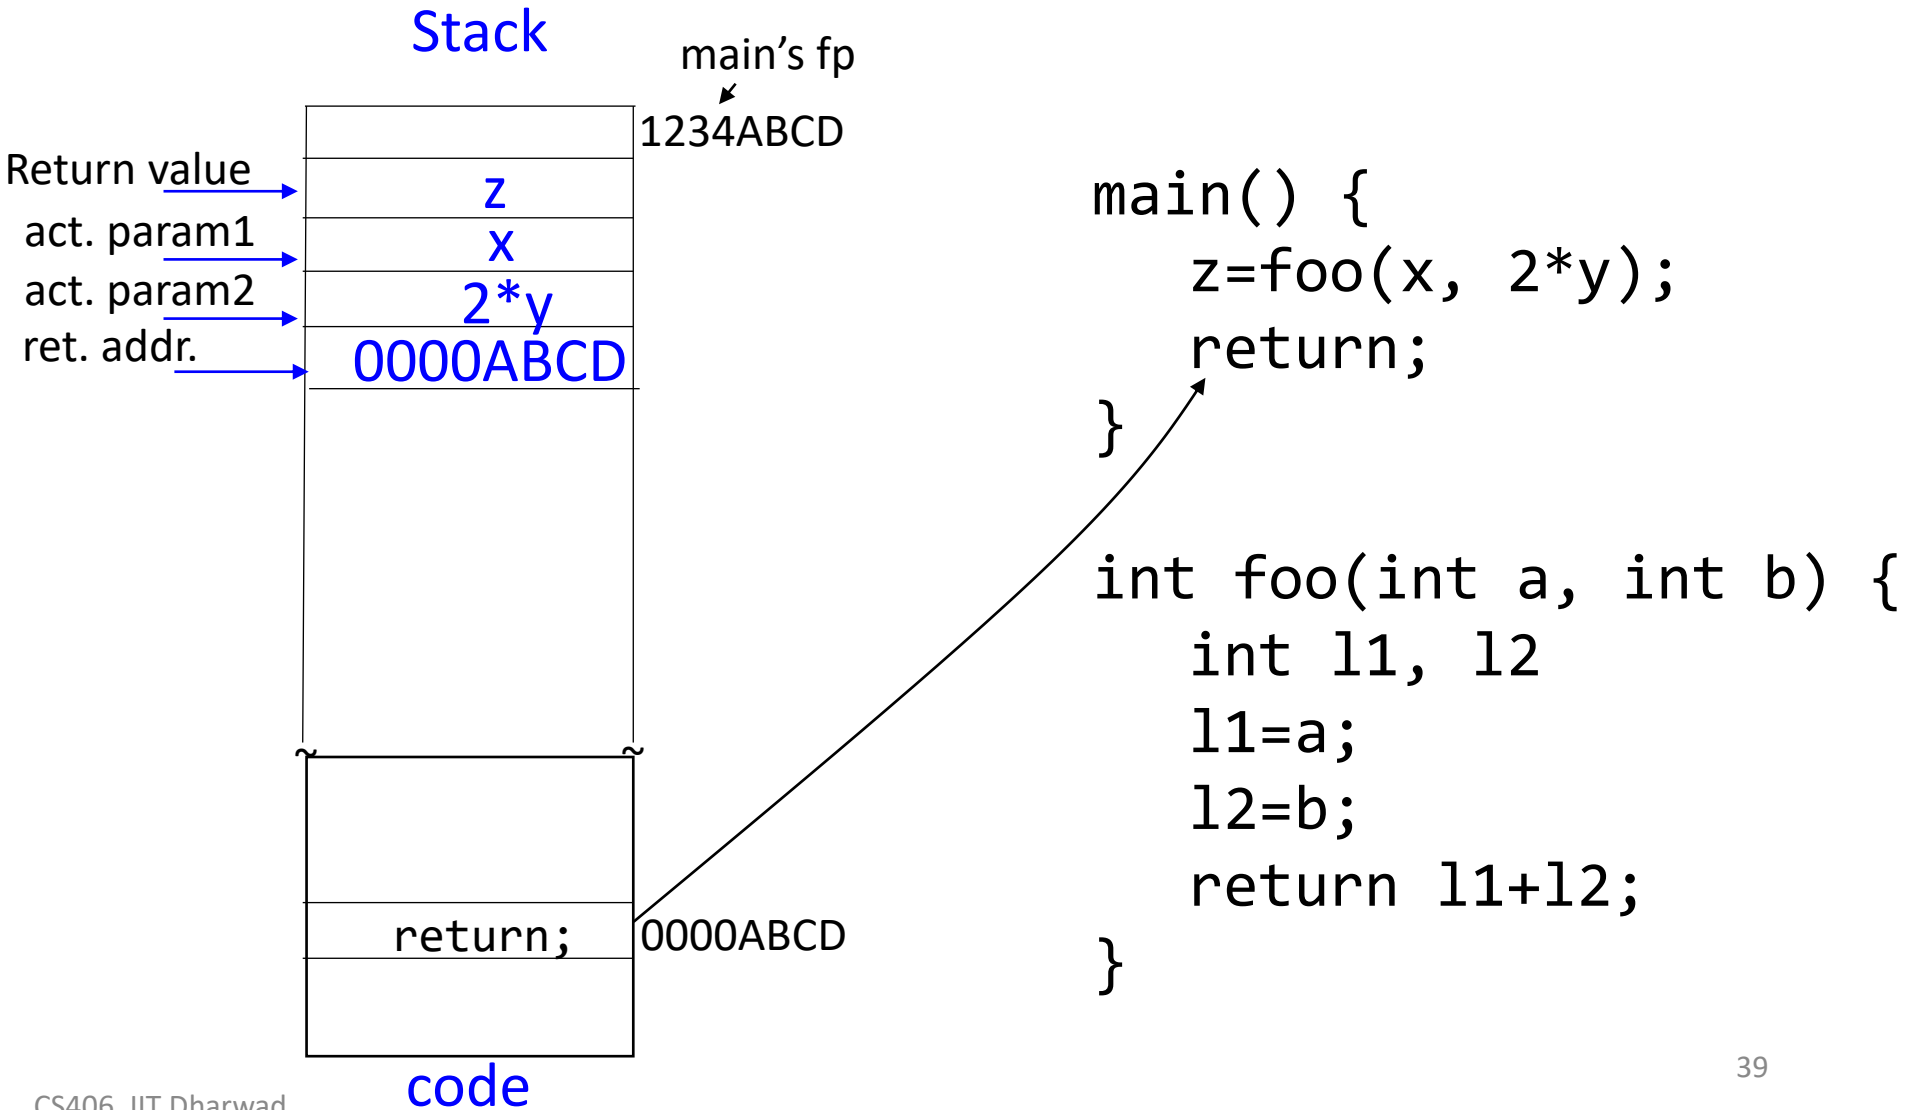
Function call: Peeking at Activation Record



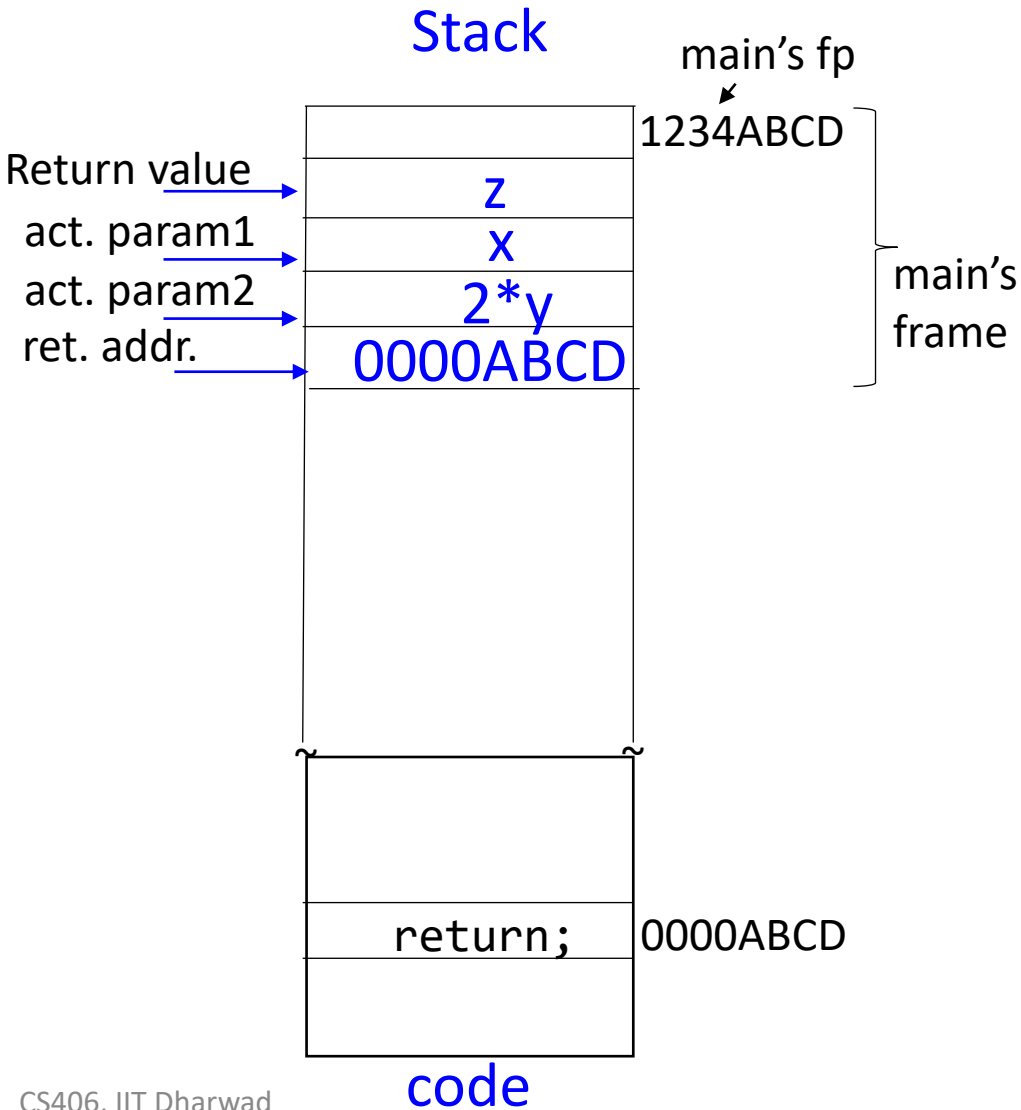
```
main() {  
    z=foo(x, 2*y);  
    return;  
}
```

```
int foo(int a, int b) {  
    int l1, l2  
    l1=a;  
    l2=b;  
    return l1+l2;  
}
```

Function call: Peeking at Activation Record



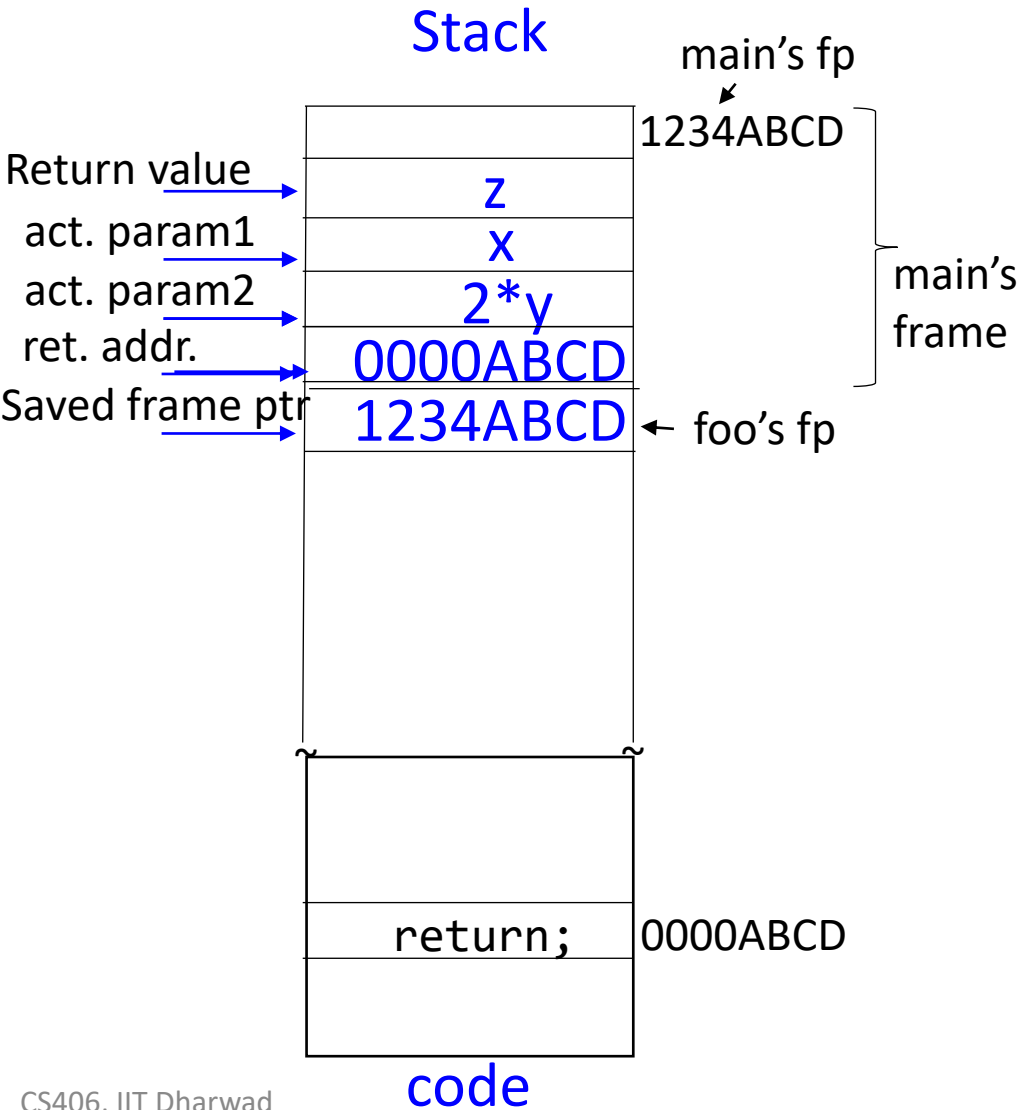
Function call: Peeking at Activation Record



```
main() {  
    z=foo(x, 2*y);  
    return;  
}
```

```
int foo(int a, int b) {  
    int l1, l2  
    l1=a;  
    l2=b;  
    return l1+l2;  
}
```

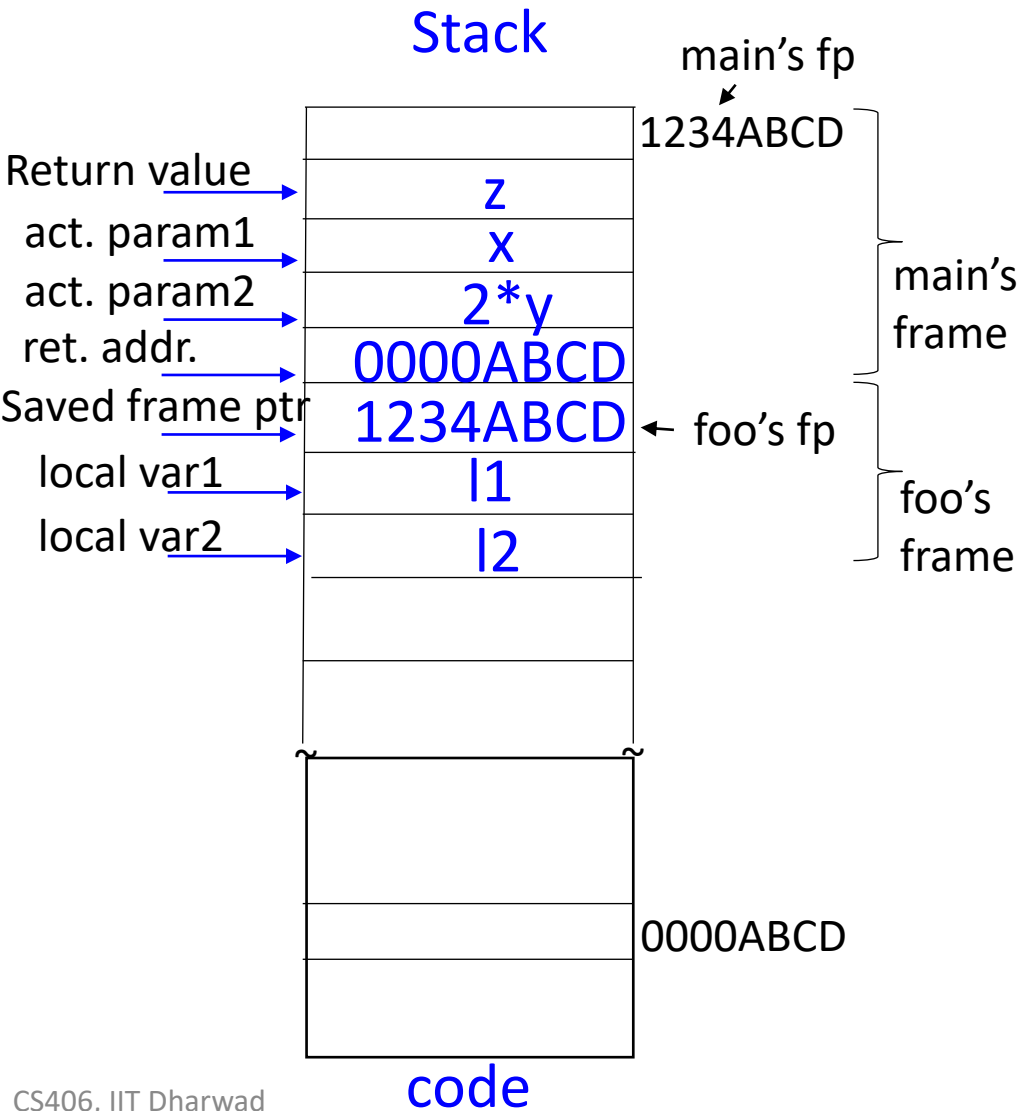

Function call: Peeking at Activation Record



```
main() {
    z=foo(x, 2*y);
    return;
}
```

```
int foo(int a, int b) {
    int l1, l2
    l1=a;
    l2=b;
    return l1+l2;
}
```

Function call: Peeking at Activation Record



```
main() {  
    z=foo(x, 2*y);  
    return;  
}
```

```
int foo(int a, int b) {  
    int l1, l2  
    l1=a;  
    l2=b;  
    return l1+l2;  
}
```

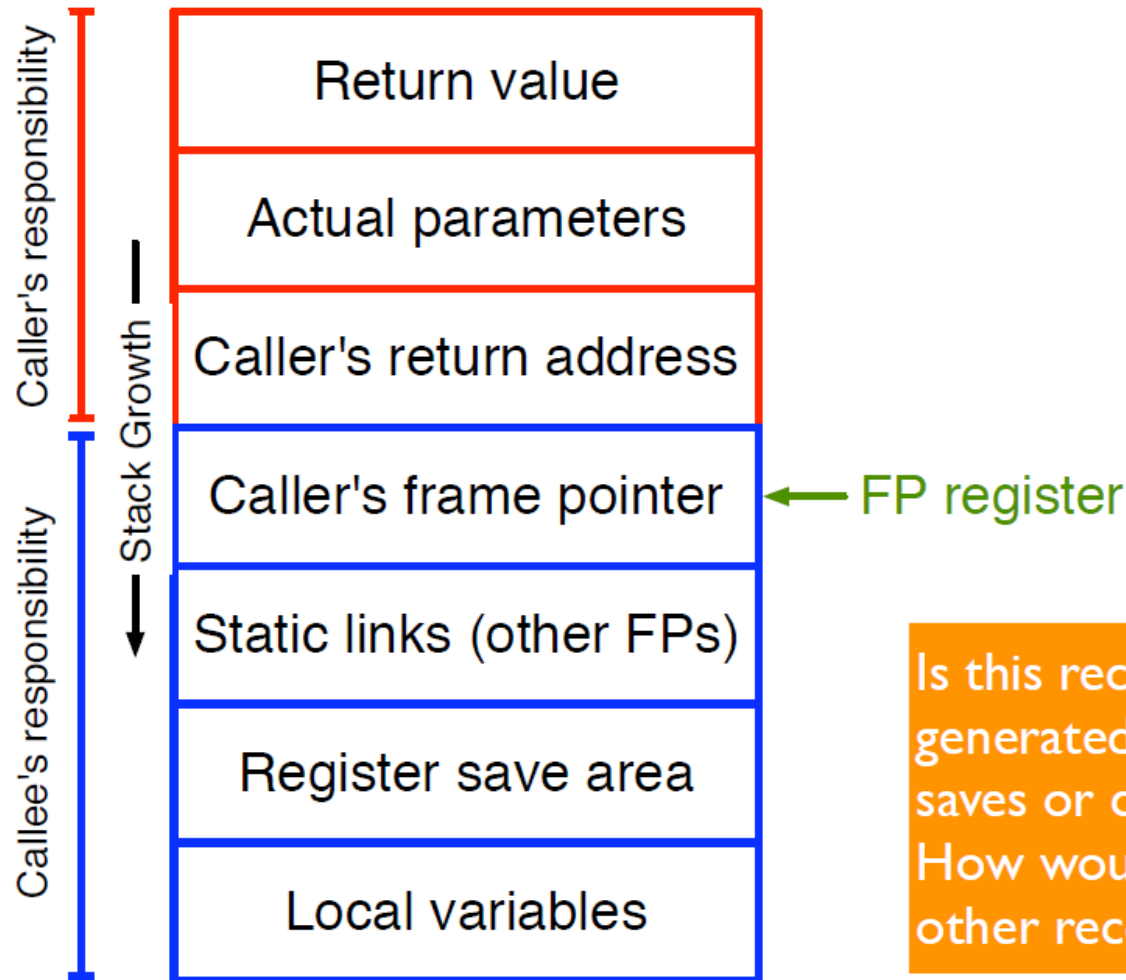
Function calls – Register Handling

- **Did not use registers** in the previous example (for parameter passing)
- Registers are faster than memory. So, compiler should keep parameters in registers whenever possible
- Modern calling convention places first few arguments in registers (arg1 in r1, arg2 in r2, arg3 in r3...) and the remaining in memory.
 - In x86 C-ABI, first 6 arguments are passed in registers
- What if callee wants to use registers r1, r2, r3 etc. for local computation? Callee must save the registers in its stack frame.

Function calls – Register Handling

- Two options: **caller saves** or **callee saves**
- Caller Saves
 - Caller pushes all the registers it is using on to the stack before calling the function
 - Restores the registers after the function returns
- Callee Saves
 - Callee pushes all the registers it is *going to use* on the stack immediately after being called
 - Restores the registers just before it returns

Activation records



Is this record generated for callee-saves or caller-saves? How would the other record look?

Activation Record – Return Address and Return Value

- Callee must be able to return to the caller when done
- Return address is the address of the instruction following the function call
- Return address can be placed on the stack or on register
- The `call` instruction on modern machines places the return address in a specific register
- Return value is placed in a specific register by the callee function

The frame pointer

- Manipulate with instructions like `link` and `unlink`
 - `Link`: push current value of FP on to stack, set FP to top of stack
 - `Unlink`: read value at current address pointed to by FP, set FP to point to that value
 - In other words: `link` pushes a new frame onto the stack, `unlink` pops it off

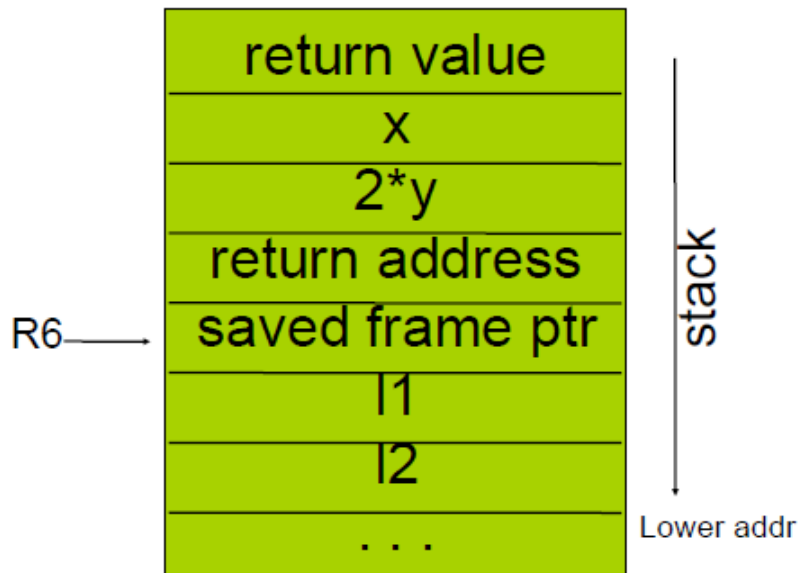
Stack Pointer

- SP is manipulated through push and pop instructions

```
Push x:  
stack_pointer--  
Memory[stack_pointer] = x
```

```
Pop x:  
x = Memory[stack_pointer]  
stack_pointer--
```


Example Subroutine Call and Stack Frame



3-address code:

```
push
push x
mul 2 y t1
push t1
jsr SubOne
pop
pop
pop z
```

assembly code:

```
push
push x
load y R1
muli 2 R1
push R1
jsr SubOne
pop
pop
pop R1
store R1 z
```

`z = SubOne(x, 2*y);`

```
int SubOne(int a, int b) {
    int l1, l2;
    l1 = a;
    l2 = b;
    return l1+l2;
};
```

```
link 3
move $P1 $L1
move $P2 $L2
add $L1 $L2 t2
move t2 $R
unlink
ret
```

```
link R6 3
load 3(R6) R1
store R1 -1(R6)
load 2(R6) R2
store R2 -2(R6)
load -1(R6) R1
add -2(R6) R1
store R1 4(R6)
unlink
```

Question ?

Where are the command-line arguments stored?

How about environment variables such as LD_LIBRARY_PATH and PATH?

Challenge Q: *are there scenarios where the activation record is required to be allocated on the heap?*


```
fun f(x) =  
  let  
    fun g(y) = x + y  
  in  
    g  
  end
```


```
val z = f(4)  
val w = z(5)
```

Local Optimizations

Naïve approach

- “Macro-expansion”
- Treat each 3AC instruction separately, generate code in isolation

ADD A, B, C  LD A, R1
LD B, R2
ADD R1, R2, R3
ST R3, C

MUL A, 4, B  LD A, R1
MOV 4, R2
MUL R1, R2, R3
ST R3, B

Why is this bad? (I)


MUL A, 4, B →
LDA A, R1
MOV 4, R2
MUL R1, R2, R3
ST R3, B

MUL A, 4, B →
LDA A, R1
MULI R1, 4, R3
ST R3, B

There is a better instruction available!

Too many instructions
Should use a different instruction type

Why is this bad? (II)



Why is this bad? (III)

ADD A, B, C → LD A, R1
LD B, R2
ADD R1, R2, R3
ST R3, C

ADD A, B, C → LD A, R1
LD B, R2
ADD R1, R2, R3
ST R3, C
ADD A, B, D LD A, R4
LD B, R5
ADD R4, R5, R6
ST R6, D

Wasting instructions recomputing $A + B$

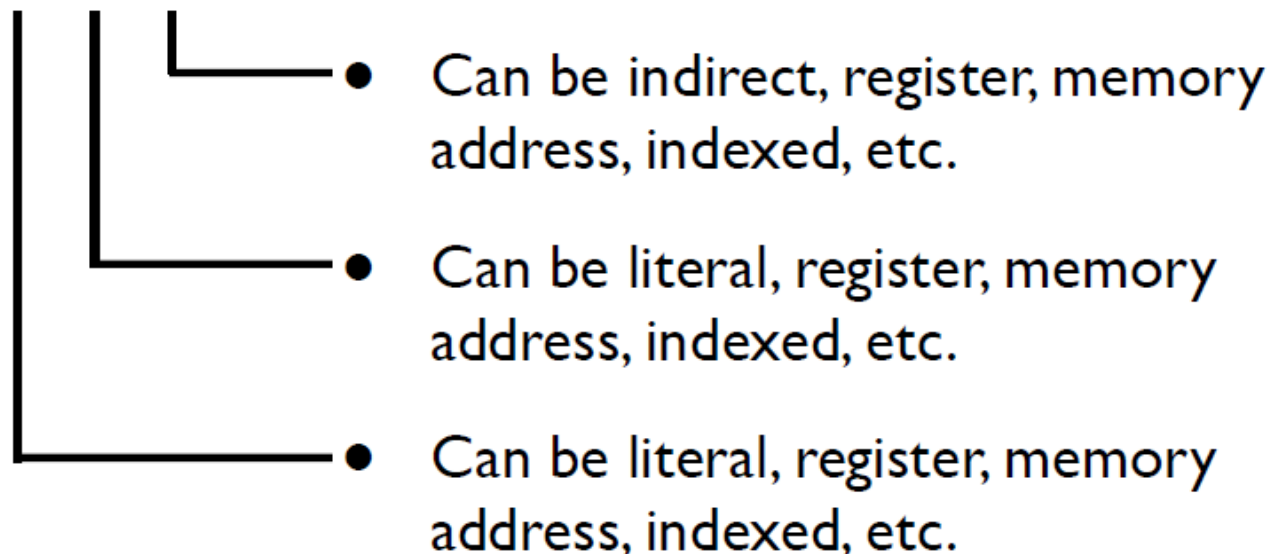
How do we address this?

- Several techniques to improve performance of generated code
 - *Instruction selection* to choose better instructions
 - *Peephole optimizations* to remove redundant instructions
 - *Common subexpression elimination* to remove redundant computation
 - *Register allocation* to reduce number of registers used

Instruction selection

- Even a simple instruction may have a large set of possible address modes and combinations

+ A B C



- Dozens of potential combinations!

More choices for instructions

- Auto increment/decrement (especially common in embedded processors as in DSPs)
 - e.g., load from this address and increment it
 - Why is this useful?
- Three-address instructions
- Specialized registers (condition registers, floating point registers, etc.)
- “Free” addition in indexed mode
MOV (R1)offset R2
 - Why is this useful?

Peephole optimizations

- Simple optimizations that can be performed by pattern matching
- Intuitively, look through a “peephole” at a small segment of code and replace it with something better
- Example: if code generator sees `ST R X; LD X R`, eliminate load
- Can recognize sequences of instructions that can be performed by single instructions

`LDI R1 R2; ADD R1 4 R1` replaced by

`LDINC R1 R2 4` //load from address in R1 then inc by 4

Peephole optimizations

- Simple optimizations that can be performed by pattern matching
- Intuitively, look through a “peephole” at a small segment of code and replace it with something better
- Example: if code generator sees `ST R X; LD X R`, eliminate load

Get the data present at address in R2 and put it in R1 be

LDI R1 R2; ADD R1 4 R1 replaced by

LDINC R1 R2 4 //load from address in R1 then inc by 4

Peephole optimizations

- Constant folding

ADD lit1, lit2, Rx \longrightarrow MOV lit1 + lit2, Rx
MOV lit1, Rx
ADD li2, Rx, Ry \longrightarrow MOV lit1 + lit2, Ry

- Strength reduction

MUL operand, 2, Rx \longrightarrow SHIFTL operand, 1, Rx
DIV operand, 4, Rx \longrightarrow SHIFTR operand, 2, Rx

- Null sequences

MUL operand, 1, Rx \longrightarrow MOV operand, Rx
ADD operand, 0, Rx \longrightarrow MOV operand, Rx

Peephole optimizations

- Combine operations

JEQ L1
JMP L2 \longrightarrow JNE L2
L1: ...

- Simplifying

SUB operand, 0, Rx \longrightarrow NEG Rx

- Special cases (taking advantage of ++/--)

ADD 1, Rx, Rx \longrightarrow INC Rx
SUB Rx, 1, Rx \longrightarrow DEC Rx

- Address mode operations

MOV A R1
ADD 0(R1) R2 R3 \longrightarrow ADD @A R2 R3

Superoptimization

- Peephole optimization/instruction selection writ large
- Given a sequence of instructions, find a different sequence of instructions that performs the same computation in less time
- Huge body of research, pulling in ideas from all across computer science
 - Theorem proving
 - Machine learning

Common subexpression elimination

- Goal: remove redundant computation, don't calculate the same expression multiple times

1: $A = B * C$

2: $E = B * C$

Keep the result of statement 1 in a temporary and reuse for statement 2

- Difficulty: how do we know when the same expression will produce the same result?

1: $A = B * C$

2: $B = \text{<new value>}$

3: $E = B * C$

B is “killed.” Any expression using B is no longer “available,” so we cannot reuse the result of statement 1 for statement 3

- This becomes harder with pointers (how do we know when B is killed?)