

CS601: Software Development for Scientific Computing

Autumn 2022

Week1: Overview

Let us listen to Jack Dongarra

- <https://youtu.be/Oe9LRKoE6L0>

Course Takeaways (intended)

- Non-CS majors:
 1. Write code and
 2. Develop software (not just write standalone code)
 - Numerical software
- CS-Majors:

In addition to the above two:

 3. Learn to face mathematical equations and implement them with confidence

What is this course about?

Software Development

+

Scientific Computing

This course **NOT** about..

- Software Engineering
 - Systematic study of Techniques, Methodology, and Tools to build correct software within time and price budget (topics covered in CS305)
 - People, Software life cycle and management etc.
- Scientific Computing
 - Rigorous exploration of numerical methods, their analysis, and theories
 - Programming models (topics covered in CS410)

Who this course is for?

- You are interested in scientific computing
- You are interested in high-performance computing
- You want to build / add to a large software system

Software Development

- *Software development is the process of **conceiving, specifying, designing, programming, documenting, testing,** and **bug fixing** involved in **creating and maintaining applications, frameworks, or other software components.***

*Software development is a process of writing and maintaining the source code, but in a broader sense, it includes all that is involved between the **conception** of the desired software through to the **final manifestation** of the software, ...*

- Wikipedia on "Software Development"

Scientific Computing

- Also called computational science
 - *Development of models to understand systems (biological, physical, chemical, engineering, humanities)*

*Collection of tools, techniques, and theories required **to solve on a computer** mathematical models of problems in science and engineering*

Why C++ ?

- C/C++/Fortran codes form the majority in scientific computing codes
- Catch a lot of errors early (e.g. at *compile-time* rather than at *run-time*)
- Has features for *object-oriented* software development
- Known to result in codes with better *performance*

Let us dive into an example....

Example - Factorial

- $$n! = n \times (n-1) \times (n-2) \times \dots \times 3 \times 2 \times 1$$
$$(n-1)! = (n-1) \times (n-2) \times \dots \times 3 \times 2 \times 1$$

therefore,

Definition1: $n! = n \times (n-1)!$

is this definition complete?

- plug 0 to n and the equation breaks.

Definition2:

$$n! = \begin{cases} n \times (n-1)! & \text{when } n \geq 1 \\ 1 & \text{when } n = 0 \end{cases}$$

Exercise 1

- Does this code implement the definition of factorial correctly?

```
int fact(int n){  
    if(n==0)  
        return 1;  
  
    return n*fact(n-1);  
  
}
```

Example - Factorial

Definition2:

$$n! = \begin{cases} n \times (n-1)! & \text{when } n \geq 1 \\ 1 & \text{when } n = 0 \end{cases}$$

is this definition complete?

- $n!$ is not defined for negative n

Solution - Factorial

```
int fact(int n){  
    if(n<0)  
        return ERROR;  
    if(n==0)  
        return 1;  
  
    return n*fact(n-1);  
  
}
```

Exercise 2

- In how many flops does the code execute?
assume 1 flop = 1 step executing **any** arithmetic operation

```
int fact(int n){  
    if(n<0)  
        return ERROR;  
    if(n==0)  
        return 1;  
  
    return n*fact(n-1);  
  
}
```

Exercise 3

- Does the code yield correct results for any n ?

```
int fact(int n){  
    if(n<0)  
        return ERROR;  
    if(n==0)  
        return 1;  
  
    return n*fact(n-1);  
}
```


Who this course is for?

- Anybody who wishes to develop "computational thinking"
 - A skill necessary for everyone, not just computer programmers
 - An approach to problem solving, designing systems, and understanding human behavior that draws on concepts fundamental to computer science.

Computational Thinking - Examples

- How difficult is the problem to solve? And what is the best way to solve?
- Modularizing something in anticipation of multiple users
- Prefetching and caching in anticipation of future use
- Thinking recursively
- Reformulating a seemingly difficult problem into one which we know how to solve by reduction, embedding, transformation, simulation
 - Are approximate solutions accepted?
 - False positives and False negatives allowed? etc.
- Using abstraction and decomposition in tackling large problem
- ...

Computational Thinking – 2 As

- Abstractions
 - Our “mental” tools
 - Includes: choosing right abstractions, operating at multiple layers of abstractions, and defining relationships among layers
- Automation
 - Our “metal” tools that amplify the power of “mental” tools
 - Is mechanizing our abstractions, layers, and relationships
 - Need precise and exact notations / models for the “computer” below (“computer” can be human or machine)

Computing - 2 As Combined

- Computing is the **automation** of our **abstractions**
- Provides us the ability to scale
 - Make infeasible problems feasible
 - E.g. SHA-1 not safe anymore
 - Improve the answer's precision
 - E.g. capture the image of a black-hole

Summary: choose the right abstraction and computer

Recap

- Need to be precise
 - recall: $n! = 1$ for $n=0$, not defined for negative n
- Choosing right abstractions
 - recall: use of recursion, correct data type
- Ability to define the complexity
 - recall: flop calculation
- Next?

Recap

- Need to be precise
 - recall: $n! = 1$ for $n=0$, not defined for negative n
- Choosing right abstractions
 - recall: use of recursion, correct data type
- Ability to define the complexity
 - recall: flop calculation
- Choose the right “computer” for mechanizing the abstractions chosen

Scientific Software - Characteristics

- The answer is not a typical yes/no, red/blue/green
- The answer varies continuously. Think of computing the value of $\pi = 3.141592\dots$
- Uses approximations. Think of *discretization*
- Employs efficient *kernels*
 - Kernels are core operations that are executed very frequently
- Should be able to adapt to change.
 - Writing everything from scratch is not an option
- Deals with large-scale problems
 - Lot of input/output data or both
 - Computationally hard

General Approach to Solving a Computational Problem

1. **Problem statement:** more precise this is, the easier it is to design and implement
2. **Solution Algorithm:** exactly how is the problem going to be solved
3. **Implementation:** breaking the algorithm into manageable pieces and putting it all together to solve the problem using a language of choice.
4. **Verification:** checking that the implementation solves the original problem.
 1. Often most difficult step, because you don't know the correct answer.

Toward Scientific Software

- Necessary Skills:
 - Understanding the mathematical problem
 - Understanding numerics
 - Designing algorithms and data structures
 - Selecting language and using libraries and tools
 - Verify the correctness of the results
 - Quick learning of new programming languages
 - E.g. [Regent](#)

Exercise


Compute root(s) of:

$$x = \cos x; \quad x \in \mathbb{R}$$

roots, also called zeros, is the value of the argument/input to the function when the function output vanishes i.e. becomes zero

Mathematical Problem

- let $y = f(x)$
 $f(x) = \cos(x) - x$
- At $x = x_n$, the value of y is $f(x_n)$. The coordinates of the point are $(x_n, f(x_n)) = \text{known point}$.
- From calculus: **derivative** of a function of single variable at a chosen input value, when it exists, is the **slope of the tangent** to the graph at that input value.
 - $f'(x_n)$ is the slope of the line that is tangent to $f(x)$ at x_n



credit: wikipedia

Mathematical Problem

- From high-school math: point-slope formula for equation of a line

$$y - y_1 = m(x - x_1),$$

given the slope m and any known point (x_1, y_1)

- Substituting with:
 - $(x_n, f(x_n))$ = known point
 - $f'(x_n)$ = slope

Equation of the tangent line to graph of $f(x)$ at x_n :

$$y - f(x_n) = f'(x_n)(x - x_n)$$

Mathematical Problem

- Interested in finding roots i.e. value of x at $y=0$ i.e. at point $(x_{np1}, 0)$.
- Substituting in the equation of the tangent line,

$$y - f(x_n) = f'(x_n)(x - x_n)$$

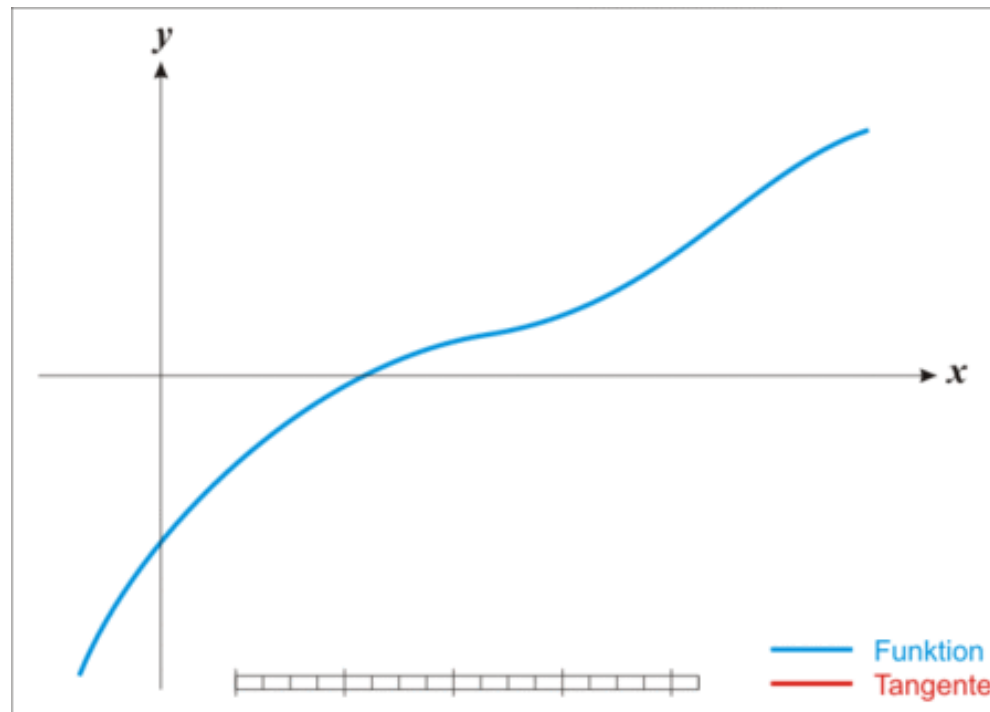
$$= -f(x_n) = f'(x_n)(x_{np1} - x_n)$$

$$= x_{np1} = x_n - f(x_n) / f'(x_n)$$

Mathematical Problem

- Visualizing

(source: https://en.wikipedia.org/wiki/Newton's_method) :



The function f is shown in blue and the tangent line is in red. We see that x_{n+1} is a better approximation than x_n for the root x of the function f .

Mathematical Problem

$$x_2 = x_1 - f(x_1) / f'(x_1)$$

$$x_3 = x_2 - f(x_2) / f'(x_2)$$

$$x_4 = x_3 - f(x_3) / f'(x_3)$$

...

Numerical Analysis

Talk to domain experts

- Choosing the initial value of x
- Does the method converge ?
- What is an acceptable approximation?
- etc.

Designing Algorithms and Data Structures

- Start with x_1

$$x_2 = x_1 - f(x_1) / f'(x_1)$$

$$x_3 = x_2 - f(x_2) / f'(x_2)$$

$$x_4 = x_3 - f(x_3) / f'(x_3)$$

...

- Repeat for up to maxIterations
- Check for $x_{n+1} - x_n$ to be “sufficiently small”
- Choose appropriate data types for x

Selecting libraries and tools

- E.g. use the math library in C++ (cmath)

Verify the correctness of results

- Compare with 'gold' code / benchmark
- Compare with empirical data

Real Numbers \mathbb{R}

- Most scientific software deal with Real numbers.
Our toy code dealt with Reals
 - Numerical software is scientific software dealing with Real numbers
- Real numbers include rational numbers (integers and fractions), irrational numbers (pi etc.)
- Used to represent values of continuous quantity such as time, mass, velocity, height, density etc.
 - Infinitely many values possible
 - But computers have limited memory. So, have to use approximations.

Representing Real Numbers

- Real numbers are stored as *floating point numbers* (floating point system is a scheme to represent real numbers)
- E.g. floating point numbers:
 - $\pi = 3.14159$,
 - 6.03×10^{23}
 - $1.60217733 \times 10^{-19}$

General format: $\pm x \times b^e$

mantissa (number ranges from: 1 to b OR 1/b to 1)

base (e.g. base 10, 8, 2, 16)

exponent

3-digit Calculator

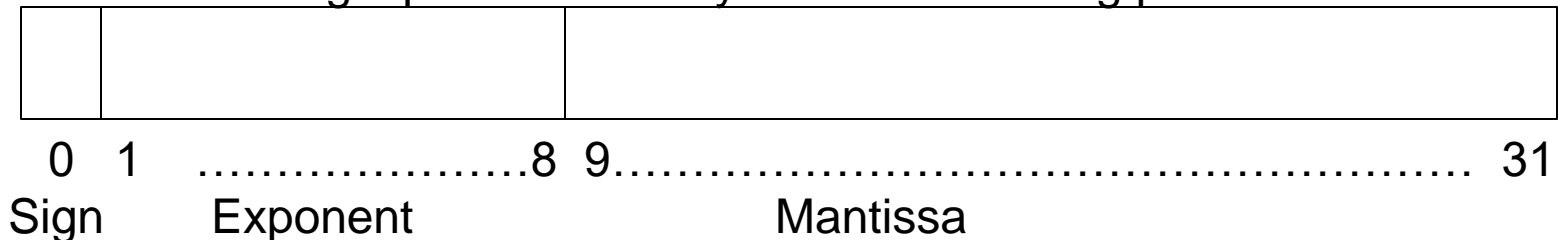
- Suppose base, $b=10$ and
- $x = \pm d_0.d_1d_2 \times 10^e$ where $\begin{cases} 1 \leq d_0 \leq 9, \\ 0 \leq d_1 \leq 9, \\ 0 \leq d_2 \leq 9 \\ -9 \leq e \leq 9 \end{cases}$
- precision = length of mantissa
 - What is the precision here?
- Exercise: What is the smallest positive number?
- Exercise: What is the largest positive number?
- Exercise: When is this representation not enough?
- Exercise: How many numbers can be represented in this format?

IEEE 754 Floating Point System

- Prescribes single, double, and extended precision formats


Precision	u	Total bits used (sign, exponent, mantissa)
Single	6×10^{-8}	32 (1, 8, 23)
Double	2×10^{-16}	64 (1, 11, 52)
Extended	5×10^{-20}	80 (1, 15, 64)

single precision binary IEEE 754 floating point format



IEEE 754 Floating Point Arithmetic

double precision binary IEEE 754 floating point format



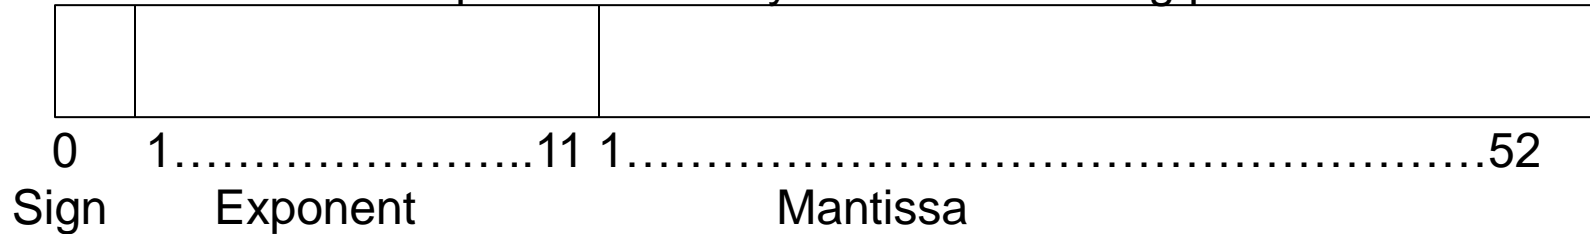
- if exponent bits e_1 - e_{11} are not all 1s or 0s, then the *normalized* number

$$n = \pm(1.m_1m_2..m_{52})_2 \times 2^{(e_1e_2..e_{11})_2 - 1023}$$

- Machine epsilon** is the gap between 1 and the next largest floating point number. $2^{-52} \approx 10^{-16}$ for double.
- Exercise: What is minimum positive *normalized* double number?
- Exercise: What is maximum positive *normalized* double number?

IEEE 754 Floating Point Arithmetic

double precision binary IEEE 754 floating point format



- if exponent bits e_1 - e_{11} are all 0s, then:
the *subnormal* number

$$n = \pm(\mathbf{0}.m_1m_2..m_{52})_2 \times 2^{(e_1e_2..e_{11})_2 - 102\mathbf{2}}$$

- if exponent bits e_1 - e_{11} are all 1s, then:
we can get $-\text{inf}$, NaN, or $+\text{inf}$ based on value of $m_1m_2..m_{52}$
 - If any m is non-zero, the number is NaN (not a number)

IEEE 754 Floating Point Arithmetic

- Order is important
 - Floating point arithmetic is *not associative*
 - $(x+y)+z$ not the same as $x+(y+z)$
- Explicit coding of textbook formula may not be the best option to solve
 - $x^2 - 2px - q = 0$ p and q are positive: p=12345678, q=1
 - **Exercise:** find the minimum of the roots.
- Subtracting approximations of two nearby numbers results in a bad approximation of the actual difference – **catastrophic cancellation**

Floating Point System - Terminology

- **Precision (p)** - Length of mantissa
 - E.g. $p=3$ in 1.00×10^{-1}
- **Machine epsilon (ϵ_{mach})** – smallest $a-1$, where a is the smallest representable number greater than 1
 - E.g. $\epsilon_{\text{mach}} = 1.001 - 1.000 = 0.001$.
- **Unit roundoff (u)** – smallest positive number where the *computed* value of $1+u$ is different from 1
 - E.g. suppose $p=4$ and we wish to compute $1.0000 + 0.0001 = 1.0001$
 - But we can't store the exact result (since $p=4$). We end up storing 1.000.
 - So, computed result of $1+u$ is same as 1
 - Suppose we tried adding 0.0005 instead. $1.0000 + 0.0005 = 1.0005$
Now, round this: 1.001

⇒ **$u = 0.0005$**

⇒ **usually $u = \frac{1}{2} * \epsilon_{\text{mach}}$**