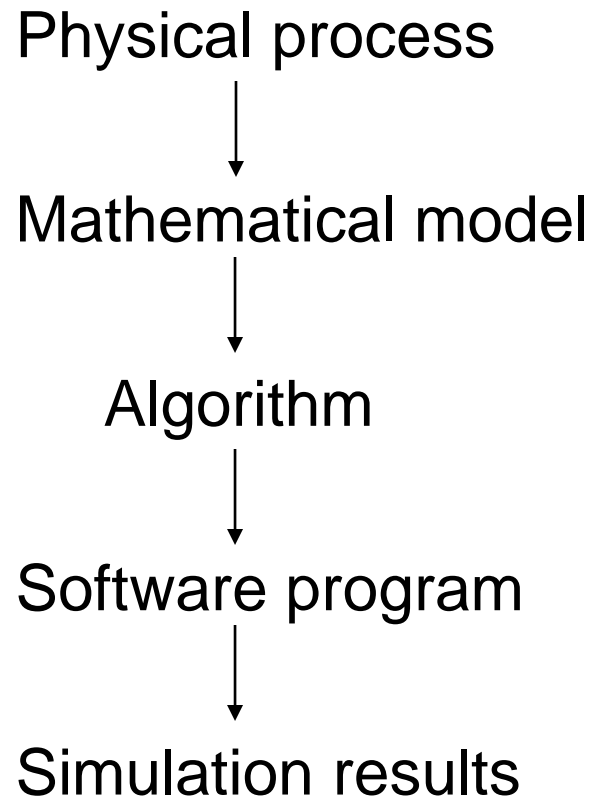


CS601: Software Development for Scientific Computing

Autumn 2023

Week2: Real Numbers, Programming
Environment, ..

Recap: Toward Scientific Software



Real Numbers \mathbb{R}

- Most scientific software deal with Real numbers.
Our toy code dealt with Reals
 - Numerical software is scientific software dealing with Real numbers
- Real numbers include rational numbers (integers and fractions), irrational numbers (pi etc.)
- Used to represent values of continuous quantity such as time, mass, velocity, height, density etc.
 - Infinitely many values possible
 - But computers have limited memory. So, have to use approximations.

Representing Real Numbers

- Real numbers are stored as *floating point numbers* (floating point system is a scheme to represent real numbers)
- E.g. floating point numbers:
 - $\pi = 3.14159$,
 - 6.03×10^{23}
 - $1.60217733 \times 10^{-19}$

General format: $\pm x \times b^e$

mantissa (number ranges from: 1 to b OR 1/b to 1)

base (e.g. base 10, 8, 2, 16)

exponent

3-digit Calculator

- Suppose base, $b=10$ and
- $x = \pm d_0.d_1d_2 \times 10^e$ where $\begin{cases} 1 \leq d_0 \leq 9, \\ 0 \leq d_1 \leq 9, \\ 0 \leq d_2 \leq 9 \\ -9 \leq e \leq 9 \end{cases}$
- precision = length of mantissa
 - What is the precision here?
- Exercise: What is the smallest positive number?
- Exercise: What is the largest positive number?
- Exercise: How many numbers can be represented in this format?
- Exercise: When is this representation not enough?

Floating Point System - Fundamentals

- **Precision (p)** - Length of mantissa
 - E.g. $p=3$ in 1.00×10^{-1}
- **Unit roundoff (u)** – smallest positive number where the *computed* value of $1+u$ is different from 1
 - E.g. suppose $p=4$ and we wish to compute $1.0000 + 0.0001 = 1.0001$
 - But we can't store the exact result (since $p=4$). We end up storing 1.000.
 - So, computed result of $1+u$ is same as 1
 - Suppose we tried adding 0.0005 instead. $1.0000 + 0.0005 = 1.0005$
Now, round this: 1.001
 - $\Rightarrow u = 0.0005$
- **Machine epsilon (ϵ_{mach})** – smallest $a-1$, where a is the smallest representable number greater than 1
 - E.g. consider $1.001 - 1.000 = 0.001$.
 - \Rightarrow **usually** $\epsilon_{\text{mach}} = 2 * u$

Floating Point System - Fundamentals

- **Forward error and backward error**

$$\text{Comp}(f(x)) = (1+\epsilon_1)f((1+\epsilon_2)x),$$

where $\epsilon_i \leq u$ (u is unit roundoff)

$\text{Comp}(f(x))$ is the computed value i.e. machine representable value of $f(x)$.

Suppose ϵ_2 is zero. Then
$$\frac{\text{Comp}(f(x)) - f(x)}{f(x)} = \epsilon_1$$

Floating Point System - Fundamentals

- **Forward error example**

Let $y = \sqrt{2}$, $z = y^2$ and

$y = \sqrt{2}$ implemented as: `y = sqrt(2);`

$z = y^2$ implemented as: `z = y * y;`

with double precision floating point system

Then forward error, $\frac{\{ \text{Comp}(f(x)) - f(x) \}}{f(x)}$, can be calculated

(note: $f(x) = z = 2$, and $\text{Comp}(f(x)) = y * y$)

```
y:1.41421356237
```

```
z:2
```

```
res1=z-2:4.4408920985e-16
```

```
res2=res1/z:2.22044604925e-16
```

**Absolute error /
relative error**

Forward error
(also happens to be u ,
unit roundoff, for
double)

Floating Point System - Fundamentals

- **Backward error example**

Let $z = \sin(2\pi)$. Then forward error is infinity!

Subtract x with a multiple of 2π to make $0 \leq x < 2\pi$

And then compute $\sin(x)$ to get the absolute error for $x \geq 2\pi$ at most $u|x|$ (u is unit roundoff)


This is *perturbing* the argument x (*argument reduction*). Instead of computing $\sin(x)$ we are computing $\sin((1 + \epsilon_2)x)$. This is example of backward error.

IEEE 754 Floating Point System

- Prescribes single, double, and extended precision formats


Precision	u	Total bits used (sign, exponent, mantissa)
Single	6×10^{-8}	32 (1, 8, 23)
Double	2×10^{-16}	64 (1, 11, 52)
Extended	5×10^{-20}	80 (1, 15, 64)

single precision binary IEEE 754 floating point format



IEEE 754 Floating Point System

double precision binary IEEE 754 floating point format




- if exponent bits e_1 - e_{11} are not all 1s or 0s, then the *normalized* number

$$n = \pm(1.m_1m_2..m_{52})_2 \times 2^{(e_1e_2..e_{11})_2 - 1023}$$

- Machine epsilon** is the gap between 1 and the next largest floating point number. $2^{-52} \approx 10^{-16}$ for double.
- Exercise: What is minimum positive *normalized* double number?
- Exercise: What is maximum positive *normalized* double number?

IEEE 754 Floating Point System

double precision binary IEEE 754 floating point format



- if exponent bits e_1 - e_{11} are all 0s, then:
the *subnormal* number

$$n = \pm(\mathbf{0}.m_1m_2..m_{52})_2 \times 2^{(e_1e_2..e_{11})_2 - 102\mathbf{2}}$$

- if exponent bits e_1 - e_{11} are all 1s, then:
we can get $-\text{inf}$, NaN, or $+\text{inf}$ based on value of $m_1m_2..m_{52}$
 - If any m is non-zero, the number is NaN (not a number)

IEEE 754 Floating Point – Misc..

- **+0, -0, Inf, and NaN –**
 - Stop your program when you see a NaN (indicative of a bug)
 - How to check if a number is NaN?
`if (x == x) is false`
Exercise: Give an example when you get a NaN?
- **Rounding modes –** Round up, Round down, Round to nearest, Round towards zero
 - Default is round to nearest. Can be set using compiler options and library methods. Avoid changing rounding modes.
 - Can use this to flush out bugs! (change round modes and results shouldn't change drastically).

IEEE 754 Floating Point Arithmetic

- Order is important
 - Floating point arithmetic is *not associative*
 - $(x+y)+z$ not the same as $x+(y+z)$
- Explicit coding of textbook formula may not be the best option to solve
 - $x^2 - 2px - q = 0$ p and q are positive: p=12345678, q=1
 - **Exercise:** find the minimum of the roots.
- Subtracting approximations of two nearby numbers results in a bad approximation of the actual difference – **catastrophic cancellation**