

CS601: Software Development for Scientific Computing

Autumn 2023

Week6: Matrix Computations with Sparse
Matrices, Tools for debugging and more

LAPACK – Linear Algebra Package

- LAPACK – uses BLAS-3 (1989 – now)
 - Ex: Obvious way to express Gaussian Elimination (GE) is adding multiples of one row to other rows – BLAS-1
 - How do we reorganize GE to use BLAS-3 ?
 - Contents of LAPACK (summary)
 - Algorithms that are (nearly) 100% BLAS-3
 - Linear Systems, Least Squares
 - Algorithms that are only $\approx 50\%$ BLAS-3
 - Eigenproblems, Singular Value Decomposition (SVD)
 - Generalized problems (eg $Ax = I Bx$)
 - Error bounds for everything
 - Lots of variants depending on A 's structure (banded, $A=A^T$, etc.)
 - How much code? (Release 3.9.0, Nov 2019) (www.netlib.org/lapack)
 - Source: 1982 routines, 827K LOC, Testing: 1210 routines, 545K LOC

Matrix Data and Efficiency

- Sparse Matrices
 - E.g. banded matrices
 - Diagonal
 - Tridiagonal etc.
 - Symmetric Matrices
- Admit optimizations w.r.t.* →
- Storage
 - Computation

Sparse Matrices - Motivation

- Matrix Multiplication with Upper Triangular Matrices
($C=C+AB$)

$$\begin{array}{ccc}
 \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ 0 & a_{22} & a_{23} \\ 0 & 0 & a_{33} \end{bmatrix} & \begin{bmatrix} b_{11} & b_{12} & b_{13} \\ 0 & b_{22} & b_{23} \\ 0 & 0 & b_{33} \end{bmatrix} & = \\
 \textcolor{blue}{A} & \textcolor{blue}{B} & \\
 & & \begin{bmatrix} a_{11}b_{11} & a_{11}b_{12}+a_{12}b_{22} & a_{11}b_{13}+a_{12}b_{23}+a_{13}b_{33} \\ 0 & a_{22}b_{22} & a_{22}b_{23}+a_{23}b_{33} \\ 0 & 0 & a_{33}b_{33} \end{bmatrix} \\
 & & \textcolor{blue}{A*B}
 \end{array}$$

The result, $A*B$, is also upper triangular.

The non-zero elements appear to be like the result of *inner-product*

Sparse Matrices - Motivation

- $C=C+AB$ when A, B, C are upper triangular, pseudocode:
for $i=1$ to N
 for $j=i$ to N
 for $k=i$ to j
 $C[i][j] = C[i][j] + A[i][k]*B[k][j]$
- Cost = $\sum_{i=1}^N \sum_{j=i}^N 2(j-i+1)$ flops (why 2?)
- Using $\sum_{i=1}^N i \approx \frac{n^2}{2}$ and $\sum_{i=1}^N i^2 \approx \frac{n^3}{3}$
- $\sum_{i=1}^N \sum_{j=i}^N 2(j-i+1) \approx \frac{n^3}{3}$, 1/3rd the number of flops required for dense matrix-matrix multiplication

Sparse Matrices

- Have lots of zeros (a *large* fraction)

X	X	0	0	X	0	0	0	X
0	X	0	0	X	0	X	0	0
0	X	X	X	0	X	0	0	X
X	0	0	X	0	0	X	0	0
0	X	0	X	X	0	0	0	X
0	X	X	0	0	0	X	X	X

- Representation
 - Many formats available
 - Compressed Sparse Row (CSR)

Implementation: Three arrays:

```
double *val;  
int *ind;  
int *rowstart;
```

Sparse Matrices - Example

- Using Arrays

A

a_{11}	a_{12}	0	0	a_{15}	0	0	0	a_{19}
0	a_{22}	0	0	a_{25}	0	a_{27}	0	0
0	a_{32}	a_{33}	a_{34}	0	a_{36}	0	0	a_{39}
a_{41}	0	0	a_{44}	0	0	a_{47}	0	0
0	a_{52}	0	a_{54}	a_{55}	0	0	0	a_{59}
0	a_{62}	a_{63}	0	0	0	a_{67}	a_{68}	a_{69}

```
double *val; //size= NNZ
int *ind; //size=NNZ
int *rowstart; //size=M=Number of rows
```

val:

a_{11}	a_{12}	a_{15}	a_{19}	a_{22}	a_{25}	a_{27}	a_{32}	a_{33}	a_{34}	a_{36}	a_{39}	a_{41}	a_{44}	a_{47}	a_{52}	a_{54}	a_{55}	a_{59}	a_{62}	a_{63}	a_{67}	a_{68}	a_{69}
----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------

ind:

1	2	5	9	2	5	7	2	3	4	6	9	1	4	7	2	4	5	9	2	3	7	8	9
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

rowstart:

0	4	7	12	15	19	
---	---	---	----	----	----	--

Nikhil Hegde

Gaxpy with Sparse Matrices: $y=y+Ax$

- Using arrays

```
for i=0 to numRows
```

```
    for j=rowstart[i] to rowstart[i+1]-1
```

```
        y[i] = y[i] + val[j]*x[ind[j]]
```

- Does the above code reuse y , x , and val ? (we want our code to reuse as much data elements as possible while they are in fast memory):
 - y ? Yes. Read and written in close succession.
 - x ? Possible. Depends on how data is scattered in val .
 - val ? Good spatial locality here. Less likely for a sparse matrix in general.

Gaxpy with Sparse Matrices: $y=y+Ax$

- Optimization strategies:

```
for i=0 to numRows
```

```
    for j=rowstart[i] to rowstart[i+1]-1
```

```
        y[i] = y[i] + val[j]*x[ind[j]]
```

- Unroll the j loop // we need to know the number of non-zeros per row
- Eliminate ind[i] and thereby the indirect access to elements of x. Indirect access is not good because we cannot predict the pattern of data access in x. //We need to know the column numbers
- Reuse elements of x //The elements of a should be e.g. located closely

These optimizations will not work for $y=y+Ax$ pseudocode in general. When you know the data pattern and metadata info as mentioned above, you can reorder computations (scheduling optimization), reorganize data for better locality.

Banded Matrices

- Special case of sparse matrices, characterized by two numbers:

- Lower bandwidth p , and upper bandwidth q
- E.g. $p=1, q=2$

for a 8×5 matrix

(x represents non-zero element)

x	x	x	0	0
x	x	x	x	0
0	x	x	x	x
0	0	x	x	x
0	0	0	x	x
0	0	0	0	x
0	0	0	0	0
0	0	0	0	0

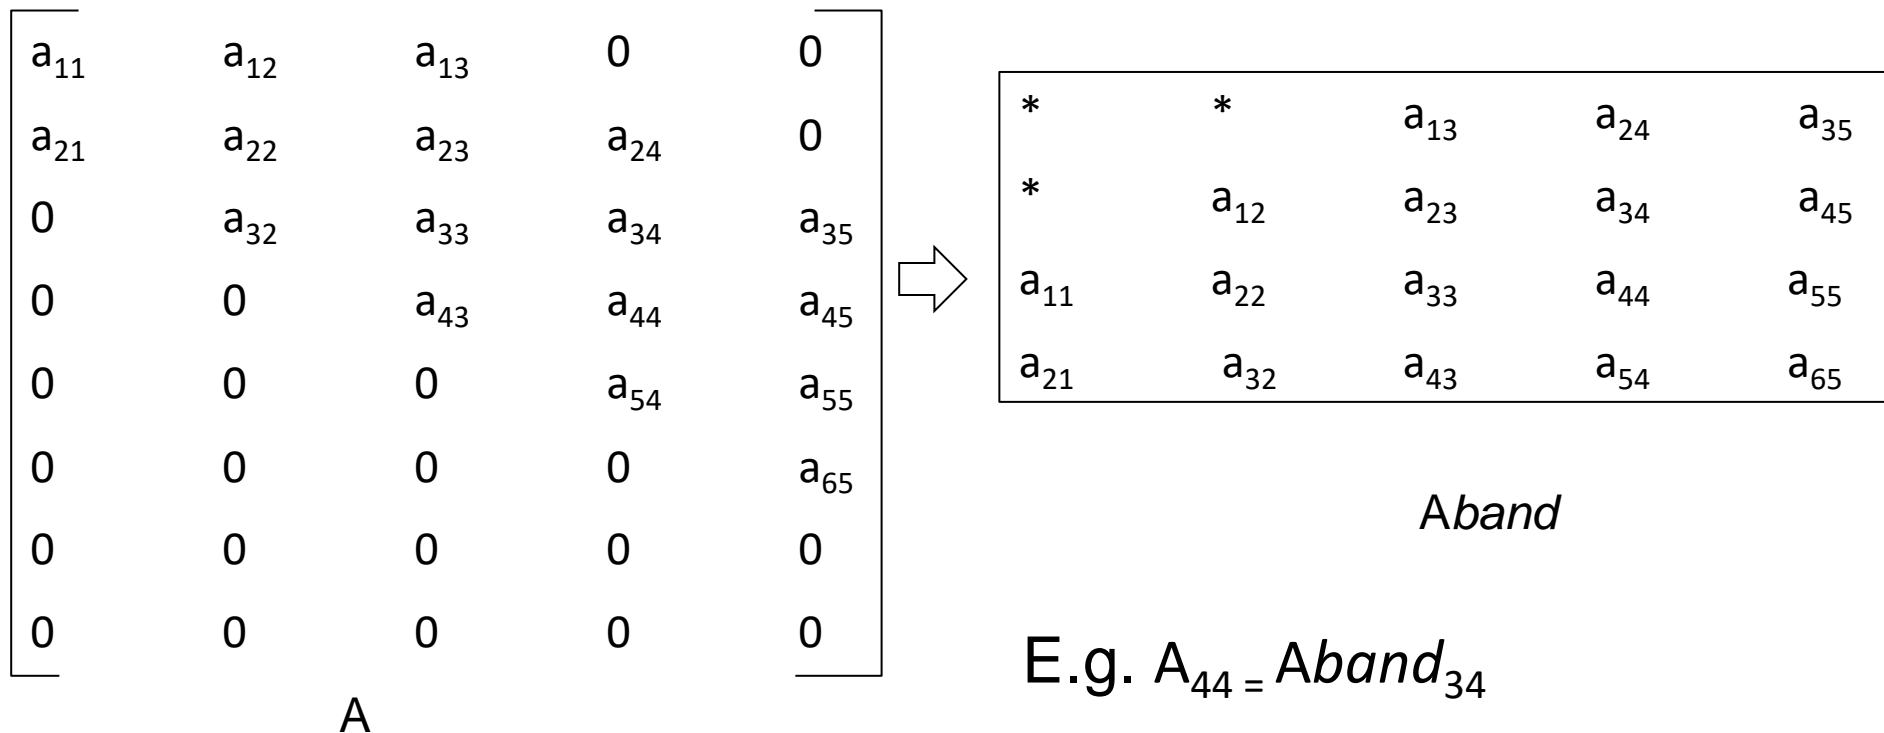
Exercise: When is $a_{ij} = 0$?

(Write the constraints in terms of i, j, p, q)

- $a_{ij} = 0$ if $i > j+p$
- $a_{ij} = 0$ if $j > i+q$

Banded Matrices - Representation

- Optimizing storage (specific to banded matrices)



Exercise: $A_{ij} = Aband(i-j+q+1, j)$

Gaxpy with Banded Matrices: $y = y + A \text{band } x$

- $A=A_{\text{band}}$: optimizing computation and storage

```
for j=1 to n
    alpha1=max(1, j-q)
    alpha2=min(n, j+p)
    beta1=max(1, q+2-j)
    for i=alpha1 to alpha2
         $y[i] = y[i] + A_{\text{band}}(\text{beta1}+i-\text{alpha1}, j) * x[j]$ 
```

- Cost? $2n(p+q+1)$ time! Much lesser than $2N^2$ time required for regular $y=y+Ax$ (assuming p and q are much smaller than n)

Tools

- Debugging
- Profiling
- Documenting

GDB

- GNU Debugger – A tool for inspecting your C/C++ programs
 - How to begin inspecting a program using gdb?
 - How to control the execution?
 - How to display, interpret, and alter memory contents of a program using gdb?
 - Misc – displaying stack frames, visualizing assembler code.