

CS601: Software Development for Scientific Computing

Autumn 2024

Week7: Tools for profiling, debugging, and more..

Valgrind

- Suite of tools for debugging and profiling
 - memcheck and cachegrind are popular ones
 - cachegrind is cache and branch-prediction profiler.
 - memcheck is a memory error detector.
- Demo of cachegrind tool with matmul
 - <https://valgrind.org/docs/manual/cg-manual.html>
- Demo of memcheck with matmul

Steps to use cachegrind

- Example: `matmul.cpp`
 1. Compile with `-g` and create a target.
 2. Run as: `valgrind --tool=cachegrind ./matmul 2048`
 3. Output of `cachegrind` is dumped in a file that has the format `cachegrind.out.xxxxxx` where `xxxxxx` is the process ID
 4. Use `cg_annotate` to get annotated output
 1. E.g. `cg_annotate cachegrind.out.12345`

cachegrind

- Visualizing cache transactions

L1 Instruction → I1 cache: 32768 B, 64 B, 8-way associative
L1 Data → D1 cache: 32768 B, 64 B, 8-way associative
Last layer → LL cache: 37748736 B, 64 B, 18-way associative
Command: ./matmul_ijk 2048
Data file: cachegrind.out.1395356
Events recorded: Ir I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Events shown: Ir I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Event sort order: Ir I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Thresholds: 0.1 100 100 100 100 100 100 100 100
Include dirs:
User annotated:
Auto-annotation: on

- Instructions read
- L1 Instruction read misses
- Last layer instruction read misses
- Data reads (total memory reads)
- L1 data read misses
- Last layer data read misses
- Data writes (total memory writes)
- L1 data write misses
- Last layer data write misses

Total last layer misses = ILmr + DLmr + DLmw

cachegrind

- Visualizing cache transactions (**ijk** loop ordering of matmul)

Ir	I1mr (L1 read miss)	ILmr (LL instruction read miss)	Dr (Data read == number of memory reads)
438,803,764,234 (100.0%)	2,267 (100.0%)	2,157 (100.0%)	189,231,226,540 (100.0%)

D1mr (L1 Data read miss)	DLmr (LL data read misses)
10,740,872,902 (100.0%)	7,827,585,951 (100.0%)

Dw (Data write = number of memory writes)	D1mw (L1 data cache write miss)	DLmw (LL data write miss)
8,674,338,548 (100.0%)	1,586,278 (100.0%)	1,582,786 (100.0%)

cachegrind

- Visualizing cache transactions (**ikj** loop ordering of matmul)

```
-----  
Ir                               I1mr (L1 read miss)          I1Lmr (LL instruction read miss)      Dr (Data read == number of memory reads)  
-----  
438,803,764,251 (100.0%) 2,267 (100.0%) 2,157 (100.0%) 189,231,226,544 (100.0%)  
  
D1mr (L1 Data read miss)          D1Lmr (LL data read misses)  
1,223,946,667 (100.0%) 1,004,088,043 (100.0%)  
  
Dw (Data write = number of memory writes)          D1mw (L1 data cache write miss)          D1Lmw (LL data write miss)  
8,674,338,550 (100.0%) 1,586,278 (100.0%) 1,582,786 (100.0%)
```

Total last layer misses are much lesser than that in ijk loop!

GNU gprof

- Usage:
 - Compile your program with `-pg` flag
 - Execute your program as normal
 - A file `gmon.out` is generated
 - `gprof <yourexecutable>`