

CS601: Software Development for Scientific Computing

Autumn 2023

Week7: Tools for debugging and profiling and
more..

Valgrind

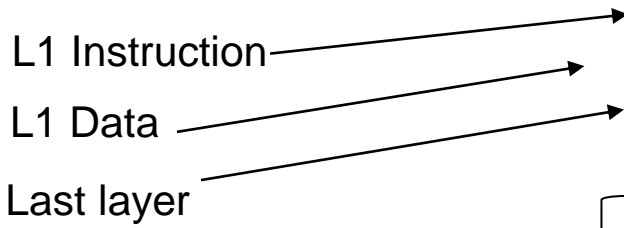
- Suite of tools for debugging and profiling
 - memcheck and cachegrind are popular ones
 - cachegrind is cache and branch-prediction profiler.
 - memcheck is a memory error detector.
- Demo of cachegrind tool with matmul
 - <https://valgrind.org/docs/manual/cg-manual.html>
- Demo of memcheck with matmul

Steps to use cachegrind

- Example: `matmul.cpp`
 1. Compile with `-g` and create a target.
 2. Run as: `valgrind --tool=cachegrind ./matmul 2048`
 3. Out of cachegrind is dumped in a file that has the format `cachegrind.out.xxxxxx` where `xxxxxx` is the process ID
 4. Use `cg_annotate` to get annotated output
 1. E.g. `cg_annotate cachegrind.out.12345`

cachegrind

- Visualizing cache transactions



```

L1 Instruction → I1 cache:
L1 Data → D1 cache:
Last layer → LL cache:
Command:
Data file:
Events recorded:
Events shown:
Event sort order:
Thresholds:
Include dirs:
User annotated:
Auto-annotation: on

```

32768 B, 64 B, 8-way associative
 32768 B, 64 B, 8-way associative
 37748736 B, 64 B, 18-way associative
 ./matmul_ijk 2048
 cachegrind.out.1395356
 Ir I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
 Ir I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
 Ir I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
 0.1 100 100 100 100 100 100 100 100
 0.1 100 100 100 100 100 100 100 100
 on

- Instructions read
- L1 Instruction read misses
- Last layer instruction read misses
- Data reads (total memory reads)
- L1 data read misses
- Last layer data read misses
- Data writes (total memory writes)
- L1 data write misses
- Last layer data write misses

Total last layer misses = ILmr + DLmr + DLmw

cachegrind

- Visualizing cache transactions (**ijk** loop ordering of matmul)

Ir	I1mr (L1 read miss)	ILmr (LL instruction read miss)	Dr (Data read == number of memory reads)
438,803,764,234 (100.0%)	2,267 (100.0%)	2,157 (100.0%)	189,231,226,540 (100.0%)

D1mr (L1 Data read miss)	DLmr (LL data read misses)
10,740,872,902 (100.0%)	7,827,585,951 (100.0%)

Dw (Data write = number of memory writes)	D1mw (L1 data cache write miss)	DLmw (LL data write miss)
8,674,338,548 (100.0%)	1,586,278 (100.0%)	1,582,786 (100.0%)

cachegrind

- Visualizing cache transactions (**ikj** loop ordering of matmul)

```
-----  
Ir                               I1mr (L1 read miss)          I1mr (LL instruction read miss)      Dr (Data read == number of memory reads)  
-----  
438,803,764,251 (100.0%) 2,267 (100.0%) 2,157 (100.0%) 189,231,226,544 (100.0%)  
  
D1mr (L1 Data read miss)          D1mr (LL data read misses)  
1,223,946,667 (100.0%) 1,004,088,043 (100.0%)  
  
Dw (Data write = number of memory writes)          D1mw (L1 data cache write miss)          D1mw (LL data write miss)  
8,674,338,550 (100.0%) 1,586,278 (100.0%) 1,582,786 (100.0%)
```

Total last layer misses are much lesser than that in ijk loop!

Memcheck – ex1

- Used for detecting memory error that include memory leaks and invalid read/write to memory

```
//Example 1
void CreateAndAddMatrices(int n){
    float *p = new float[n*n]; // allocate a matrix, p, of float elements
    for(int i=0;i<n*n;i++){
        p[i]=i;
    }
    float *q = new float[n*n]; // allocate a matrix, q, of float elements
    for(int i=0;i<n*n;i++){
        q[i]=i;
    }
    float *r = new float[n*n]; // allocate a matrix, r, of float elements
    for(int i=0;i<n*n;i++)
        r[i]=p[i]+q[i]; //do r = p + q

    return ;
}

int main(int argc, char* argv[]){
    //Example 1
    CreateAndAddMatrices(16); //this function leaks memory. Exercise: fix the leak.
```

memcheck - ex2

//Example 2

```
float* CreateAndAddMatricesV2(int n){
    float *p = new float[n*n]; // allocate a matrix, p, of float elements
    for(int i=0;i<n*n;i++){
        p[i]=i;
    }
    float *q = new float[n*n]; // allocate a matrix, q, of float elements
    for(int i=0;i<n*n;i++){
        q[i]=i;
    }
    float *r = new float[n*n]; // allocate a matrix, r, of float elements
    for(int i=0;i<n*n;i++)
        r[i]=p[i]+q[i]; //do r = p + q

    delete [] p;
    delete [] q;
    delete [] r;

    return r;
}
```

```
int main(int argc, char* argv){
```

//Example 2

```
float* result=CreateAndAddMatricesV2(16); //this function releases memory to early. Exercise: fix the error.
```


memcheck – ex3

//Example 3

```
float** CreateAndAddMatricesU3(int n){
    float *p = new float[n*n]; // allocate a matrix, p, of float elements
    for(int i=0;i<n*n;i++){
        p[i]=i;
    }
    float *q = new float[n*n]; // allocate a matrix, q, of float elements
    for(int i=0;i<n*n;i++){
        q[i]=i;
    }
    float *r = new float[n*n]; // allocate a matrix, r, of float elements
    for(int i=0;i<n*n;i++)
        r[i]=p[i]+q[i]; //do r = p + q

    float **s = new float*; // allocate an element to store the handle for matrix r
    *s = r;

    delete [] p;
    delete [] q;
    //not sure if I should release the memory allocated for r or not.

    return s; //s is not released because it is being returned.
}

int main(int argc, char* argv[]){
    //Example 3
    float** result2=CreateAndAddMatricesU3(16); //In this example, we do not know whether it is safe to release memory
    (*result2)[0]=1.234; //sets the (0,0) element of matrix r to 1.234.
    //assume that you are done using the r matrix.
    (*result2)=NULL; //reset so that result can hold a handle to some other matrix. This is a problem. Exercise: fix the error.
```

memcheck - Usage

- Compile with `-g` option and create a target
- Execute with `valgrind`

```
valgrind -tool=memcheck -leak-check=full mytarget
```

<https://valgrind.org/docs/manual/mc-manual.html>

memcheck - Demo

- From week7 code samples, run:

```
make -f memchkMakefile example1
```

```
==664== LEAK SUMMARY:
==664==    definitely lost: 3,072 bytes in 3 blocks
==664==    indirectly lost: 0 bytes in 0 blocks
==664==    possibly lost: 0 bytes in 0 blocks
==664==    still reachable: 0 bytes in 0 blocks
==664==    suppressed: 0 bytes in 0 blocks
==664== total heap usage: 4 allocs, 1 frees, 75,776 bytes allocated
==664==
==664== 1,024 bytes in 1 blocks are definitely lost in loss record 1 of 3
==664==    at 0x483C583: operator new[](unsigned long) (in /usr/lib/x86_64-linux-g
nu/valgrind/vgpreload_memcheck-amd64-linux.so)
==664==    by 0x1091DA: CreateAndAddMatrices(int) (memerrors.cpp:10)
==664==    by 0x10932F: main (memerrors.cpp:79)
==664==
==664== 1,024 bytes in 1 blocks are definitely lost in loss record 2 of 3
==664==    at 0x483C583: operator new[](unsigned long) (in /usr/lib/x86_64-linux-g
nu/valgrind/vgpreload_memcheck-amd64-linux.so)
==664==    by 0x10923D: CreateAndAddMatrices(int) (memerrors.cpp:14)
==664==    by 0x10932F: main (memerrors.cpp:79)
==664==
==664== 1,024 bytes in 1 blocks are definitely lost in loss record 3 of 3
==664==    at 0x483C583: operator new[](unsigned long) (in /usr/lib/x86_64-linux-g
nu/valgrind/vgpreload_memcheck-amd64-linux.so)
==664==    by 0x1092A0: CreateAndAddMatrices(int) (memerrors.cpp:18)
==664==    by 0x10932F: main (memerrors.cpp:79)
```

memcheck - Demo

- From week7 code samples, run:

`make -f memchkMakefile example3`

```
==671== LEAK SUMMARY:
==671==    definitely lost: 1,032 bytes in 2 blocks
==671==    indirectly lost: 0 bytes in 0 blocks
==671==    possibly lost: 0 bytes in 0 blocks
==671==    still reachable: 0 bytes in 0 blocks
==671==    suppressed: 0 bytes in 0 blocks

==671== HEAP SUMMARY:
==671==    in use at exit: 1,032 bytes in 2 blocks
==671==    total heap usage: 5 allocs, 3 frees, 75,784 bytes allocated
==671==
==671== 8 bytes in 1 blocks are definitely lost in loss record 1 of 2
==671==    at 0x483BE63: operator new(unsigned long) (in /usr/lib/x86_64-linux-gnu/libc.so.6)
==671==    by 0x109359: CreateAndAddMatricesV3(int) (memerrors.cpp:65)
==671==    by 0x1093B1: main (memerrors.cpp:87)
==671==
==671== 1,024 bytes in 1 blocks are definitely lost in loss record 2 of 2
==671==    at 0x483C583: operator new[](unsigned long) (in /usr/lib/x86_64-linux-gnu/libc.so.6)
==671==    by 0x1092E0: CreateAndAddMatricesV3(int) (memerrors.cpp:61)
==671==    by 0x1093B1: main (memerrors.cpp:87)
```

memcheck - Demo

- From week7 code samples, run:

```
make -f memchkMakefile example4
```

```
==678== Invalid write of size 1
==678==    at 0x483F0BE: strcpy (in /usr/lib/x86_64-linux-gnu/valgrind/vgpreload_m
emcheck-amd64-linux.so)
==678==    by 0x109231: main (memerrors.cpp:96)
==678== Address 0x4da7c85 is 0 bytes after a block of size 5 alloc'd
==678==    at 0x483C583: operator new[](unsigned long) (in /usr/lib/x86_64-linux-g
nu/valgrind/vgpreload_memcheck-amd64-linux.so)
==678==    by 0x10921A: main (memerrors.cpp:95)
==678==
==678==
==678== HEAP SUMMARY:
==678==    in use at exit: 0 bytes in 0 blocks
==678== total heap usage: 2 allocs, 2 frees, 72,709 bytes allocated
==678==
==678== All heap blocks were freed -- no leaks are possible
```

memcheck - Demo

- From week7 code samples, run:

```
make -f memchkMakefile example5
```

```
==685== Invalid read of size 1
==685==    at 0x483EF54: strlen (in /usr/lib/x86_64-linux-gnu/valgrind/vgpreload_m
memcheck-amd64-linux.so)
==685==    by 0x4AB0E94: __vfprintf_internal (vfprintf-internal.c:1688)
==685==    by 0x4A99EBE: printf (printf.c:33)
==685==    by 0x109208: main (memerrors.cpp:102)
==685== Address 0x4da7c81 is 0 bytes after a block of size 1 alloc'd
==685==    at 0x483BE63: operator new(unsigned long) (in /usr/lib/x86_64-linux-gnu
/valgrind/vgpreload_memcheck-amd64-linux.so)
==685==    by 0x1091E5: main (memerrors.cpp:100)
==685==
printing p: A
==685==
==685== HEAP SUMMARY:
==685==    in use at exit: 0 bytes in 0 blocks
==685==    total heap usage: 3 allocs, 3 frees, 73,729 bytes allocated
==685==
==685== All heap blocks were freed -- no leaks are possible
```