

CS601: Software Development for Scientific Computing

Autumn 2021

Week13:

Hierarchical Methods (FMM) and Sparse Matrices

Course Progress..

- Last week
 - Tree-based codes (hierarchical methods)
 - Barnes-Hut
 - Fast Multipole Method (FMM)
- This week
 - FMM
 - Sparse matrices and
 - PA4 discussion


FMM Algorithm

1. Build the quadtree containing all the points.
2. Traverse the quadtree from bottom to top, computing $\text{Outer}(n)$ for each square n in the tree.
3. Traverse the quadtree from top to bottom, computing $\text{Inner}(n)$ for each square in the tree.
4. For each leaf, add the contributions of nearest neighbors and particles in the leaf to $\text{Inner}(n)$

what is $\text{Outer}(n)$ and $\text{Inner}(n)$?

Well Separated Regions

- Compute the influence of all particles in source region (B) on every particle in target region (A)
(assumption: A and B are well-separated)



- At each point p_i in A, compute potential:


$$\Phi(x_i, y_i) = \sum_{p_j \in B} m_i \log |p_i - p_j|$$
$$i = 1 \text{ to } N_A, \quad j = 1 \text{ to } N_B$$

- Cost: $O(N_A N_B)$

Well Separated Regions

- Compute the influence of all particles in source region (B) on every particle in target region (A)

$$\Phi(x_{p_i}, y_{p_i}) = \sum_{p_j \in B} m_j \log |p_i - p_j|, p_i \in A$$



Applying the 3-step Approximation


- In N-body simulation every point serves as source as well as target.

How to identify source and target (boxes A and B in previous slide) i.e. well-separated regions?


Hierarchical decomposition

Hierarchical Decomposition

- *Level-0 decomposition*



- *Level-1 decomposition*



No well-separated boxes

Hierarchical Decomposition

- *Level-2 decomposition*

N1	N2	N3	A1
N4	B	N5	A2
N6	N7	N8	A3
A7	A6	A5	A4



Well-separated from B

Can approximate the influence of points in B on points in A_i s

What do we do about **B**'s influence on N_i s?

Hierarchical Decomposition

- *Level-3 decomposition*

N1	N2	N3	A1
N4	B1 B2 B3 B4	N5	A2
N6	N7	N8	A3
A7	A6	A5	A4



Influence of points in B_i s on those in A_i s
already computed at the previous level
(level-2)

Hierarchical Decomposition

- Level-3 decomposition*

n1	n2	n5	n6	n9	n10				
n3	n4	n7	n8	n11	n12			A1	
n13	n14	B1	B2		n27			A2	
n15	n16	B3	B4		n26				
n17	n18				n25			A3	
n19	n20	n21	n22	n23	n24				
								A4	
A7									
		A6							
				A5					



Influence of points in B_i s on those in A_i s already computed at the previous level (level-2)



Well-separated from B_4

Influence of B_4 's points on n_x 's points can be approximated

n_x 's constitute the interaction list for B_4 .

What is the max size of interaction list? i.e. max number of n_x s that we can have for any B_i ?

Hierarchical Decomposition

- *Level-3 decomposition*

n1	n2	n5	n6	n9	n10		
n3	n4	n7	n8	n11	n12		
n13	n14	B1	B2		n27		
n15	n16	B3	B4		n26		
n17	n18				n25		
n19	n20	n21	n22	n23	n24		



Influence of points in B_i s on those in A_i s already computed at the previous level (level-2)




Well-separated from B_4



Influence of B_4 's points on n_x 's points can be approximated

What do we do about **B4**'s influence on its neighbors (white/unshaded boxes)?

Hierarchical Decomposition

- *Level-4 decomposition*



Any unshaded box outside  can be the *target* for computing the influence of points in  (*source*)

1. Computing Potential for Well-Separated Regions

```
1. for level L=2 to last_level
2.   for each Box B at level L
3.     iList = GetInteractionList(B)
4.     for each well-separated box A in iList
5.       //Compute potential
6.       potential =  $m_B \log |C_A - C_B|$ 
7.       //Accumulate potential
8.        $\Phi(x_{C_A}, y_{C_A}) += \text{potential}$ 
```

Cost?

1. Computing Potential for Well-Separated Regions

```
1. for level L=2 to last_level
2.   for each Box B at level L
3.     iList = GetInteractionList(B)
4.     for each well-separated box A in iList
5.       //Compute potential
6.       potential =  $m_B \log |C_A - C_B|$ 
7.       //Accumulate potential
8.        $\Phi(x_{C_A}, y_{C_A}) += \text{potential}$ 
```

Prereqs: we need m_B, C_A, C_B details. (step 0)

2. Assigning Potential to Points

1. **for** each Box A at level $L=0$ to last_level
2. $\Phi_{p_i} = \Phi_{p_i} + \Phi_{C_A}$ (where $p_i \in A$ and C_A is A 's CM)

Cost?

3. Assigning Potential to Points (last level)

1. **for** each Box B at `last_level`
2. $\Phi_{p_i} = \Phi_{p_i} + \sum_{p_j \in \text{Neighbors}(B)} m_B \log |p_i - p_j|$ (where $p_i \in B$)

Cost?

0. Computing Prereqs

1. **for** each Box B at level L=0 to last_level
2. $m_B = \sum_{p_j \in B} m_j$
3. //similarly compute C_B

Cost?

Total Cost (steps 0 + 1 + 2 + 3)

$$O(N \log N) + O(N) + O(N \log N) + O(N)$$

Can we do better?

0'. Computing Prereqs

- Traverse the tree bottom up instead of top-down
for each Box B starting from `last_level` to $L=0$
 if B is a leaf box

$$m_B = \sum_{p_j \in B} m_j$$

 else

$$m_B = m_{B_1} + m_{B_2} + m_{B_3} + m_{B_4}$$

 // B_1 - B_4 are children of B

Cost?

2'. Assigning Potential to Points

1. **for** each Box A at level $L=0$ to last_level

2. **if** A is a leaf box

$$\Phi_{p_i} = \Phi_{p_i} + \Phi_{C_A} \quad (\text{where } p_i \in A \text{ and } C_A \text{ is } A\text{'s CM})$$

else

$$\Phi_{A_1} = \Phi_{A_1} + \Phi_A$$

$$\Phi_{A_2} = \Phi_{A_2} + \Phi_A$$

$$\Phi_{A_3} = \Phi_{A_3} + \Phi_A$$

$$\Phi_{A_4} = \Phi_{A_4} + \Phi_A$$

 // A_1 - A_4 are children of A

Cost?

Total Cost (steps 0' + 1 + 2' + 3)


$$O(N) + O(N) + O(N) + O(N)$$

Problem: low accuracy if source (A) and target (B) are not far away from each other

Solution: more accurate representations for m_B and $\Phi(x_{C_A}, y_{C_A})$

Multipole expansion

- Like a Taylor series expansion that is accurate when $x^2 + y^2$ is large (x, y are cartesian coordinates of the point)
- For a quadtree box B centered at (x_{C_B}, y_{C_B}) , we compute and store the terms: $\{m_B, \alpha_1, \alpha_2, \dots, \alpha_p, z_{C_B}\}$




$$\alpha_j = \sum_{i=1}^{N_B} m_i \left(\frac{z_i^j}{j} \right)$$

z_i means $|z_i| = |(x_i, y_i)|$

Multipole expansion

- We approximate the potential at point z due to B by:




$$\Phi(x_z, y_z) = m_B \log(z - C_B) + \frac{\alpha_1}{z - C_B} + \frac{\alpha_2}{(z - C_B)^2} + \dots + \frac{\alpha_p}{(z - C_B)^p}$$

- Because $\{m_B, \alpha_1, \alpha_2, \dots, \alpha_p, z_{C_B}\}$ is used to compute potential outside B , it is called outer expansion

Multipole expansion

- Similarly, we have the inner expansion $\{m_B, \beta_1, \beta_2, \dots, \beta_p, z_{C_B}\}$ for computing the potential inside the Box due to all other points outside the box
- Computing outer expansions starts from leaf nodes and proceeds upwards in the tree.
- Computing inner expansions starts from root node and proceeds downwards in the tree.

3-Step Approximation (accurate)



FMM Algorithm

1. Build the quadtree containing all the points.
2. Traverse the quadtree from bottom to top, computing $\text{Outer}(n)$ for each square n in the tree.
3. Traverse the quadtree from top to bottom, computing $\text{Inner}(n)$ for each square in the tree.
4. For each leaf, add the contributions of nearest neighbors and particles in the leaf to $\text{Inner}(n)$

Multipole expansion

- How to obtain the expression for alpha, beta ?
- What is the value of p ?
- How to compute alpha and beta?
- Further reading:
<https://people.eecs.berkeley.edu/~demmel/cs267/lecture27/lecture27.html>

Matrix Algebra and Efficient Computation

- Pic source: the Parallel Computing Laboratory at U.C. Berkeley: A Research Agenda Based on the Berkeley View (2008)

<i>Motif</i>	Embed	Desktop	Games	DB	ML	HPC	Medicine	Music	Speech	CBIR	Browser	Motif	Desktop	Games	DB	ML	HPC	Medicine	Music	Speech	CBIR	Browser	
1 Finite State Mach.												9 N-Body											
2 Combinational												10 MapReduce											
3 Graph Traversal												11 Backtrack/B&B											
4 Structured Grid												12 Graphical Models											
5 Dense Matrix												13 Unstructured Grid											
6 Sparse Matrix												<i>Temperature Chart of Need</i>			DB = database								
7 Spectral (FFT)												<i>Hot</i>	<i>Warm</i>	<i>Med</i>	<i>Cool</i>	ML = machine learning							
8 Dynamic Prog																	HPC = High Perf. Comp.						

Figure 4. Temperature Chart of the 13 Motifs. It shows their importance to each of the original six application areas and then how important each one is to the five compelling applications of Section 3.1. More details on the motifs can be found in (Asanovic, Bodik et al. 2006).

Seen earlier
Next..

Matrix Multiplication

- Why study?
 - An important “kernel” in many linear algebra algorithms
 - Most studied kernel in high performance computing
 - Simple. Optimization ideas can be applied to other kernels
- Matrix representation
 - Matrix is a 2D array of elements. Computer memory is inherently linear
 - C++ and Fortran allow for definition of 2D arrays. 2D arrays stored row-wise in C++. Stored column-wise in Fortran. E.g.
`// stores 10 arrays of 20 doubles each in C++`
`double** mat = new double[10][20];`

Storage Layout - Example

- Matrix (**2D**): $A = \begin{bmatrix} A(0,0) & A(0,1) & A(0,2) \\ A(1,0) & A(1,1) & A(1,2) \\ A(2,0) & A(2,1) & A(2,2) \end{bmatrix}$

$A(i, j) = A(\text{row}, \text{column})$ refers to the matrix element in the i^{th} row and the j^{th} column

- Row-wise (/Row-major) storage in memory:

$A(0,0)$	$A(0,1)$	$A(0,2)$	$A(1,0)$	$A(1,1)$	$A(1,2)$	$A(2,0)$	$A(2,1)$	$A(2,2)$
----------	----------	----------	----------	----------	----------	----------	----------	----------


- Column-wise (/Column-major) storage in memory:

$A(0,0)$	$A(1,0)$	$A(2,0)$	$A(0,1)$	$A(1,1)$	$A(2,1)$	$A(0,2)$	$A(1,2)$	$A(2,2)$
----------	----------	----------	----------	----------	----------	----------	----------	----------

- Generalizing data storage order for ND:** last index changes fastest in row-major. Last index changes slowest in col-major.

Storage Layout - Exercise

- For a 3D array (tensor) assume $A(i, j, k) = A(\text{row}, \text{column}, \text{depth})$



- What is the offset of $A(1, 2, 1)$? as per row-major storage?
- What is the offset of $A(1, 2, 1)$? as per col-major storage?

Storage Layout

- Layout format itself doesn't influence efficiency (i.e. no general answer to “is column-wise layout better than row-wise?”)
- However, knowing the layout format is critical for good performance
 - *Always traverse the data in the order in which it is laid out*

How good performance?

Run on (12 X 2592.01 MHz CPU s)

CPU Caches:

L1 Data 32 KiB (x6)

L1 Instruction 32 KiB (x6)

L2 Unified 256 KiB (x6)

L3 Unified 12288 KiB (x1)

Load Average: 0.07, 0.02, 0.07


Source code: <https://github.com/eliben/code-for-blog/tree/master/2015/benchmark-row-col-major>

Benchmark	Time	CPU	Iterations	UserCounters...
BM_AddByRow/64/64	693 ns	693 ns	1042737	items_per_second=5.91004G/s
BM_AddByRow/128/128	2464 ns	2464 ns	271766	items_per_second=6.64813G/s
BM_AddByRow/256/256	11134 ns	11133 ns	63210	items_per_second=5.88639G/s
BM_AddByRow/512/512	44353 ns	44353 ns	15576	items_per_second=5.91041G/s
BM_AddByCol/64/64	3270 ns	3270 ns	212929	items_per_second=1.25254G/s
BM_AddByCol/128/128	39741 ns	39741 ns	17617	items_per_second=412.272M/s
BM_AddByCol/256/256	314880 ns	314878 ns	2241	items_per_second=208.132M/s
BM_AddByCol/512/512	1276733 ns	1276723 ns	545	items_per_second=205.326M/s

```
des/week13_codesamples$ ./a.out 4096
Rowwise time n=4096 (us): 18967
Colwise time n=4096 (us): 158608
nikhilh@ndhpc01:/mnt/c/temp/Nikhil/Cou
des/week13_codesamples$ ./a.out 2048
Rowwise time n=2048 (us): 4860
Colwise time n=2048 (us): 32158
nikhilh@ndhpc01:/mnt/c/temp/Nikhil/Cou
des/week13_codesamples$ ./a.out 1024
Rowwise time n=1024 (us): 1125
Colwise time n=1024 (us): 1980
```



Matrix-Matrix Addition benchmarking
([Source code and further reading](#))



Matvec execution time
(we used the [source code](#) as a
basic example to demonstrate row_major vs.
col_major storage.)

Matrix Multiplication

- Three fundamental ways to think of the algorithm

- Dot product

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \times \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix} = \begin{bmatrix} 1.5 + 2.7 & 1.6 + 2.8 \\ 3.5 + 4.7 & 3.6 + 4.8 \end{bmatrix}$$

- Linear combination of left matrix

$$\begin{bmatrix} \textcolor{blue}{1} & \textcolor{green}{2} \\ \textcolor{blue}{3} & \textcolor{green}{4} \end{bmatrix} \times \begin{bmatrix} \textcolor{red}{5} & \textcolor{blue}{6} \\ \textcolor{red}{7} & \textcolor{red}{8} \end{bmatrix} = \left[\textcolor{red}{5} \begin{bmatrix} \textcolor{blue}{1} \\ \textcolor{blue}{3} \end{bmatrix} + 7 \begin{bmatrix} \textcolor{green}{2} \\ \textcolor{green}{4} \end{bmatrix} \quad \textcolor{blue}{6} \begin{bmatrix} \textcolor{blue}{1} \\ \textcolor{blue}{3} \end{bmatrix} + \textcolor{red}{8} \begin{bmatrix} \textcolor{green}{2} \\ \textcolor{green}{4} \end{bmatrix} \right]$$

- Sum of outer products

$$\begin{bmatrix} \textcolor{blue}{1} & \textcolor{green}{2} \\ \textcolor{blue}{3} & \textcolor{green}{4} \end{bmatrix} \times \begin{bmatrix} \textcolor{red}{5} & \textcolor{blue}{6} \\ \textcolor{red}{7} & \textcolor{red}{8} \end{bmatrix} = \left[\begin{bmatrix} \textcolor{blue}{1} \\ \textcolor{blue}{3} \end{bmatrix} \begin{bmatrix} \textcolor{red}{5} & \textcolor{blue}{6} \end{bmatrix} + \begin{bmatrix} \textcolor{green}{2} \\ \textcolor{green}{4} \end{bmatrix} \begin{bmatrix} \textcolor{red}{7} & \textcolor{red}{8} \end{bmatrix} \right]$$