

# CS601: Software Development for Scientific Computing

Autumn 2021

Week13:

Hierarchical Methods (FMM) and Sparse Matrices

# Course Progress..

- Last week
  - Tree-based codes (hierarchical methods)
    - Barnes-Hut
    - Fast Multipole Method (FMM)
- This week
  - FMM
  - Sparse matrices and
  - PA4 discussion


# FMM Algorithm

1. Build the quadtree containing all the points.
2. Traverse the quadtree from bottom to top, computing  $\text{Outer}(n)$  for each square  $n$  in the tree.
3. Traverse the quadtree from top to bottom, computing  $\text{Inner}(n)$  for each square in the tree.
4. For each leaf, add the contributions of nearest neighbors and particles in the leaf to  $\text{Inner}(n)$

*what is  $\text{Outer}(n)$  and  $\text{Inner}(n)$  ?*

# Well Separated Regions

- Compute the influence of all particles in source region (B) on every particle in target region (A)  
(assumption: A and B are well-separated)



- At each point  $p_i$  in A, compute potential:


$$\Phi(x_i, y_i) = \sum_{p_j \in B} m_j \log |p_i - p_j|$$
$$i = 1 \text{ to } N_A, \quad j = 1 \text{ to } N_B$$

- Cost:  $O(N_A N_B)$

# Well Separated Regions

- Compute the influence of all particles in source region (B) on every particle in target region (A)

$$\Phi(x_{p_i}, y_{p_i}) = \sum_{p_j \in B} m_j \log |p_i - p_j|, p_i \in A$$



# Applying the 3-step Approximation


- In N-body simulation every point serves as source as well as target.

How to identify source and target (boxes A and B in previous slide) i.e. well-separated regions?


*Hierarchical decomposition*

# Hierarchical Decomposition

- *Level-0 decomposition*



- *Level-1 decomposition*



*No well-separated boxes*

# Hierarchical Decomposition

- *Level-2 decomposition*

N1	N2	N3	A1
N4	<b>B</b>	N5	A2
N6	N7	N8	A3
A7	A6	A5	A4



Well-separated from B

Can approximate the influence of points in B on points in  $A_i$  s

What do we do about **B**'s influence on  $N_i$  s?



# Hierarchical Decomposition

- *Level-3 decomposition*

N1	N2	N3	A1
N4	B1 B2 B3 B4	N5	A2
N6	N7	N8	A3
A7	A6	A5	A4



Influence of points in  $B_i$  s on those in  $A_i$  s  
already computed at the previous level  
(level-2)

# Hierarchical Decomposition

- Level-3 decomposition*

n1	n2	n5	n6	n9	n10			
n3	n4	n7	n8	n11	n12			A1
n13	n14	B1	B2		n27			A2
n15	n16	B3	B4		n26			
n17	n18				n25			A3
n19	n20	n21	n22	n23	n24			
								A4
A7		A6		A5				



Influence of points in  $B_i$  s on those in  $A_i$  s already computed at the previous level (level-2)



Well-separated from  $B_4$

Influence of  $B_4$ 's points on  $n_x$ 's points can be approximated

$n_x$ 's constitute the interaction list for  $B_4$ .

*What is the max size of interaction list? i.e. max number of  $n_x$  s that we can have for any  $B_i$ ?*

# Hierarchical Decomposition

- *Level-3 decomposition*

n1	n2	n5	n6	n9	n10		
n3	n4	n7	n8	n11	n12		
n13	n14	B1	B2		n27		
n15	n16	B3	<b>B4</b>		n26		
n17	n18				n25		
n19	n20	n21	n22	n23	n24		



Influence of points in  $B_i$  s on those in  $A_i$  s already computed at the previous level (level-2)




Well-separated from  $B4$



Influence of  $B4$ 's points on  $n_x$ 's points can be approximated

What do we do about **B4**'s influence on its neighbors (white/unshaded boxes)?

# Hierarchical Decomposition

- *Level-4 decomposition*



Any unshaded box outside  can be the *target* for computing the influence of points in  (*source*)

# 1. Computing Potential for Well-Separated Regions

```
1. for level L=2 to last_level
2.   for each Box B at level L
3.     iList = GetInteractionList(B)
4.     for each well-separated box A in iList
5.       //Compute potential
6.       potential =  $m_B \log |C_A - C_B|$ 
7.       //Accumulate potential
8.        $\Phi(x_{C_A}, y_{C_A}) += \text{potential}$ 
```

Cost?

# 1. Computing Potential for Well-Separated Regions

```
1. for level L=2 to last_level
2.   for each Box B at level L
3.     iList = GetInteractionList(B)
4.     for each well-separated box A in iList
5.       //Compute potential
6.       potential =  $m_B \log |C_A - C_B|$ 
       //Accumulate potential
        $\Phi(x_{C_A}, y_{C_A}) += \text{potential}$ 
```

**Prereqs:** we need  $m_B, C_A, C_B$  details. (step 0)

## 2. Assigning Potential to Points

1. **for** each Box  $A$  at level  $L=0$  to  $\text{last\_level}$
2.  $\Phi_{p_i} = \Phi_{p_i} + \Phi_{C_A}$  (where  $p_i \in A$  and  $C_A$  is  $A$ 's CM)

Cost?

### 3. Assigning Potential to Points (last level)

1. **for** each Box  $B$  at `last_level`
2.  $\Phi_{p_i} = \Phi_{p_i} + \sum_{p_j \in \text{Neighbors}(B)} m_B \log |p_i - p_j|$  (where  $p_i \in B$ )

Cost?



# 0. Computing Prereqs

1. **for** each Box B at level L=0 to last\_level
2.  $m_B = \sum_{p_j \in B} m_j$
3. //similarly compute  $C_B$

Cost?

# Total Cost (steps 0 + 1 + 2 + 3)

$$O(N \log N) + O(N) + O(N \log N) + O(N)$$

Can we do better?

# 0'. Computing Prereqs

- Traverse the tree bottom up instead of top-down  
**for** each Box  $B$  starting from `last_level` to  $L=0$   
    **if**  $B$  is a leaf box

$$m_B = \sum_{p_j \in B} m_j$$

**else**

$$m_B = m_{B_1} + m_{B_2} + m_{B_3} + m_{B_4}$$

//  $B_1$ - $B_4$  are children of  $B$

Cost?

## 2'. Assigning Potential to Points

1. **for** each Box  $A$  at level  $L=0$  to  $\text{last\_level}$

2.     **if**  $A$  is a leaf box

$$\Phi_{p_i} = \Phi_{p_i} + \Phi_{C_A} \quad (\text{where } p_i \in A \text{ and } C_A \text{ is } A\text{'s CM})$$

**else**

$$\Phi_{A_1} = \Phi_{A_1} + \Phi_A$$

$$\Phi_{A_2} = \Phi_{A_2} + \Phi_A$$

$$\Phi_{A_3} = \Phi_{A_3} + \Phi_A$$

$$\Phi_{A_4} = \Phi_{A_4} + \Phi_A$$

    // $A_1$ - $A_4$  are children of  $A$

Cost?

# Total Cost (steps 0' + 1 + 2' + 3)


$$O(N) + O(N) + O(N) + O(N)$$

**Problem:** low accuracy if source (A) and target (B) are not far away from each other

**Solution:** more accurate representations for  $m_B$  and  $\Phi(x_{C_A}, y_{C_A})$

# Multipole expansion

- Like a Taylor series expansion that is accurate when  $x^2 + y^2$  is large ( $x, y$  are cartesian coordinates of the point)
- For a quadtree box B centered at  $(x_{C_B}, y_{C_B})$ , we compute and store the terms:  $\{m_B, \alpha_1, \alpha_2, \dots, \alpha_p, z_{C_B}\}$




$$\alpha_j = \sum_{i=1}^{N_B} m_i \left( \frac{z_i^j}{j} \right)$$

$z_i$  means  $|z_i| = |(x_i, y_i)|$

# Multipole expansion

- We approximate the potential at point  $z$  due to  $B$  by:



$$\begin{aligned} \Phi(x_z, y_z) = & m_B \log(z - C_B) + \\ & \frac{\alpha_1}{z - C_B} + \\ & \frac{\alpha_2}{(z - C_B)^2} + \\ & \dots \\ & + \\ & \frac{\alpha_p}{(z - C_B)^p} \end{aligned}$$


- Because  $\{m_B, \alpha_1, \alpha_2, \dots, \alpha_p, z_{C_B}\}$  is used to compute potential outside  $B$ , it is called outer expansion

# Multipole expansion

- Similarly, we have the inner expansion  $\{m_B, \beta_1, \beta_2, \dots, \beta_p, z_{C_B}\}$  for computing the potential inside the Box due to all other points outside the box
- Computing outer expansions starts from leaf nodes and proceeds upwards in the tree.
- Computing inner expansions starts from root node and proceeds downwards in the tree.



# 3-Step Approximation (accurate)



# FMM Algorithm

1. Build the quadtree containing all the points.
2. Traverse the quadtree from bottom to top, computing  $\text{Outer}(n)$  for each square  $n$  in the tree.
3. Traverse the quadtree from top to bottom, computing  $\text{Inner}(n)$  for each square in the tree.
4. For each leaf, add the contributions of nearest neighbors and particles in the leaf to  $\text{Inner}(n)$

# Multipole expansion

- How to obtain the expression for  $\alpha$ ,  $\beta$  ?
- What is the value of  $p$ ?
- How to compute  $\alpha$  and  $\beta$ ?
- Further reading:  
<https://people.eecs.berkeley.edu/~demmel/cs267/lecture27/lecture27.html>

# Matrix Algebra and Efficient Computation

- Pic source: the Parallel Computing Laboratory at U.C. Berkeley: A Research Agenda Based on the Berkeley View (2008)

<i>Motif</i>	Embed	Desktop	Games	DB	ML	HPC	Medicine	Music	Speech	CBIR	Browser	Motif	Desktop	Games	DB	ML	HPC	Medicine	Music	Speech	CBIR	Browser	
1 Finite State Mach.												9 N-Body											
2 Combinational												10 MapReduce											
3 Graph Traversal												11 Backtrack/B&B											
4 Structured Grid												12 Graphical Models											
5 Dense Matrix												13 Unstructured Grid											
6 Sparse Matrix												<i>Temperature Chart of Need</i>				DB = database							
7 Spectral (FFT)												Hot	Warm	Med	Cool	ML = machine learning							
8 Dynamic Prog																	HPC = High Perf. Comp.						

**Figure 4. Temperature Chart of the 13 Motifs.** It shows their importance to each of the original six application areas and then how important each one is to the five compelling applications of Section 3.1. More details on the motifs can be found in (Asanovic, Bodik et al. 2006).

Seen earlier  
Next..

# Matrix Multiplication

- Why study?
  - An important “kernel” in many linear algebra algorithms
  - Most studied kernel in high performance computing
  - Simple. Optimization ideas can be applied to other kernels
- Matrix representation
  - Matrix is a 2D array of elements. Computer memory is inherently linear
  - C++ and Fortran allow for definition of 2D arrays. 2D arrays stored row-wise in C++. Stored column-wise in Fortran. E.g.  
`// stores 10 arrays of 20 doubles each in C++`  
`double** mat = new double[10][20];`

# Storage Layout - Example

- Matrix (**2D**):  $A = \begin{bmatrix} A(0,0) & A(0,1) & A(0,2) \\ A(1,0) & A(1,1) & A(1,2) \\ A(2,0) & A(2,1) & A(2,2) \end{bmatrix}$

$A(i, j) = A(\text{row}, \text{column})$  refers to the matrix element in the  $i^{\text{th}}$  row and the  $j^{\text{th}}$  column

- Row-wise (/Row-major) storage in memory:

$A(0,0)$	$A(0,1)$	$A(0,2)$	$A(1,0)$	$A(1,1)$	$A(1,2)$	$A(2,0)$	$A(2,1)$	$A(2,2)$
----------	----------	----------	----------	----------	----------	----------	----------	----------

- Column-wise (/Column-major) storage in memory:

$A(0,0)$	$A(1,0)$	$A(2,0)$	$A(0,1)$	$A(1,1)$	$A(2,1)$	$A(0,2)$	$A(1,2)$	$A(2,2)$
----------	----------	----------	----------	----------	----------	----------	----------	----------

- Generalizing data storage order for ND:** last index changes fastest in row-major. Last index changes slowest in col-major.

# Storage Layout - Exercise

- For a 3D array (tensor) assume  $A(i, j, k) = A(\text{row}, \text{column}, \text{depth})$



- What is the offset of  $A(1, 2, 1)$  ? as per row-major storage?
- What is the offset of  $A(1, 2, 1)$  ? as per col-major storage?

# Storage Layout

- Layout format itself doesn't influence efficiency (i.e. no general answer to “is column-wise layout better than row-wise?” )
- However, knowing the layout format is critical for good performance
  - *Always traverse the data in the order in which it is laid out*

How good performance?



Run on (12 X 2592.01 MHz CPU s)

CPU Caches:

L1 Data 32 KiB (x6)

L1 Instruction 32 KiB (x6)

L2 Unified 256 KiB (x6)

L3 Unified 12288 KiB (x1)

Load Average: 0.07, 0.02, 0.07


Source code: <https://github.com/eliben/code-for-blog/tree/master/2015/benchmark-row-col-major>

Benchmark	Time	CPU	Iterations	UserCounters...
BM_AddByRow/64/64	693 ns	693 ns	1042737	items_per_second=5.91004G/s
BM_AddByRow/128/128	2464 ns	2464 ns	271766	items_per_second=6.64813G/s
BM_AddByRow/256/256	11134 ns	11133 ns	63210	items_per_second=5.88639G/s
BM_AddByRow/512/512	44353 ns	44353 ns	15576	items_per_second=5.91041G/s
BM_AddByCol/64/64	3270 ns	3270 ns	212929	items_per_second=1.25254G/s
BM_AddByCol/128/128	39741 ns	39741 ns	17617	items_per_second=412.272M/s
BM_AddByCol/256/256	314880 ns	314878 ns	2241	items_per_second=208.132M/s
BM_AddByCol/512/512	1276733 ns	1276723 ns	545	items_per_second=205.326M/s

```
des/week13_codesamples$ ./a.out 4096
Rowwise time n=4096 (us): 18967
Colwise time n=4096 (us): 158608
nikhilh@ndhpc01:/mnt/c/temp/Nikhil/Cou
des/week13_codesamples$ ./a.out 2048
Rowwise time n=2048 (us): 4860
Colwise time n=2048 (us): 32158
nikhilh@ndhpc01:/mnt/c/temp/Nikhil/Cou
des/week13_codesamples$ ./a.out 1024
Rowwise time n=1024 (us): 1125
Colwise time n=1024 (us): 1980
```



Matrix-Matrix Addition benchmarking  
([Source code and further reading](#) )



Matvec execution time  
(we used the [source code](#) as an  
inconclusive example for benchmarking)

# Linear Algebra Software

- Use optimized kernels from libraries whenever possible.
- E.g. BLAS, LAPACK, SuperLU, Trilinos, OpenBLAS etc.