

# CS601: Software Development for Scientific Computing

Autumn 2021

Week14:  
Matrix Algebra

# Course Progress..

- Last week: FMM, PA4, Matrix Algebra
  - FMM ideas - applying 3-step approximation (decomposition), optimizing (reuse computation), better approximation (multipole expansion), Cost.
  - PA4 discussion
  - Matrix algebra
    - Overview: matrix-matrix multiplication (motivation), program representation of a matrix, storage layout and performance implications.
- This week: Matrix algebra contd.

# Matrix Multiplication

- Three fundamental ways to think of the computation

1. Dot product

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \times \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix} = \begin{bmatrix} 1.5 + 2.7 & 1.6 + 2.8 \\ 3.5 + 4.7 & 3.6 + 4.8 \end{bmatrix}$$

2. Linear combination of the columns of the left matrix

$$\begin{bmatrix} \textcolor{blue}{1} & \textcolor{green}{2} \\ \textcolor{blue}{3} & \textcolor{green}{4} \end{bmatrix} \times \begin{bmatrix} \textcolor{red}{5} & \textcolor{blue}{6} \\ \textcolor{red}{7} & \textcolor{red}{8} \end{bmatrix} = \left[ \textcolor{red}{5} \begin{bmatrix} \textcolor{blue}{1} \\ \textcolor{blue}{3} \end{bmatrix} + 7 \begin{bmatrix} \textcolor{green}{2} \\ \textcolor{green}{4} \end{bmatrix} \quad \textcolor{blue}{6} \begin{bmatrix} \textcolor{blue}{1} \\ \textcolor{blue}{3} \end{bmatrix} + \textcolor{red}{8} \begin{bmatrix} \textcolor{green}{2} \\ \textcolor{green}{4} \end{bmatrix} \right]$$

3. Sum of outer products

$$\begin{bmatrix} \textcolor{blue}{1} & \textcolor{green}{2} \\ \textcolor{blue}{3} & \textcolor{green}{4} \end{bmatrix} \times \begin{bmatrix} \textcolor{red}{5} & \textcolor{blue}{6} \\ \textcolor{red}{7} & \textcolor{red}{8} \end{bmatrix} = \left[ \begin{bmatrix} \textcolor{blue}{1} \\ \textcolor{blue}{3} \end{bmatrix} \begin{bmatrix} 5 & 6 \end{bmatrix} + \begin{bmatrix} \textcolor{green}{2} \\ \textcolor{green}{4} \end{bmatrix} \begin{bmatrix} 7 & 8 \end{bmatrix} \right]$$

# Dot Product

- Vector  $x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$ , Vector  $y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$   $x_i, y_i \in \mathbb{R}$
- $x^T = [x_1 \quad x_2 \quad \dots \quad x_n]$
- Dot Product or Inner Product:  $c = x^T y$   $x^T \in \mathbb{R}^{1 \times n}, y \in \mathbb{R}^{n \times 1}, c$  is scalar

$$[x_1 \quad x_2 \quad \dots \quad x_n] \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = [x_1 y_1 + x_2 y_2 + \dots + x_n y_n]$$

- E.g.  $[1 \quad 2 \quad 3] \begin{bmatrix} 4 \\ 5 \\ 6 \end{bmatrix} = [1 \times 4 + 2 \times 5 + 3 \times 6] = 32$

# AXPY

- Computing the more common (a times x plus y):  $y = y + ax$

- $$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} y \\ y_2 \\ \vdots \\ y_n \end{bmatrix} + a \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$


```
..  
for i=1 to n  
    c[i] = c[i] + x[i]*y[i]  
..
```

- Cost? n multiplications and n additions = **2n** or **O(n)**

# Matrix Vector Product

- Computing Matrix-Vector product:  $c = c + Ax$ ,  $A \in \mathbb{R}^{m \times r}$ ,  $x \in \mathbb{R}^{r \times 1}$

$$\begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} + \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1r} \\ a_{21} & a_{22} & \cdots & a_{2r} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mr} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_r \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} + \begin{bmatrix} a_{11}x_1 + a_{12}x_2 + \cdots + a_{1r}x_r \\ a_{21}x_1 + a_{22}x_2 + \cdots + a_{2r}x_r \\ \vdots \\ a_{m1}x_1 + a_{m2}x_2 + \cdots + a_{mr}x_r \end{bmatrix}$$



- Rewriting Matrix-Vector product using dot products:

$$\begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} + \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1r} \\ a_{21} & a_{22} & \cdots & a_{2r} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mr} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_r \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} + \begin{bmatrix} a_1^T x \\ a_2^T x \\ \vdots \\ a_m^T x \end{bmatrix}$$

- Cost?  $m$  rows involving dot products and having the form  $c_i = c_i + x^T y$  (Per row cost =  $2r$  (because  $a_i, x \in \mathbb{R}^r$ ), Total cost =  $2mr$  or  $O(mr)$ )

# Matrix-Matrix Product

- Computing Matrix-Matrix product  $C = C + AB$ ,  $A \in \mathbb{R}^{m \times r}$ ,  $B \in \mathbb{R}^{r \times n}$ ,  $C \in \mathbb{R}^{m \times n}$

$$\begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1n} \\ c_{21} & c_{22} & \cdots & c_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mn} \end{bmatrix} = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1n} \\ c_{21} & c_{22} & \cdots & c_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mn} \end{bmatrix} + \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1r} \\ a_{21} & a_{22} & \cdots & a_{2r} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mr} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{r1} & b_{r2} & \cdots & b_{rn} \end{bmatrix}$$

- Consider the  $AB$  part first.

$$\begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1r} \\ a_{21} & a_{22} & \cdots & a_{2r} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mr} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{r1} & b_{r2} & \cdots & b_{rn} \end{bmatrix}$$

# Matrix-Matrix Product

$$\begin{array}{c} \text{A} \end{array} \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1r} \\ a_{21} & a_{22} & \dots & a_{2r} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mr} \end{bmatrix} \begin{array}{c} \text{B} \end{array} \begin{bmatrix} b_{11} & b_{12} & \dots & b_{1n} \\ b_{21} & b_{22} & \dots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{r1} & b_{r2} & \dots & b_{rn} \end{bmatrix}$$

$$= \begin{bmatrix} a_{11}b_{11} + a_{12}b_{21} + \dots + a_{1r}b_{r1} & \dots & a_{11}b_{1n} + a_{12}b_{2n} + \dots + a_{1r}b_{rn} \\ \vdots & \ddots & \vdots \\ a_{m1}b_{11} + a_{m2}b_{21} + \dots + a_{mr}b_{r1} & \dots & a_{m1}b_{1n} + a_{m2}b_{2n} + \dots + a_{mr}b_{rn} \end{bmatrix}$$

$$= \begin{bmatrix} a_1^T b_1 & \dots & a_1^T b_n \\ \vdots & \ddots & \vdots \\ a_m^T b_1 & \dots & a_m^T b_n \end{bmatrix}$$

$a_i^T \in \mathbb{R}^{1 \times r}, b_j \in \mathbb{R}^{r \times 1}$   
 $i$  ranges from 1 to  $m$   
 $j$  ranges from 1 to  $n$



# Matrix-Matrix Product using Dot Product Formulation

- Pseudocode - Matrix-Matrix product:  $C = C + AB$ ,  $A \in \mathbb{R}^{m \times r}$ ,  $B \in \mathbb{R}^{r \times n}$ ,  $C \in \mathbb{R}^{m \times n}$

..

for i=1 to m

for j=1 to n

//compute updates involving dot products

$$c_{ij} = c_{ij} + a_i^T b_j$$


- Expanded: ..

for i=1 to m

for j=1 to n

for k=1 to r

$$c_{ij} = c_{ij} + a_{ik} b_{kj}$$



Elements of C matrix are computed from top to bottom, left to right. Per element computation, you need a row of A and a column of B.

# Matrix-Matrix Product using Dot Product Formulation

- Pseudocode - Matrix-Matrix product:  $C = C + AB$ ,  $A \in \mathbb{R}^{m \times r}$ ,  $B \in \mathbb{R}^{r \times n}$ ,  $C \in \mathbb{R}^{m \times n}$   
..  
for i=1 to m  
for j=1 to n  
//compute updates involving dot products  
 $c_{ij} = c_{ij} + a_i^T b_j$
- Cost?
  - Per dot-product cost =  $2r$  ( $a_i, b_j \in \mathbb{R}^r$ ) Total cost =  $2mnr$  or  $O(mnr)$

# Common Computational Patterns

Some patterns that we see while doing Matrix-Matrix product:

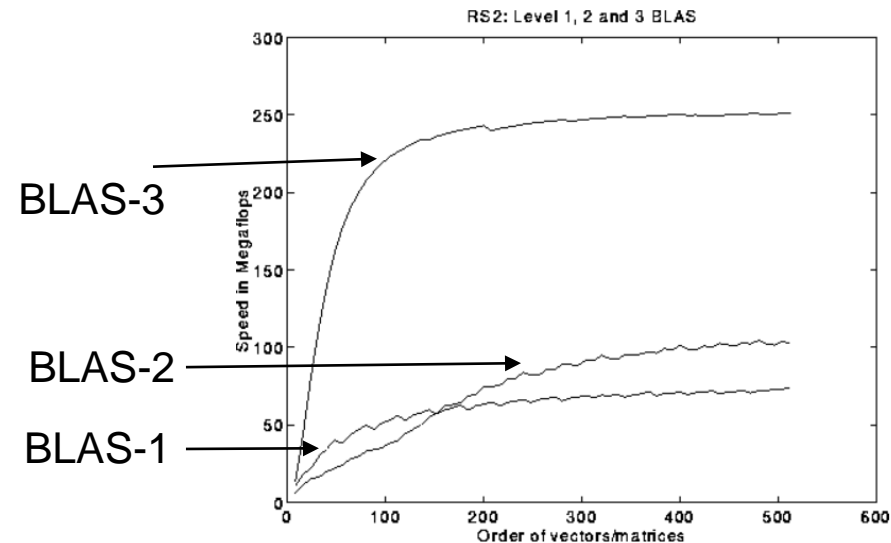
- Dot Product or Inner Product:  $x^T y$  ← Slide 4, Method 1
- Scalar **a** times **x** plus **y**:  $y = y + ax$  OR  $axpy$  ← Slide 4, Method 2
- Scalar times **x**:  $\alpha x$
- **Matrix** times **x** plus **y**:  $y = y + Ax$  ← Slide 4, Method 1
  - generalized  $axpy$  OR  $gaxpy$
- Outer product:  $C = C + xy^T$  ← Slide 4, Method 3
- **Matrix** times **Matrix** plus **Matrix**
  - GEMM or generalized matrix multiplication

# BLAS – Basic Linear Algebra Subroutines

- Level-1 or BLAS-1 (46 operations, routines operating on vectors mostly)
  - axpy, dot product, rotation, scale, etc.
  - 4 versions each: **Single-precision**, **double-precision**, **complex**, **complex-double (z)**
  - E.g. saxpy, daxpy, caxpy etc.
  - **Do  $O(n)$  operations on  $O(n)$  data.**
- Level-2 or BLAS-2 (25 operations, routines operating on matrix-vectors mostly)
  - E.g. GEMV ( $\alpha A \cdot x + \beta y$ ), GER (Rank-1 update  $A = A + y \cdot x^T$ ), Triangular solve ( $y = T \cdot x, T$  is a triangular matrix) etc.
  - 4 versions each, **do  $O(n^2)$  operations on  $O(n^2)$  data.**

# BLAS – Basic Linear Algebra Subroutines

- Level-3 or BLAS-3 (9 basic operations, routines operating on matrix-matrix mostly)
  - GEMM ( $C = \alpha A \cdot B + \beta C$ ),
  - Multiple triangular solve ( $Y = TX$ ,  $T$  is triangular,  $X$  is rectangular)
  - **Do  $O(n^3)$  operations on  $O(n^2)$  data.**
- *Why categorize as BLAS-1, BLAS-2, BLAS-3?*
  - *Performance*



source: <http://people.eecs.berkeley.edu/~demmel/cs267/lecture02.html>

# Computational Intensity

- Average number of operations performed per data element (word) read/written from slow memory
  - E.g. Read/written  $m$  words from memory. Perform  $f$  operations on  $m$  words.
  - Computational Intensity  $q = f/m$  (*flops per word*).
- We want to *maximize* the computational intensity
- What is  $q$  for axpy? Matrix-vector product? Matrix-Matrix product?

# Computational Intensity - axpy

$$\begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_n \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_n \end{bmatrix} + [x_1 \quad x_2 \quad \cdots \quad x_n] \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_n \end{bmatrix} + \begin{bmatrix} x_1 \times y_1 \\ x_2 \times y_2 \\ \vdots \\ x_n \times y_n \end{bmatrix}$$

Read(x) //read x from slow memory

Read(y) //read y from slow memory

Read(c) //read c from slow memory

for i=1 to n

    c[i] = c[i] + x[i]\*y[i] //do arithmetic on data read

Write(c) //write c back to slow memory

- Number of memory operations = 4n (assuming one word of storage for each component  $(x_i, y_i, c_i)$  of vectors x, y, c resp.)
- Number of arithmetic operations = 2n (one addition and one multiplication per row.)
- **$q=2n/4n = 1/2$**

# Computational Intensity – matrix-vector

- Assume  $m=r=n =n$

$$\begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} + \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1r} \\ a_{21} & a_{22} & \cdots & a_{2r} \\ & \vdots & & \\ a_{m1} & a_{m2} & \cdots & a_{mr} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_r \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_m \end{bmatrix} + \begin{bmatrix} a_{11}x_1 + a_{12}x_2 + \cdots + a_{1r}x_r \\ a_{21}x_1 + a_{22}x_2 + \cdots + a_{2r}x_r \\ \vdots \\ a_{m1}x_1 + a_{m2}x_2 + \cdots + a_{mr}x_r \end{bmatrix}$$

- Number of memory operations =  $n^2 + 3n = n^2 + O(n)$
- Number of arithmetic operations =  $2n^2$
- $q \approx 2n^2/n^2 = 2$



# Computational Intensity – matrix-matrix

```
for i=1 to n
  //Read row i of A into fast memory →  $n^2$  words read: each row of A read
  for j=1 to n                                     once for each i. Assume that the row
  //Read C(i,j) into fast memory                    read stays in fast memory during the
  //Read column j of B into fast memory →  $n^3$  words read: each column      execution of inner two loops.
  for k=1 to n                                     of B read  $n^2$  times
    C(i,j)=C(i,j) + A(i,k)*B(k,j)
  //Write C(i,j) back to slow memory →  $2n^2$  words read: read/write each
  entry of C to memory once.
```

- Number of memory operations =  $n^3 + 3n^2 = n^3 + O(n^2)$
- Number of arithmetic operations =  $2n^3$
- $q \approx 2n^3/n^3 = 2$ . Same as matrix-vector?
- What if the fast memory has space to hold entire B matrix, a row of A matrix, and one element of C matrix?

# Blocked Matrix Multiply

- For  $N=4$ :

$$\begin{bmatrix} C1 & C2 & C3 & C4 \end{bmatrix} = \begin{bmatrix} C1 & C2 & C3 & C4 \end{bmatrix} + \begin{bmatrix} A \end{bmatrix} * \begin{bmatrix} B1 & B2 & B3 & B4 \end{bmatrix}$$

$$\begin{bmatrix} Cj \end{bmatrix} = \begin{bmatrix} Cj \end{bmatrix} + \begin{bmatrix} A \end{bmatrix} * \begin{bmatrix} Bj \end{bmatrix} = \begin{bmatrix} Cj \end{bmatrix} + \sum_{k=1}^n \begin{bmatrix} A(:,k) \end{bmatrix} * \begin{bmatrix} Bj(k,:) \end{bmatrix}$$

```

for j=1 to N
//Read column j of B into fast memory
//Read column j of C into fast memory
  for k=1 to n
    //Read column k of A into fast memory
    C(*,j)=C(*,j) + A(*,k)*Bj(k,*) //outer-product
  //Write C(i,j) back to slow memory
  
```

# Computational Intensity - Blocked Matrix Multiply

```
for j=1 to N
//Read column j of B into fast memory →  $n^2$  words read: each column of B read once.
//Read column j of C into fast memory
  for k=1 to n
    //Read column k of A into fast memory →  $Nn^2$  words read: each column of A read N times
    C(*,j)=C(*,j) + A(*,k)*Bj(k,*) //outer-product
    //Write C(i,j) back to slow memory
```

- Number of arithmetic operations =  $2n^3$  →  $2n^2$  words read: read/write each entry of C to memory once.
- $q = 2n^3 / (N + 3)n^2 = 2n/N$ . **Good!**