

ECE264: Advanced C Programming

Summer 2019

Week 1: Tools, Program Layout, Data Types and Structs



Git

- Version Control System
 - Manage versions of your code – access to different versions when needed
 - Lets you collaborate
- ‘Repository’ – virtual storage
 - Local and Remote Repository
 - Local is working copy

Git – Initializing Repositories

- Getting started with local working copies:
 - git init

Terminal

```
[ecegrid-thin4:~/ECE264/dem0] hegden$ls -a
.  ..
[ecegrid-thin4:~/ECE264/dem0] hegden$git init
Initialized empty Git repository in /home/min/a/hegden/ECE264/dem0/.git/
[ecegrid-thin4:~/ECE264/dem0] hegden$ls -a
.  .. .git
[ecegrid-thin4:~/ECE264/dem0] hegden$
```

- git clone (when a remote repository on github.com exists)

Terminal

```
[ecegrid-thin4:~/ECE264] hegden$ls -a
.  ..
[ecegrid-thin4:~/ECE264] hegden$git clone git@github.com:ece264summer2019/dem0.git
Cloning into 'dem0'...
warning: You appear to have cloned an empty repository.
[ecegrid-thin4:~/ECE264] hegden$ls -a
.  .. dem0
[ecegrid-thin4:~/ECE264] hegden$cd dem0/
[ecegrid-thin4:~/ECE264/dem0] hegden$ls -a
.  .. .git
[ecegrid-thin4:~/ECE264/dem0] hegden$
```

Git – Adding Content

- Staging

Terminal

```
[ecegrid-thin4:~/ECE264/dem0] hegden$echo "This repository is created for demo purposes." > README.txt  
[ecegrid-thin4:~/ECE264/dem0] hegden$git add README.txt
```

- Commit (save changes in local repository)

```
[ecegrid-thin4:~/ECE264/dem0] hegden$git commit -m "My first commit"  
[master ab680c6] My first commit  
1 file changed, 1 insertion(+)  
create mode 100644 README.txt
```

- Save changes in remote repository (guard against accidental deletes)

```
[ecegrid-thin4:~/ECE264/dem0] hegden$git push  
Enumerating objects: 4, done.  
Counting objects: 100% (4/4), done.  
Writing objects: 100% (3/3), 290 bytes | 145.00 KiB/s, done.  
Total 3 (delta 0), reused 0 (delta 0)  
To github.com:ece264summer2019/dem0.git  
3dccc4f..ab680c6 master -> master
```



Git – Releasing Code

- Tagging

- make sure there are no unsaved changes in local repository

```
[ecegrid-thin4:~/ECE264/dem0] hegden$git status .  
On branch master  
Your branch is up to date with 'origin/master'.  
  
nothing to commit, working tree clean
```

```
[ecegrid-thin4:~/ECE264/dem0] hegden$git tag -a RELEASE_V0.1 -m "First release"
```

- Save tags in remote repository

```
[ecegrid-thin4:~/ECE264/dem0] hegden$git push --tags  
Enumerating objects: 1, done.  
Counting objects: 100% (1/1), done.  
Writing objects: 100% (1/1), 176 bytes | 176.00 KiB/s, done.  
Total 1 (delta 0), reused 0 (delta 0)  
To github.com:ece264summer2019/dem0.git  
* [new tag]          RELEASE_V0.1 -> RELEASE_V0.1
```

Git – Recap..

- Please read <https://git-scm.com/book/en/v2> for details

1. `git clone` (creating a local working copy)
2. `git add` (staging the modified local copy)
3. `git commit` (saving local working copy)
4. `git push` (saving to remote repository)
5. `git tag` (Naming the release with a label)
6. `git push --tags` (saving the label to remote)

Makefile

- Is a file, contains instructions for the 'make' program to generate a target (executable).
- Generating a target involves:
 1. Preprocessing (e.g. strips comments, conditional compilation etc.)
 2. Compiling (.c -> .s files, .s -> .o files)
 3. Linking (e.g. making printf available)
- A Makefile typically contains directives on how to do steps 1, 2, and 3.



Makefile - Format

- Contains series of 'rules'-

```
target: dependencies  
[TAB] system command(s)
```

Note that it is important that there be a TAB character before the command (not spaces).

Example,

```
testgen: testgen.c  
        gcc testgen.c -o testgen
```

- And Macro/Variable definitions -

```
CFLAGS = -std=c99 -g -Wall -Wshadow --pedantic -Wvla -Werror  
GCC = gcc
```


Makefile - Usage

- The 'make' command (Assumes that a file by name 'makefile' or 'Makefile'. exists)

```
[ecegrid-thin4:~/ECE264/dem0] hegden$cat makefile
GCC=gcc
CFLAGS=-std=c99 -g -Wall -Wshadow --pedantic -Wvla -Werror
testgen: testgen.c
    $(GCC) $(CFLAGS) testgen.c -o testgen
clean:
    rm testgen
[ecegrid-thin4:~/ECE264/dem0] hegden$make
gcc -std=c99 -g -Wall -Wshadow --pedantic -Wvla -Werror testgen.c -o testgen
[ecegrid-thin4:~/ECE264/dem0] hegden$
```


- To know more, please read:
https://www.gnu.org/software/make/manual/html_node/index.html#Top

Sorting - Selection sort


- Repeatedly find the minimum element in the unsorted array and put it at the beginning.
 - Divides the input array into 2 pieces - **sorted** and **rest** - such that *all elements* in **sorted** are smaller than *any element* in the **rest**
 - Works by growing **sorted** and shrinking **rest**

Selection sort - example


- A cursor dividing sorted and rest




Selection sort - example



Selection sort - example



Selection sort - example



Selection sort - example

-1	0	2	3	4	9
----	---	---	---	---	---



sorted

Selection sort - pseudocode

```
1 int input[N] = //input
2 int cursor = 0 //initial position of the cursor
3 for(cursor = 0; cursor < N; cursor++)
4     //sorted list from [0,cursor)
5     //rest of the list from [cursor, N)
6     for(i = cursor; i < N; i++)
7         //search the rest of the list to find the smallest value
8         //swap the smallest value with the value at input[cursor]
```


Selection sort - Analysis

```
1 int input[N] = //input
2 int cursor = 0 //initial position of the cursor
3 for(cursor = 0; cursor < N; cursor++)
4     //sorted list from [0,cursor)
5     //rest of the list from [cursor, N)
6     for(i = cursor; i < N; i++)
7         //search the rest of the list to find the smallest value
8         //swap the smallest value with the value at input[cursor]
```

- Outer loop (line 3) is moving the cursor, inner loop (line 6) is finding minimum.

How many times does inner loop execute?

Selection sort - Analysis

```
1 int input[N] = //input
2 int cursor = 0 //initial position of the cursor
3 for(cursor = 0; cursor < N; cursor++)
4     //sorted list from [0,cursor)
5     //rest of the list from [cursor, N)
6     for(i = cursor; i < N; i++)
7         //search the rest of the list to find the smallest value
8         //swap the smallest value with the value at input[cursor]
```

- inner loop runs N times, (N - cursor) iterations every time.

$$\begin{aligned} &= \sum_{i=0}^{N-1} N - i \\ &= \sum_{i=1}^N i = \frac{N(N+1)}{2} \end{aligned}$$



Selection sort - Analysis

- outer loop runs for N iterations
- inner loop runs for $\sim N(N+1)/2$ iterations
 - inner loop dominates

1. Approximately how many array write operations occur?

2. Double the input, how long does Selection sort take?

Program Layout in Memory

- A program's memory space is divided into four segments:
 1. Text
 - source code of the program
 2. Data
 - Broken into uninitialized and initialized segments; contains space for global and static variables. E.g. `int x = 7; int y;`
 3. Heap
 - Memory allocated using `malloc/calloc/realloc`
 4. Stack
 - Function arguments, return values, local variables, [special registers](#).

Detour - Stacks

Real Stack



Hardware Stack




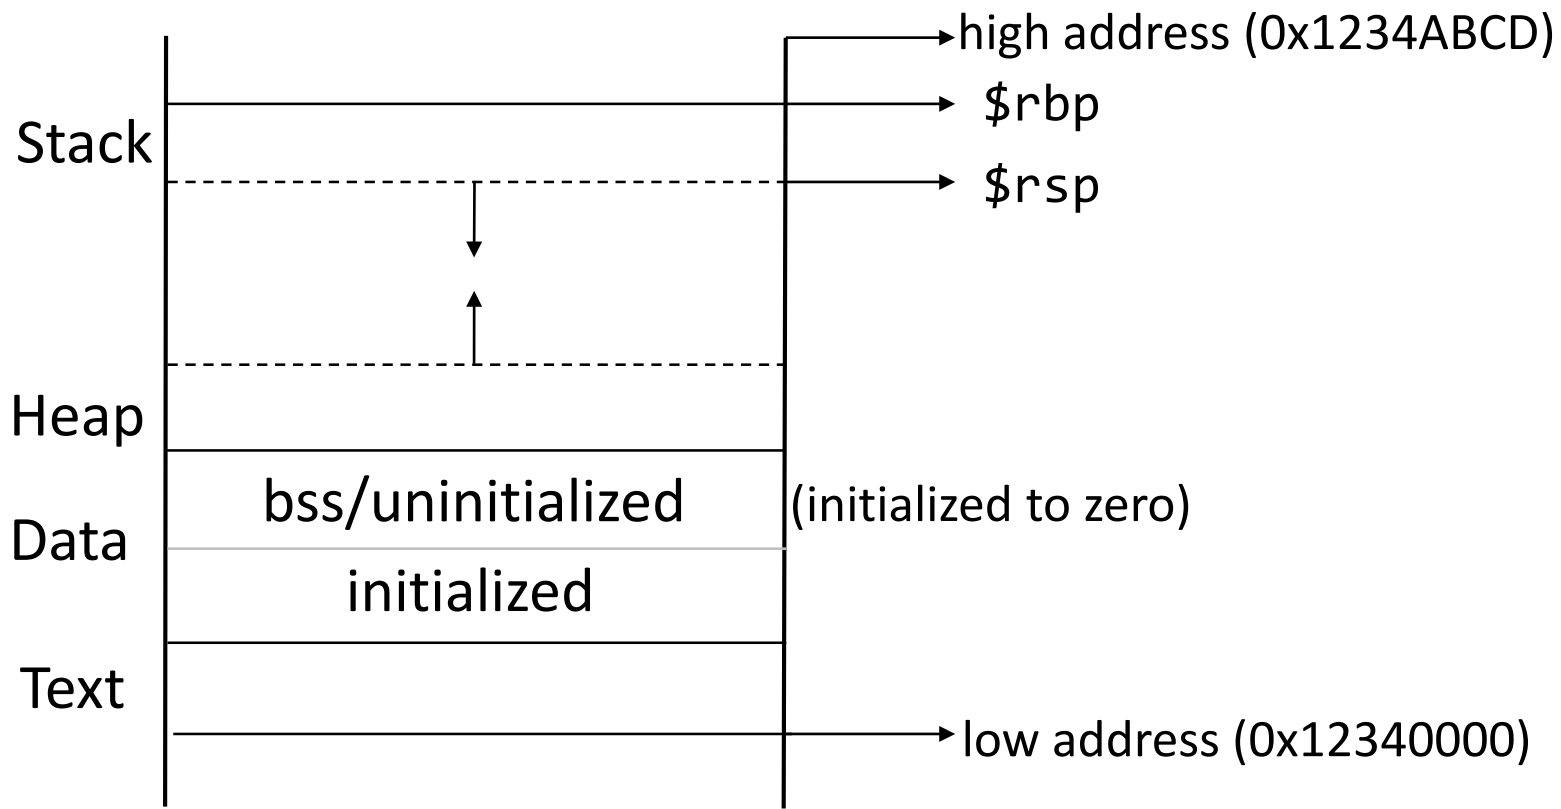
Image source: <https://eli.thegreenplace.net/2011/02/04/where-the-top-of-the-stack-is-on-x86/>



Stack Frame

- A sub-segment of memory on the stack space
 - Special registers `$rbp` and `$rsp` track the bottom and top of the stack frame.
 - Example: when `main` calls function `foo`
 1. The following are pushed on to stack:
 - `foo`'s arguments
 - Space for `foo`'s return value
 - Address of the next instruction executed (in `main`) when `foo` returns
 - Current value of `$rbp`
 2. `$rsp` is automatically updated (decremented) to point to current top of the stack.
 3. `$rbp` is assigned the value of `$rsp`

Program Layout in Memory





Question ?

Where are the command-line arguments stored?

Number Bases

- We use base-10 (decimal), Computers use base-2 (Binary).
- Binary is difficult to read. So, we use base-16 (Hexadecimal).
- Hexadecimal uses 16 digits: 0 to 9 and A to F. A to F represent decimal numbers 10 to 15.
- A digit in hexadecimal needs 4 bits. Therefore, a byte of information (8 bits) represents two digits.
- Example:

Decimal	Binary	Hexadecimal
10	1010	0xA
16	1 0000	0x10
43981	1010 1011 1100 1101	0xABCD



How are Numbers Stored in Memory? - Endianness

- Assume an integer needs 4 bytes of storage
 - E.g. 1193 in Hexadecimal = 0x4A9 = 0x 00 00 04 A9 when stored in 4 bytes of memory.
 - How are those 4 bytes ordered in memory? – Endianness
- Two popular formats: Big-Endian and Little-Endian

Big-Endian

- Most-significant-byte (MSB) at low-address and least-significant-byte (LSB) at high-address
 - E.g. 1193 = 0x00 00 04 A9 ($= 4 * 16^2 + A * 16 + 9$)
 - MSB (0x00) is written at lower address, LSB (0xA9) is written at higher address.

0000 0000	0000 0000	0000 0100	1010 1001
Address: 0x00000001	0x00000002	0x00000003	0x00000004

- Motorola 68000 Series, IBM-Z Mainframes.

Little-Endian

- Most-significant-byte (MSB) at high-address and least-significant-byte (LSB) at low-address
 - E.g. 1193 = 0x00 00 04 A9 ($= 4 * 16^2 + A * 16 + 9$)
 - MSB (0x00) is written at higher address, LSB (0xA9) is written at lower address.

1010 1001	0000 0100	0000 0000	0000 0000
Address: 0x00000001	0x00000002	0x00000003	0x00000004

- Intel x86 Architecture

Endianness

- Fortunately, we don't have to worry about endianness.
 - You don't have to reverse bytes when you read an integer.
 - Compiler and the processor do the job for you.
- However, you need to be aware of endianness only when inspecting memory contents.
 - E.g. when using GDB while debugging.