

ECE264: Advanced C Programming

Summer 2019

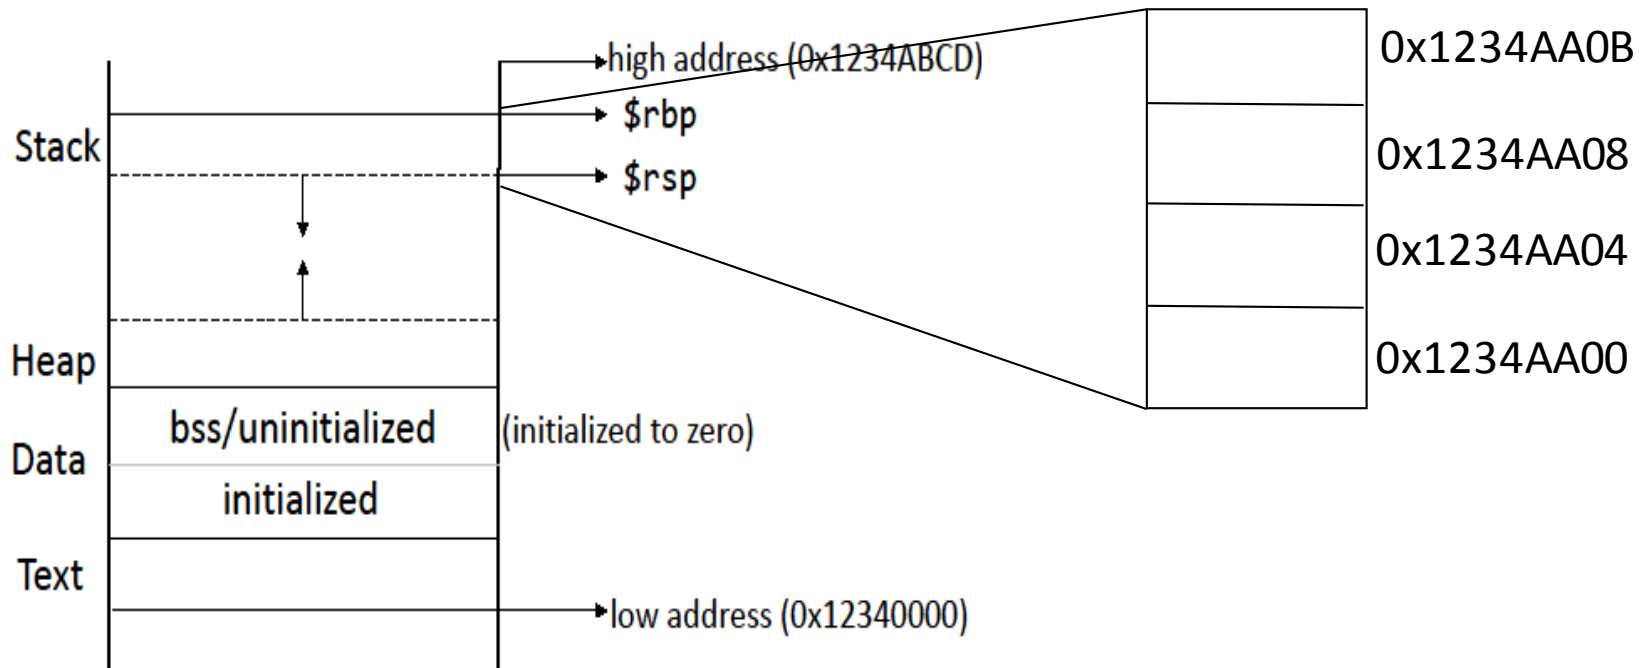
Week 2: Addresses, Pointers, Pointer Arithmetic

Addresses

- Humans are not good at remembering numerical addresses.
 - What are the GPS coordinates (latitude and longitude) of your residence?
- Addresses in computer programs are just numbers.

- Addresses in computer programs identify memory locations.
- Computer programs think and live in terms of memory locations.

Program Memory Layout - Revisited



- Every memory location is a box holding data
- Each box has an address

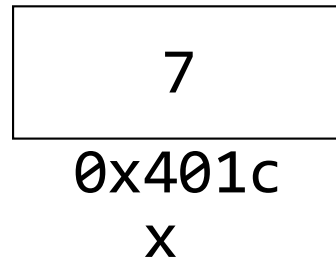
Addresses Contd..

- A program navigates by visiting one address after another.
- We (humans) choose convenient ways to identify addresses so that we can give directions to a program
 - Variables

Handles to Addresses

- What is a variable?
 - Its just a handle to an address / program memory location

- `int x = 7;`



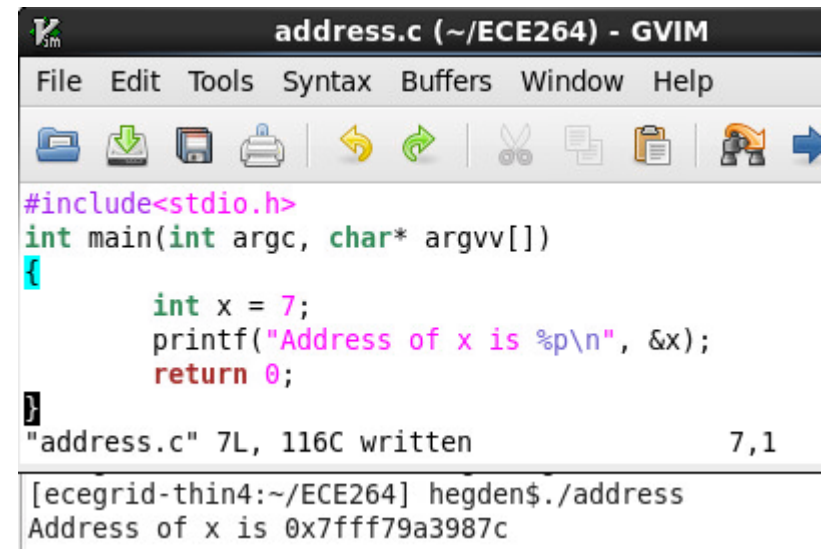
- Read x => Read the content at address 0x401C
- Write x=> Write at address 0x401C

Visualizing Addresses

- The *address of* (&) operator fetches a variable's address in C.
- &x would return the address 0x401C.
- Format specifier 'p' :

```
printf("%p\n",&x)
```

prints the Hexadecimal
address of x



The screenshot shows a Gvim window titled "address.c (~/ECE264) - Gvim". The menu bar includes File, Edit, Tools, Syntax, Buffers, Window, and Help. Below the menu is a toolbar with icons for file operations and editing. The code in the editor is as follows:

```
#include<stdio.h>
int main(int argc, char* argv[])
{
    int x = 7;
    printf("Address of x is %p\n", &x);
    return 0;
}
```

At the bottom of the window, a status line indicates "address.c" 7L, 116C written. Below the editor, a terminal window shows the command and output:

```
[ecegrid-thin4:~/ECE264] hegden$ ./address
Address of x is 0x7fff79a3987c
```

Pointers

- Pointer is a data type that *holds an address*.

`<type>* <pointer_name>;`

We read it as “**pointer to** <type>”

- Example:

- `int* p;`

is a variable named p whose type is pointer to int
OR p is an integer pointer

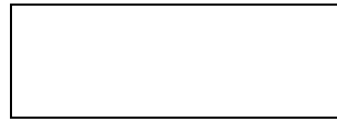
Note that the variable declared is p, *not* *p

- A pointer always stores an address
- `<type>` of the pointer tells us what kind of data is stored at that address
- Example:
 - `int* p;`

declares a pointer variable `p` holding an address, which identifies a memory location capable of storing an integer.

- `int* p;`

Remember `p` is a variable and all variables are just names identifying addresses.

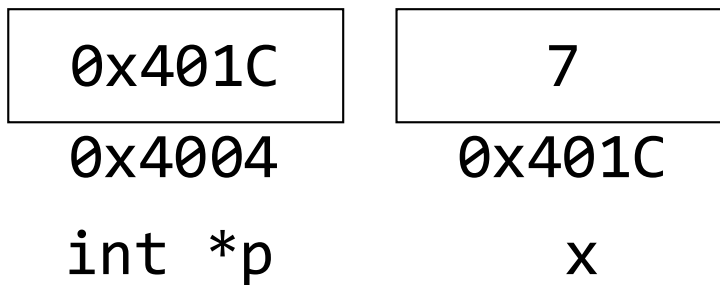


`0x4004`
`int *p`

Initializing Pointers

- `int* p=&x;`

//p holds the address of a memory location that stores an integer.



- We say *p points to x*

- Cannot assign arbitrary addresses to pointers.
- Example:
`int* p=5;`
- Operating system determines addresses available to each program.

The NULL address

- NULL is a special address

- Example

```
int* p=NULL; //p points to nowhere
```

- Useful when it is not yet known where p points to.
- Uninitialized pointers store garbage addresses

Using Pointers

- The *dereference* operator (*)
 - Lets us access the memory location at the address stored in the pointer

```
int x=7;  
int *p = &x; //p now points to x  
*p = 10; //this is the same as x=10  
int y=*p; //this is the same as y=x
```

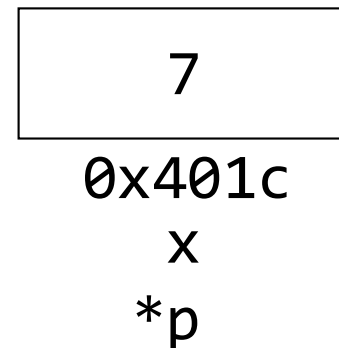
The expression `*p` is equivalent to `x`

- Pointers as alternate names to memory locations

```
int x=7;  
int *p = &x; //p now points to x  
*p = 10; //this is the same as x=10  
int y=*p; //this is the same as y=x
```

The expression `*p` is equivalent to `x`

`x` is the name for an address
`*p` is the name for an address



- Pointers as “dynamic” names to memory locations

```
int x=7;
```



0x401c

x

//x always names the location 0x401C

```
int *p = &x; // *p is now another name for x
```

```
int y = *p //like saying y=x
```

```
p = &y; // *p is now another name for y
```

```
*p=8; //like saying y=8
```


The swap function

```
int a = 8;  
int b = 10;
```

```
void swap(int x, int y) {  
    int tmp = x;  
    x = y;  
    y = tmp;  
}
```

```
void main() {  
    swap(a, b); //a is still 8, b is still 10  
}
```

Pass by value

- C functions operate on ***copies*** of arguments.
- Change the data inside the function, you change the copy. Not the original.
- In swap, x and y are names of memory locations that are copies of a and b

What if x and y held addresses of a and b?

- *x and *y would name the same memory locations that a and b did.

The swap function

```
int a = 8;  
int b = 10;
```

```
void swap(int* x, int* y) {  
    int tmp = *x; //tmp = whatever is in the  
    location that x points to.  
    *x = *y;  
    *y = tmp;  
}
```

```
void main() {  
    //remember, we have to pass addresses now,  
    not ints.  
    swap(&a, &b); //a is now 10, b is 8  
}
```

Pointers to Different Types

- What can pointers point to? any data type!
 - Basic data types,
 - Structures,
 - Functions, and
 - even Pointers!

Pointer Chains

```
int x = 7;  
int *p = x; //p points to x; *p is same  
as x.
```

```
int * * q; //q is a pointer to pointer  
to int
```

*q is same as p.

*(*q) is the same as *p, which is same as x

Pointers to Structures

```
typedef struct {  
    int year;  
    char model;  
    float acceleration; //0-60mph in seconds  
}Car;
```

```
Car t1 = {.year = 2017, .model = 'S',  
    .acceleration = 2.8 };
```

```
Car * pt1 = &t1; //now you can use *pt1  
anywhere you use t1
```

```
(*pt1).acceleration = 2.3;  
(*pt1).year = 2019;  
(*pt1).model = 'X';  
float avg_acceleration = ((*pt1).acceleration  
+ (*pt2).acceleration) / 2.0;
```

We can also use the -> operator to access structure members.

```
pt1->acceleration = 2.3;  
pt1->year = 2019;  
pt1->model = 'X'  
float avg_acceleration = (pt1->acceleration +  
pt2->acceleration) / 2.0;
```

Address of (&) operator and Type

- Adding & to a variable adds * to its type
- Example:
 - if a is an int, then &a is an int*
 - if b is an int*, then &b is an int**
 - if c is an int**, then &c is an int***
 - ...

Dereference (*) operator and Type

- Adding * to a variable subtracts * from its type
- Example:
 - if a is an `int*`, then `*a` is an `int`
 - if b is an `int**`, then `*b` is an `int*`
 - if c is an `int***`, then `*c` is an `int**`
 - ...

Pointers to Functions (Function Pointers)

- Every function in a C program refers to a specific address (remember disassembling code during buffer overflow attack)
- Function pointers store addresses of functions
- Syntax:

```
typedef type (*name) (argument types)
```

Function Pointers - Example

```
typedef void (*myfuncptr) (int, int)
```

- `myfuncptr` is a pointer to a function that returns a `void` and accepts two arguments of type `int`.

Function Pointers - Example

```
void swap(int x, int y) {  
    int tmp = x;  
    x = y;  
    y = tmp;  
}
```

```
myfuncptr ptrswap = swap; //initialization.
```

```
int main(int argc, char* argv[]) {  
    int a=10;  
    int b=20;  
    ptrswap(a,b); //swap called by a function  
    pointer  
}
```

Function Pointers

How about these?

```
(*ptrswap)(a,b);
```

```
(****ptrswap)(a, b)
```

C says dereferencing a function pointer returns a function pointer. Behavior different from normal '&' and '' operators.*

Pointer Arithmetic

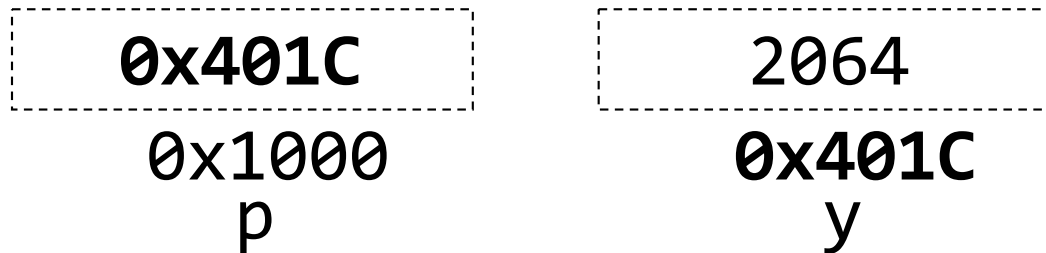
```
int y = 1040;  
int* p = &y;
```

- What does `*(p+1)` mean?
 - Data at “one element past” `p`
- What does “one element past” mean?
 - `p` is a pointer, so holds the address of a memory location
 - `p` is an `int` pointer, so that memory location holds an integer
 - `p+1` is interpreted as **address of the next integer**

Pointer Arithmetic

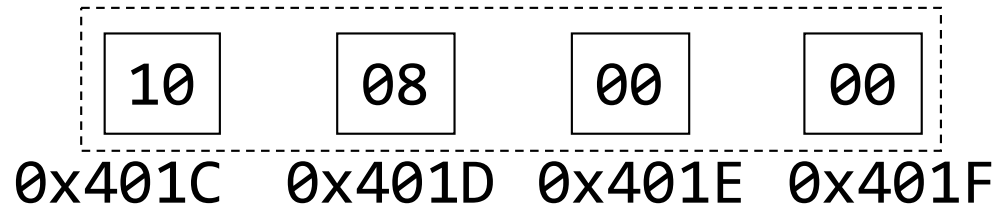
- Our representation of

```
int y=2064;  
int* p = &y;
```



Pointer Arithmetic

- ints occupy 4 bytes. 0x401C is the address of the first byte*:

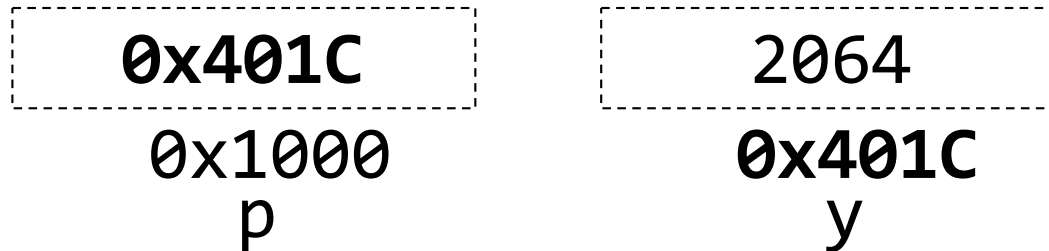


*2064 = 0x810 (=0x00,00,08,10 when written using 8 digits and x86 is little-endian)

- (*p) = data at 0x401C
 - returns the correct value of 2064 and not 0x10. Why?*

Pointer Arithmetic

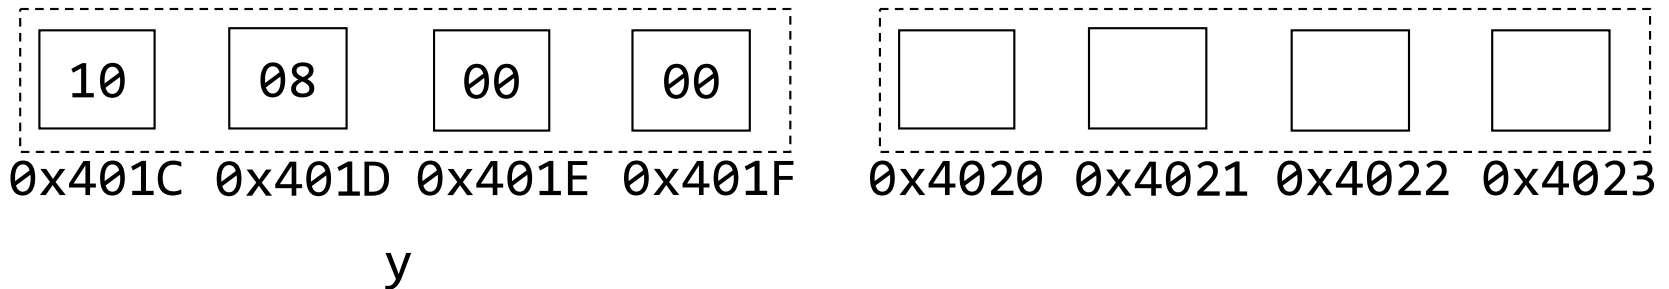
- $(p+1)$ gets the “address of the next integer”



What is the address of the next integer?

Pointer Arithmetic

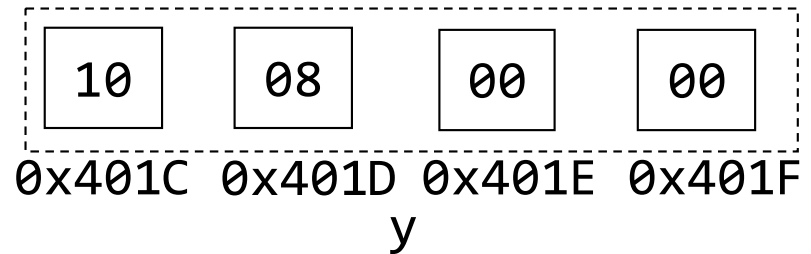
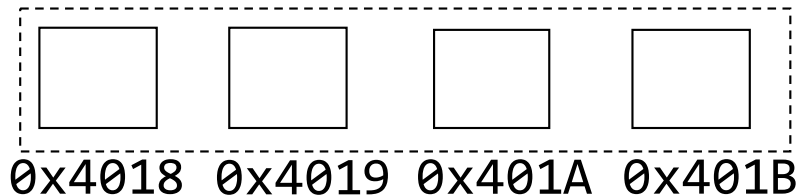
- What is the address of the next integer?
 - Add 4 to current value of p (0x401C) = 0x4020



Pointer Arithmetic

- $(p-1)$ computes the address before y

```
int y=2064;  
int* p = &y;
```



subtract 4 from the current value of p ($0x401C$) = $0x4018$

- Similarly we can add/subtract any number to/from a pointer variable.
- Compare to a specific address (E.g. `if(p == NULL)`)

Pointer Arithmetic

- Pointer to double (double occupies 8 bytes)

```
double pi=3.1428;  
double* ptrPi = &pi;
```

0x401C
0x1000
ptrPi

3.1428
0x401C
pi

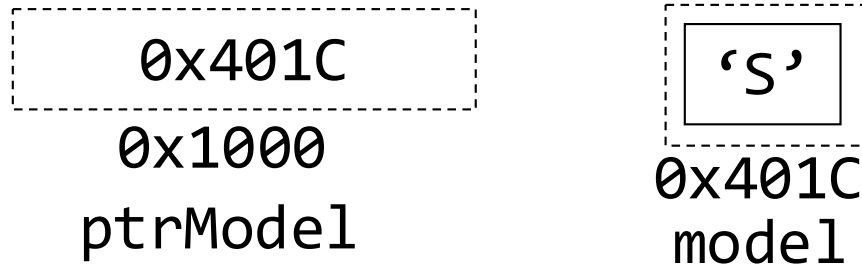
What is the address computed for (ptrPi+1)? 0x4024

What is the address computed for (ptrPi-1)? 0x4014

Pointer Arithmetic

- Pointer to char

```
char model='S';  
char* ptrModel = &model;
```



What is the address computed when we do `(ptrModel+1)`?

Pointer Arithmetic

- Pointer to pointer

```
char model='S';  
char* ptrModel = &model;  
char** doublePtr = &ptrModel;
```

0x1000

0x0500

doublePtr

0x401C

0x1000

ptrModel

'S'

0x401C
model

Bonus: what is the address computed when we do
(doublePtr+1)? (assuming we are using 32-bit machines)

Pointer Arithmetic

- Pointer to struct

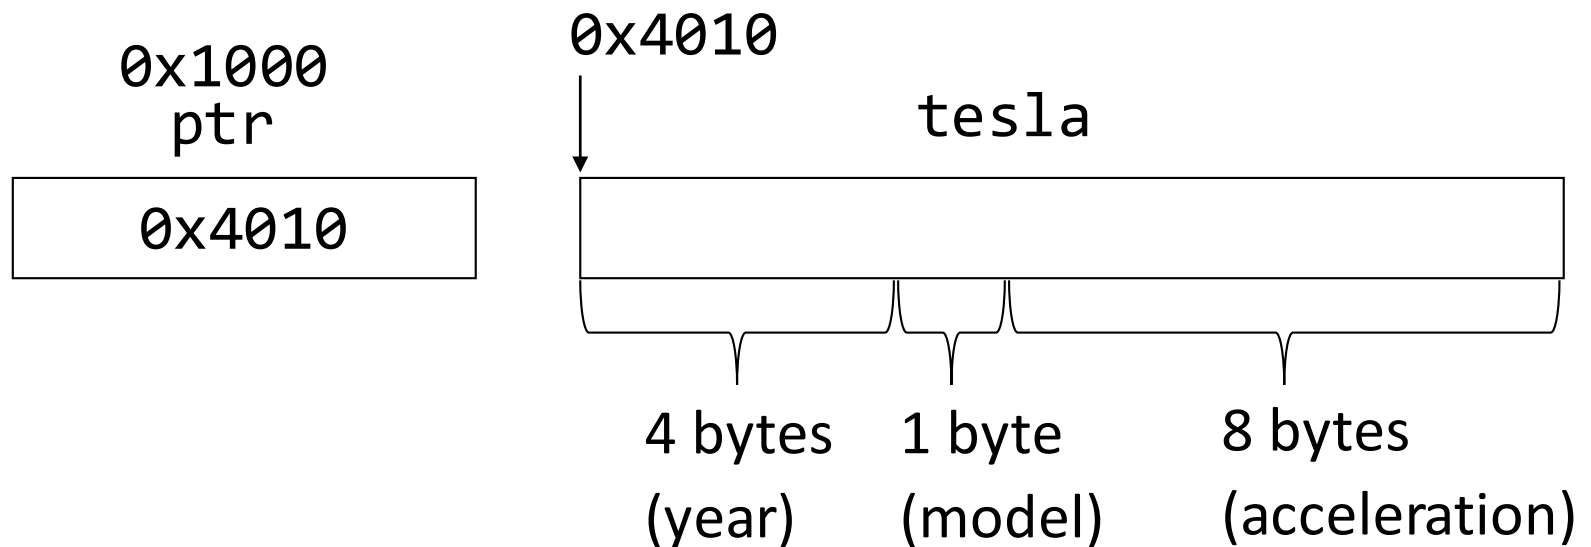
```
typedef struct {  
    int year;  
    char model;  
    double acceleration; //0-60mph in seconds  
}Car;
```

```
Car tesla = {.year = 2017, .model = 'S',  
    .acceleration = 2.8 };
```

```
Car* ptr = &tesla;
```

Pointer Arithmetic

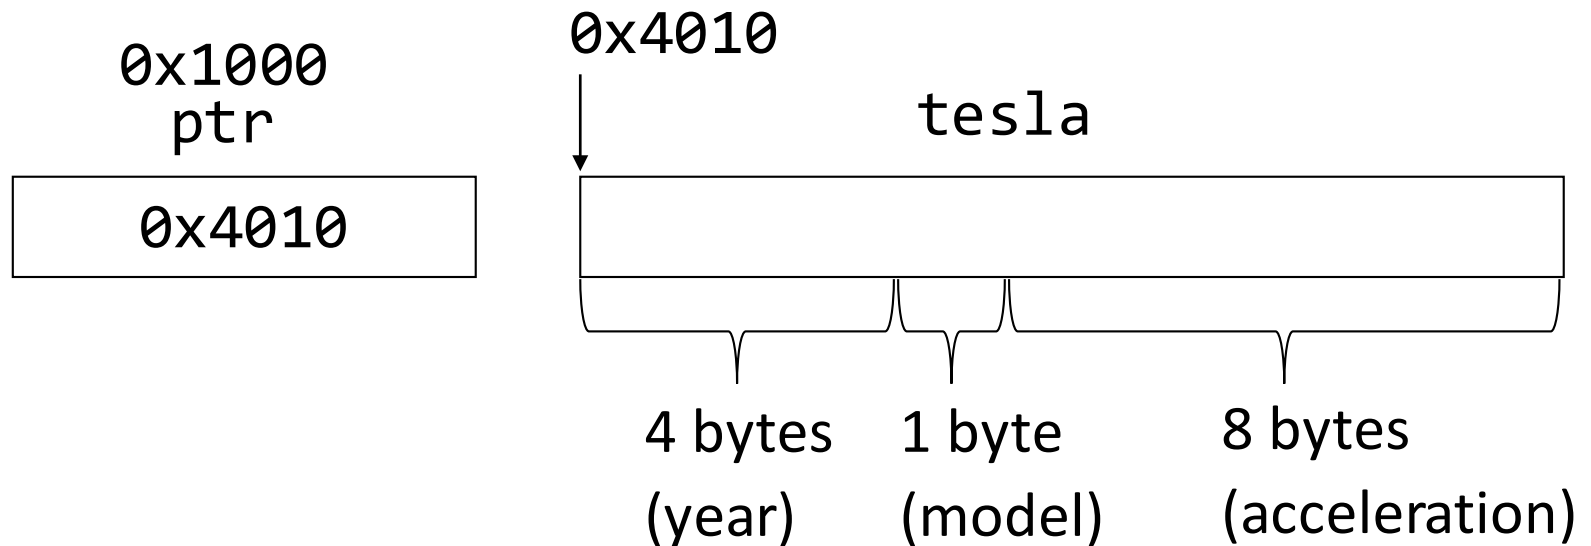
- Pointer to struct



- With `#pragma pack(1)`

Pointer Arithmetic

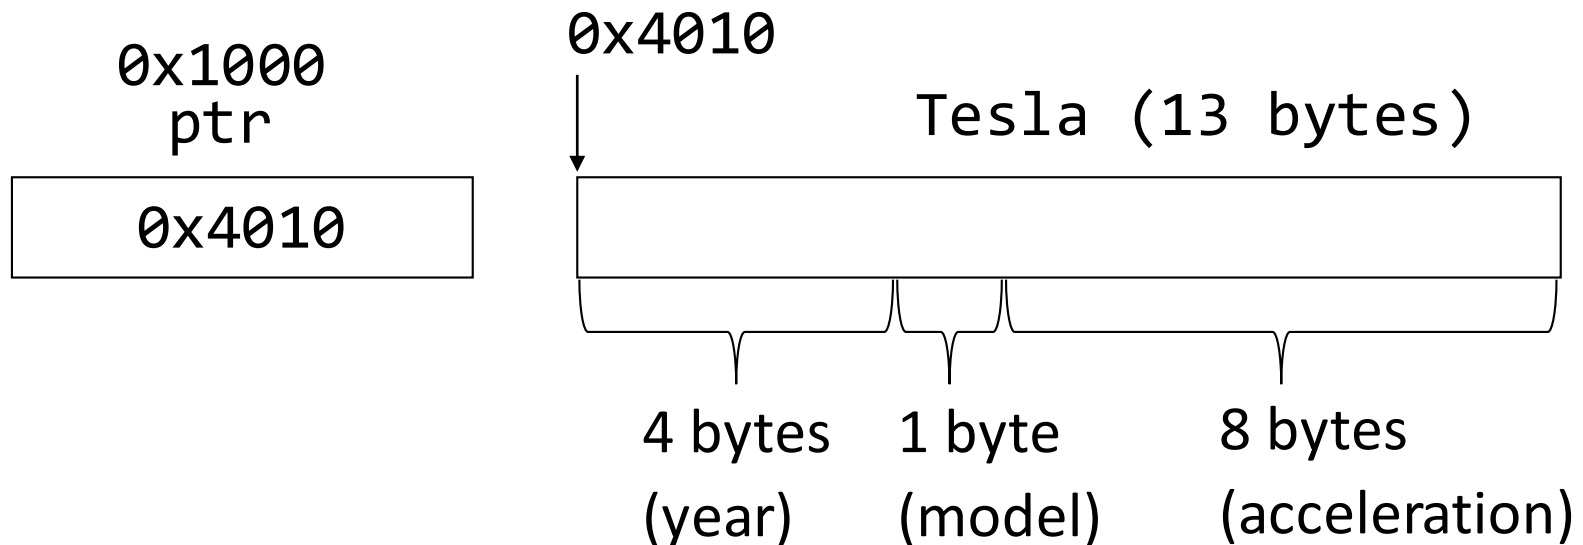
- What address does $(ptr+1)$ evaluate to?
 - Add 13 ($4+1+8$) to the value at ptr



- $ptr+1 = 0x401D$

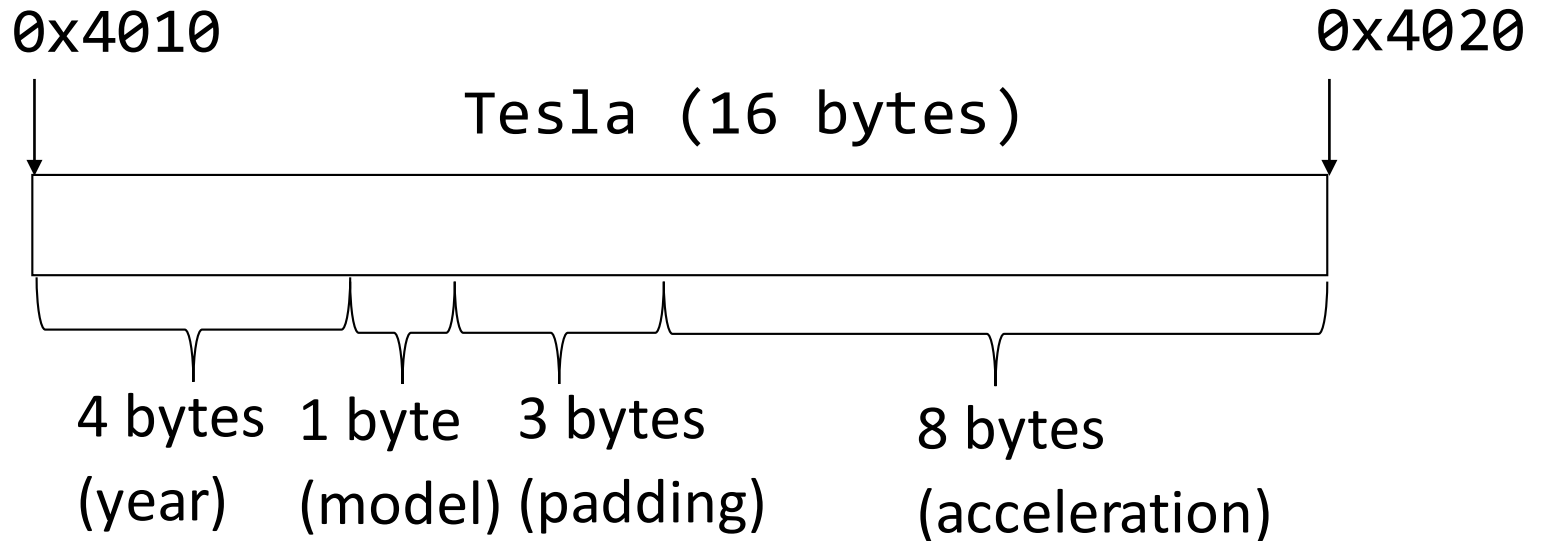
Detour - #pragma pack

- Preprocessor directive (starts with '#')
 - Preprocessor specifies instructions for the compiler on how to *pack* structure members in memory.
 - Varies from compiler to compiler



#pragma pack

- Normally (without #pragma pack) structure members are padded to create an alignment of the structure size with memory addresses.



Arrays

- Another data type!
 - Array of ints, structs etc.
 - Array of chars (strings in C)
- Work a little bit like pointers

```
int a[10]={1,2,3,4,5,6,7,8,9,10};  
//array of 10 integers
```

1	2	3	4	5	6	7	8	9	10
a[0]	a[1]	a[2]	a[3]	a[4]	a[5]	a[6]	a[7]	a[8]	a[9]

10 elements **guaranteed** to be next to each other in memory

Arrays

```
int a[10]={1,2,3,4,5,6,7,8,9,10};
```

1	2	3	4	5	6	7	8	9	10
a[0]	a[1]	a[2]	a[3]	a[4]	a[5]	a[6]	a[7]	a[8]	a[9]



a

0x4001

- 0x4001 is starting address of the array = address of a[0] = **&a[0]**
- Fetch the address of a = &a = 0x4001

Arrays

- Array name in C is the address of the first element of the array

```
int a[10]={1,2,3,4,5,6,7,8,9,10};
```

Therefore, `a == &a[0]`

a, &a, &a[0] are the same and have values 0x4001.

Arrays

- Array name in C is the address of the first element of the array

Array names are converted to pointers (in most cases) but a's type is not a pointer.

```
int* ptr=a; //ptr holds the address of the  
first element of the array (also &a[0]).
```

```
ptr[1] gets a[1]  
ptr[2] gets a[2]  
...
```

How is this possible?

Arrays

- Array dereferencing operator [] is implemented in terms of pointers.
 - $a[3]$ means: start at the address a , go forward 3 elements, fetch the *data at* that address.
 - In pointer arithmetic syntax, this is equivalent to:

$*(a+3)$

So,

$a[0]$ really means: $*(a+0)$

$a[1]$ really means: $*(a+1)$

Arrays

- So, when

```
int* ptr = a;
```

- `ptr[0]` really means `*(ptr+0)`, which is the same as `*(a+0)`, which is `a[0]`
- `ptr[1]` really means `*(ptr+1)`, which is the same as `*(a+1)`, which is `a[1]`

...

Exercise

```
char s[3] = "Hi";
```

```
char *t = "Si";
```

```
int u[3] = {5, 6, 7};
```

```
int n=8;
```

Expression	Type	Comments
s	char[3]	array of 3 chars
t	char*	address of a char
u	int[3]	array of 3 ints
&u[0]	int*	address of an int

Exercise

```
char s[3] = "Hi";
```

```
char *t = "Si";
```

```
int u[3] = {5, 6, 7};
```

```
int n=8;
```

Expression	Type	Comments
*&n	int	value at n
*t	char	data at address Held by t

Exercise

- Array initializers:

1. `int u[3] = {5, 6};`

Is this valid?

If yes, what is the value held in the third element `u[2]`?

2. `int u[3] = {5, 6, 7, 8};`

Is this valid?

3. `char s1[]="Hi";`

What is the size of `s1`? (how many bytes are reserved for `s1`)

4. `char s2[3]="Si";`

Is this valid?